

# Taming floating-point rounding errors with proofs

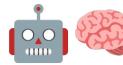
Laura Titolo  
Code Metal, USA

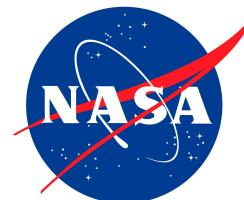
16th International Conference on Interactive Theorem Proving - ITP 2025  
Reykjavík, 28th September 2025

# Outline

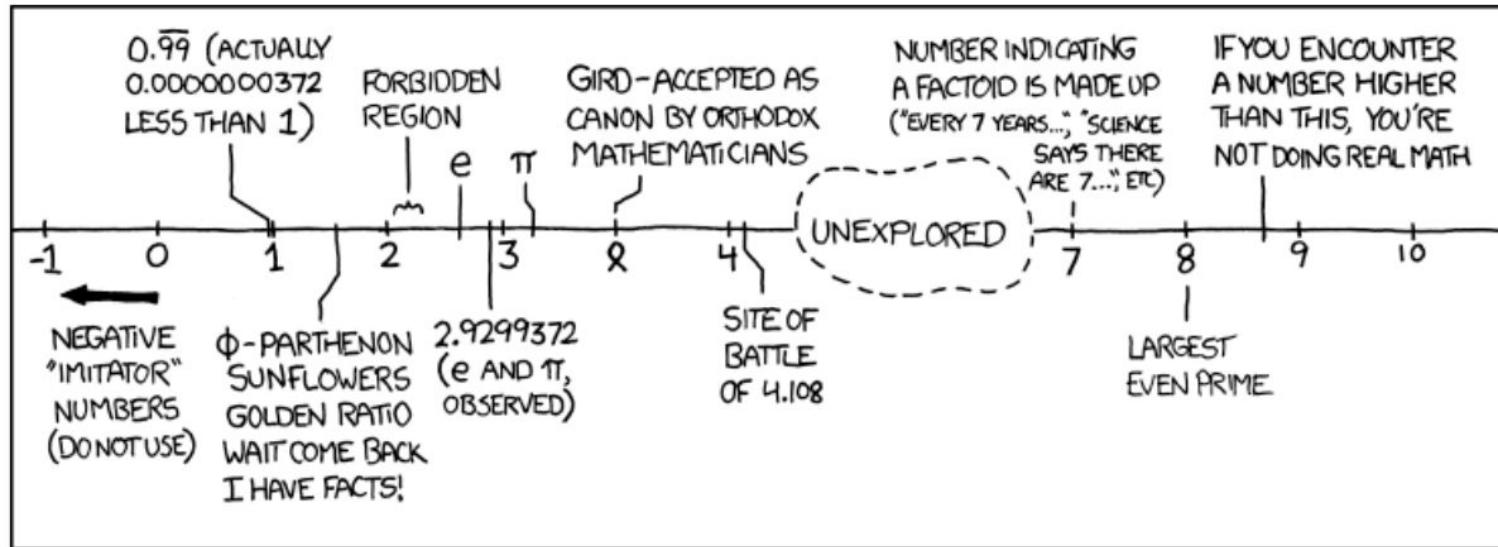
 application of  
ITP + other formal methods  
for **floating-point** program  
analysis

 ITP at **NASA**

 ITP for **AI** and AI for ITP in  
industry



# Representing Real Numbers in Digital Computers



Credits by XKCD

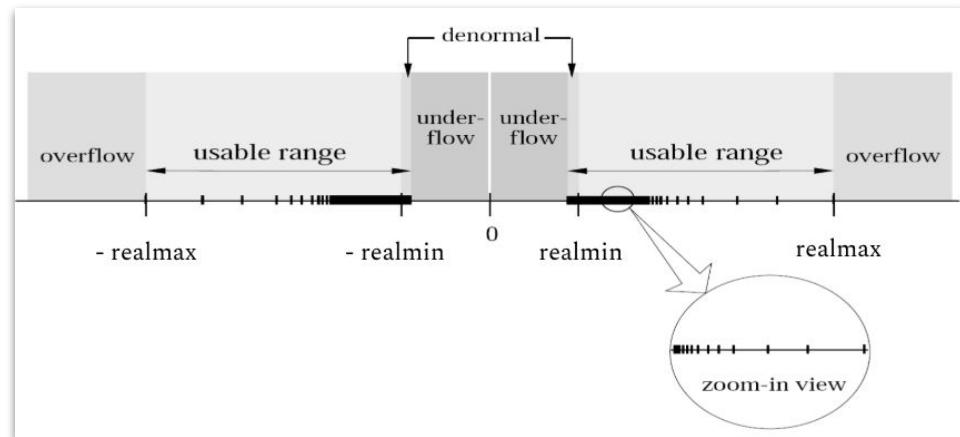
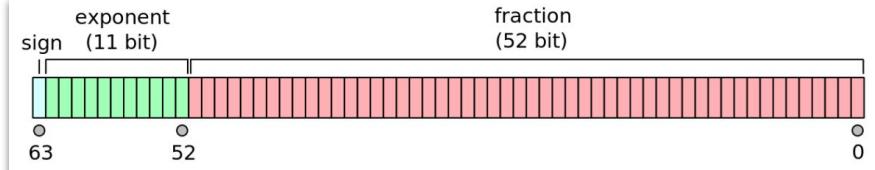
# Floating-Point Numbers and Rounding errors

 Floating-point (FP) = **finite representation** of real numbers

$\text{FP} \subset \text{Reals}$

Example: 0.1 is not a FP  
(rounded to 3602879701896397 /  
36028797018963968)

⚠ Round-off error =  $| r - \text{round}(r) |$



Credits by illinois.edu

# Writing Floating-Point Code is hard!

```
Prelude> (4/3 -1) * 3 - 1
```

# Writing Floating-Point Code is hard!

```
Prelude> (4/3 -1) * 3 - 1 ←  
-2.220446049250313e-16
```

Result is **0** if evaluated in  
exact real number  
arithmetic

# Writing Floating-Point Code is hard!

```
Prelude> (4/3 -1) * 3 - 1
-2.220446049250313e-16
Prelude> floor((4/3 -1) * 3 - 1)
-1
```

Result is **0** if evaluated in  
exact real number  
arithmetic

# Writing Floating-Point Code is hard!

```
Prelude> (4/3 -1) * 3 - 1
-2.220446049250313e-16
Prelude> floor((4/3 -1) * 3 - 1)
-1
Prelude> if (floor((4/3 -1) * 3 - 1)) < 0 then 100 else 1
100
Prelude> []
```

Result is **0** if evaluated in exact real number arithmetic

The Boolean guard is evaluated to **false** in real arithmetic and to **true** in floating-point arithmetic

# Writing Floating-Point Code is hard!

```
Prelude> (4/3 -1) * 3 - 1  
-2.220446049250313e-16  
Prelude> floor((4/3 -1) * 3 - 1)  
-1  
Prelude> if (floor((4/3 -1) * 3 - 1)) < 0 then 100 else 1  
100  
Prelude> □
```

Accumulated  
round-off error = 99



# Floating-point behavior is difficult to predict!



Floating-point numbers are ubiquitous



“Many **developers** do not understand core floating point behavior particularly well, yet believe they do.”

! Reals arithmetic  $\neq$  FP arithmetic  
(associativity, commutativity, ...)



Rounding errors, overflows, underflows...

2018 IEEE International Parallel and Distributed Processing Symposium

## Do Developers Understand IEEE Floating Point?

Peter Dinda      Conor Hetland  
*Northwestern University*

**Abstract**—Floating point arithmetic, as specified in the IEEE standard, is used extensively in programs for science and engineering. This use is expanding rapidly into other domains, for example with the growing application of machine learning everywhere. While floating point arithmetic often appears to be arithmetic using real numbers, or at least numbers in scientific notation, it actually has a wide range of gotchas. Compiler and hardware implementations of floating point inject additional surprises. This complexity is only increasing as different levels of precision are becoming more common and there are even proposals to automatically reduce program precision (reducing power/energy and increasing performance) when results are deemed “good enough.” Are software developers who depend on floating point aware of these issues? Do they understand how floating point can bite them? To find out, we conducted an anonymous study of different groups from academia, national labs, and industry. The participants in our sample did only slightly better than chance in correctly identifying key unusual behaviors of the floating point standard, and poorly understood which compiler and architectural optimizations were non-standard. These surprising results and others strongly suggest caution in the face of the expanding complexity and use of floating point arithmetic.

**Keywords**—floating point arithmetic, software development, user studies, correctness, IEEE 754

bounds [13], and approximate computing [9] where performance/energy and output quality can be traded off. More immediately, any programmer using a modern compiler is faced with dozens of flags that control floating point optimizations which could affect results. Optimizing a program across the space of flags has itself become a subject of research [5].

The floating point arithmetic experienced by a software developer via a particular hardware implementation, language, and compiler, is swelling in complexity at the very same time that the demand for such developers is also growing. This may be setting the stage for increasing problems with numeric correctness in an increasing range of programs. Numeric issues can produce major effects. Recall that Lorenz’s insight, a cornerstone of chaos theory, was triggered by a seemingly innocuous rounding error [10]. Arguably, modern applications, certainly those that model systems with chaotic dynamics, could see small errors in developer understanding of floating point become amplified into bad overall results.

Do software developers understand the core properties of floating point arithmetic? Do they grasp which optimizations might result in non-compliance with the IEEE standard?

# Floating-point errors may be dangerous and costly

## Patriot missile failure (1991)

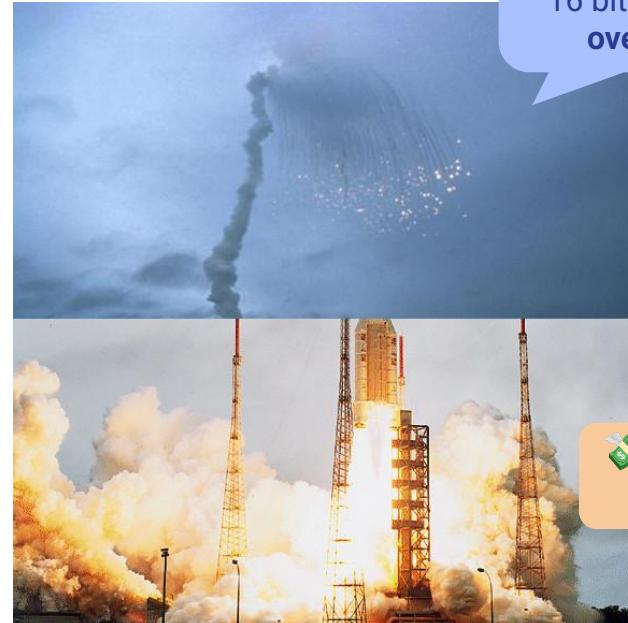


Accumulated **round-off error** in the time window calculation (1/10 constant used)

28 soldiers were killed and 100 injured

Picture: Bernd vdB, Public domain, via Wikimedia Commons

## Ariane V (1996)



64 bits floating-point to 16 bits integer **overflow**

US\$370 million loss

Photographs credits: 1996 © ESA/CNES; 1996 © ESA; 1996 © ESA.

# Automatic Dependent Surveillance - Broadcast (ADS-B)

- **NextGen** (Next generation of air traffic management systems) to enhance radar technologies
- Aircraft periodically **broadcasts** surveillance information
- **Automatic** – no pilot intervention
- **Mandatory** from Jan 1, 2020 (in USA and Europe)
- **Thousands of aircraft** currently equipped with ADS-B



# ADS-B: Pro and Cons

## ✓ Pros: Broadcast vs. radar

- 🎯 More precise

- 🌐 More coverage



## ✗ Cons: Makes use of **existing hardware**

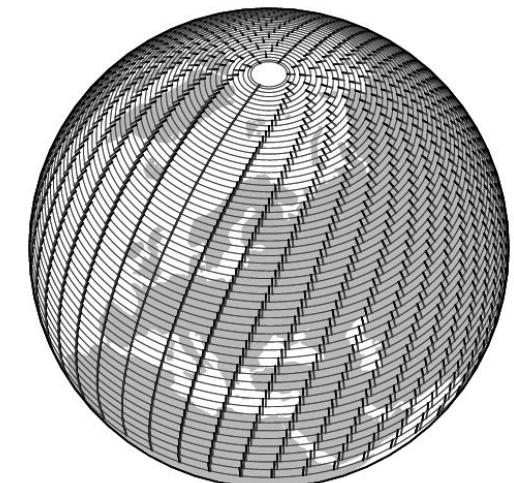
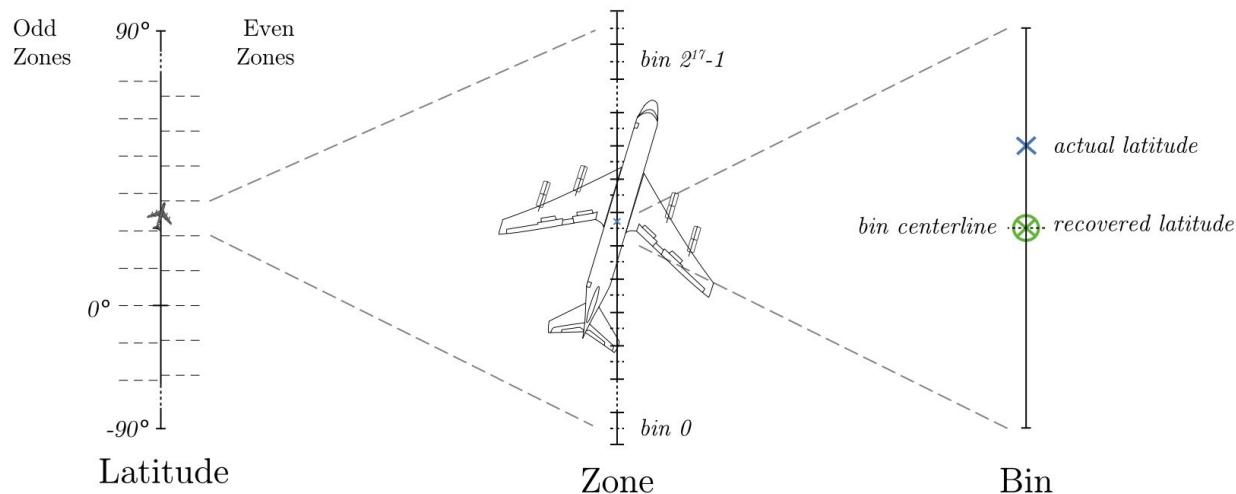
- 📡 TCAS transponders

- 📏 35 bits for position in the broadcast message (lat + lon)

- 🌐 Too coarse granularity (~300 mt) for raw positions

# ADS-B CPR: Compact Position Reporting Algorithm

- ➊ Divide the globe into 59 (odd) or 60 (even) equally sized **zones**
- ➋ Divide each zone in  $2^{17}$  **bins**
- ➌ Zone number + Bin number = Position  $\pm \approx 5.1\text{mt}$



# CPR Encoding

$$lat\_enc(i, lat) = \text{mod} \left( \left\lfloor 2^{17} \frac{\text{mod}(lat, dlat_i)}{dlat_i} + \frac{1}{2} \right\rfloor, 2^{17} \right)$$

 The encoding computes the **bin number**

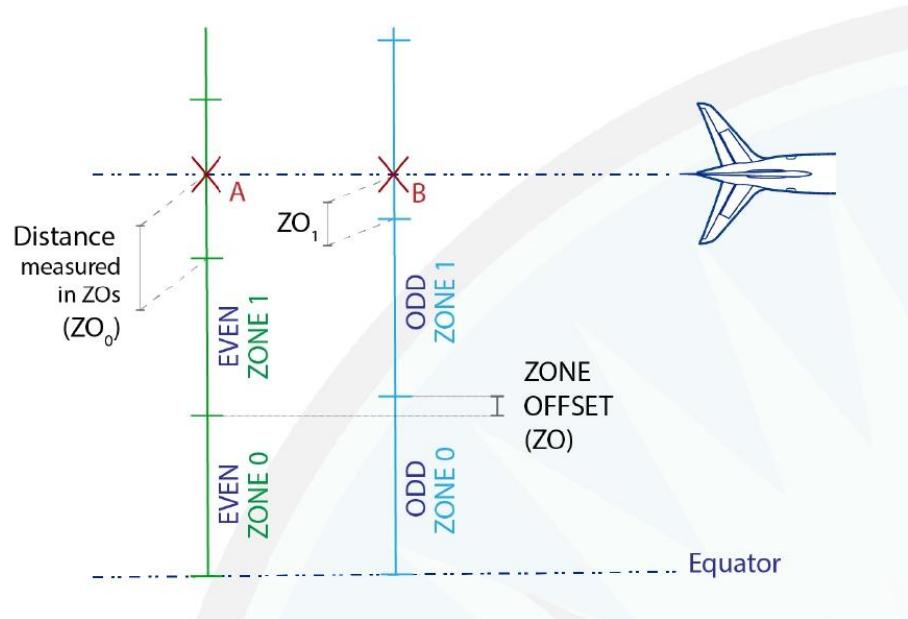
 Only the bin numbers and encoding type are transmitted

 Not the zone!



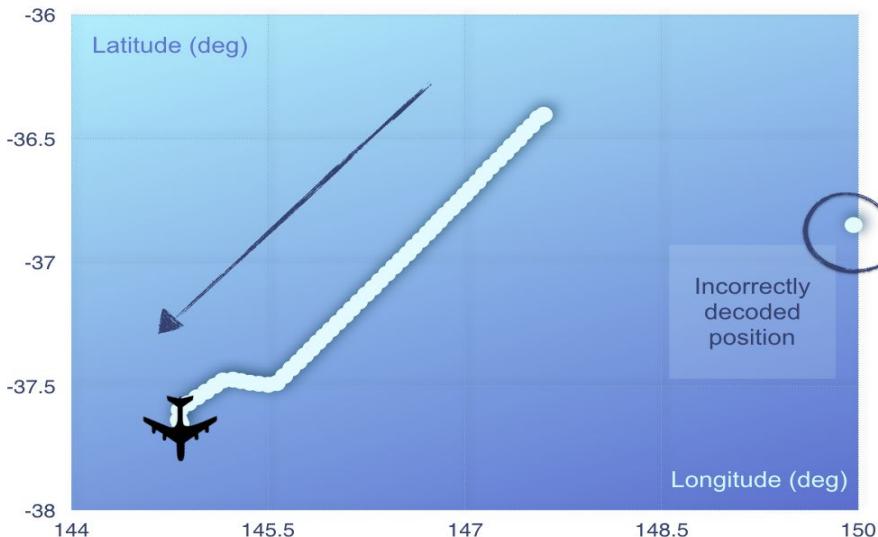
# CPR Decoding

- Both **zone** and **bin number** are necessary to recover the position
- CPR needs a **reference position or an odd and an even message sent sufficiently close to each other to compute the decoded zone**
- **Zone offset number** determines the zone number



# CPR issues

## ADS-B Compact Position Reporting Algorithm - CPR (2016)



Satellite dish icon | Situation as reported by Airservices Australia | Airplane icon

| actual position - decoded position  
|  
= 241 nautical miles  
= 446 km  
= 277 miles

# What was the issue with CPR?



Dutle A., Moscato M.,  
Titolo L., Muñoz C. A  
*Formal Analysis of the  
Compact Position  
Reporting Algorithm.*  
VSTTE 2017.

**Requirements** were **not enough** to guarantee the intended precision even assuming exact real-number arithmetic.

We proposed a slightly **tightened set of requirements** and formally proved the algorithm correct in PVS **assuming exact arithmetic**

$| \text{original\_position} \text{ } \text{📍} - \text{recovered\_position} \text{ } \text{💻} | \leq 5.1\text{mt}$  = algorithm accuracy 

What about **finite-precision** implementations?

 Heavy use of **modulus** and **floor**  $\Rightarrow$  huge accumulated round-off error

 Using single-precision floating-point, the recovered position of was off by approx **1500 nautical miles!**

# Formally verified finite-precision implementation of CPR

- Propose simpler formulation reducing numerical complexity (**CPR\***)
- Use a suite of existing **formal methods** tools to provide a verified and correct implementation of CPR\* with finite precision arithmetic:

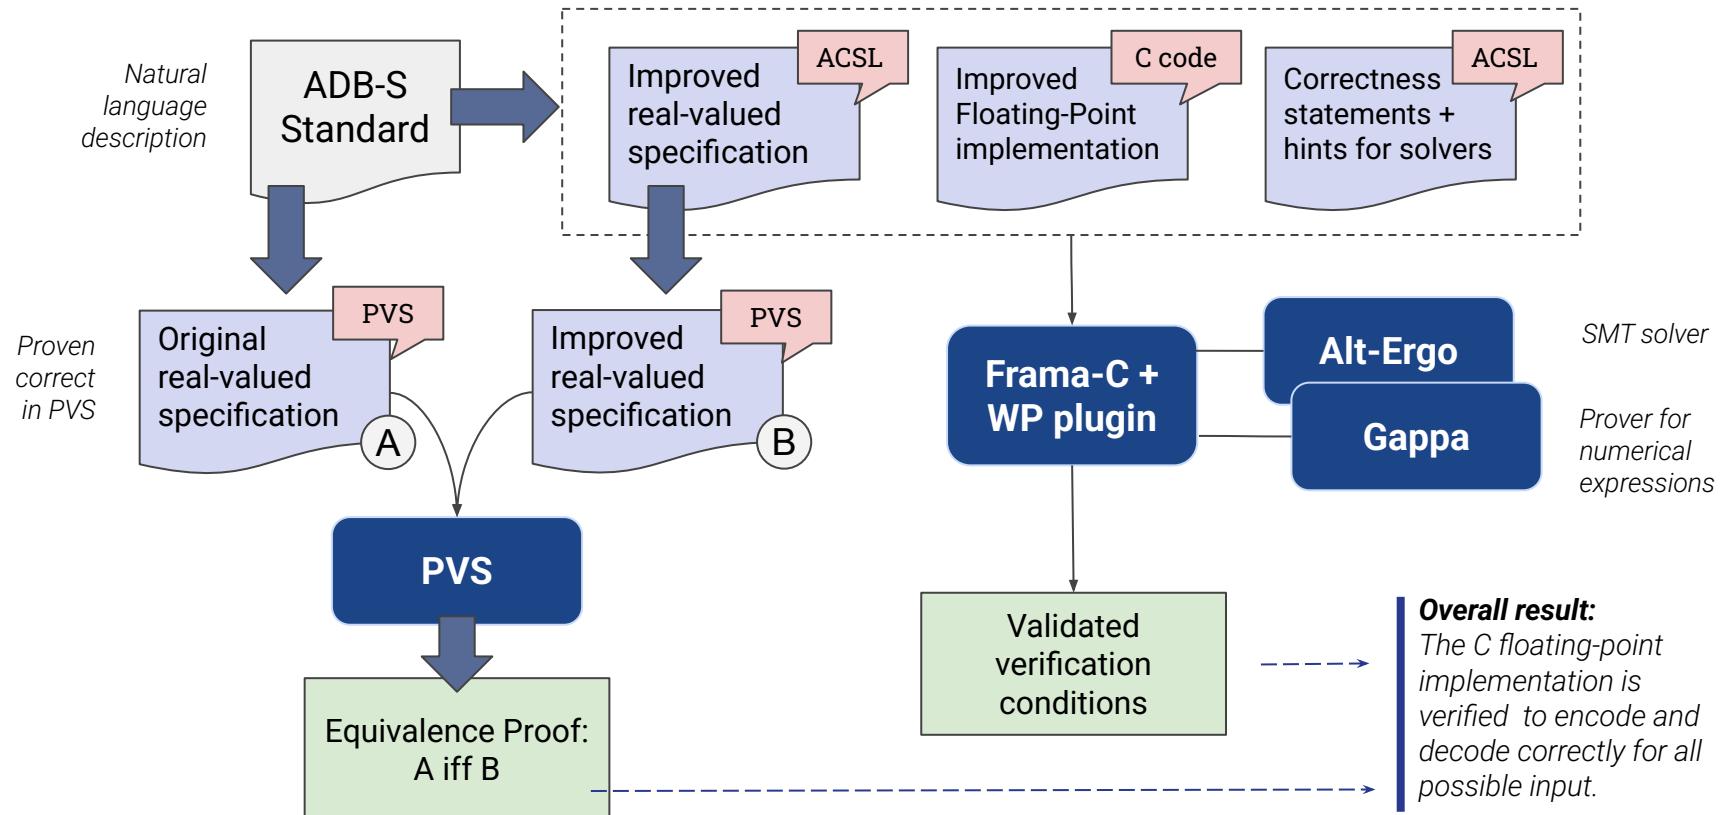


Software Analyzers



**Titolo L., Moscato M., Muñoz C., Dutle A., Bobot F. A**  
*Formally Verified Floating-Point Implementation of the Compact Position Reporting Algorithm.* FM 2018.

# CPR floating-point implementation verification

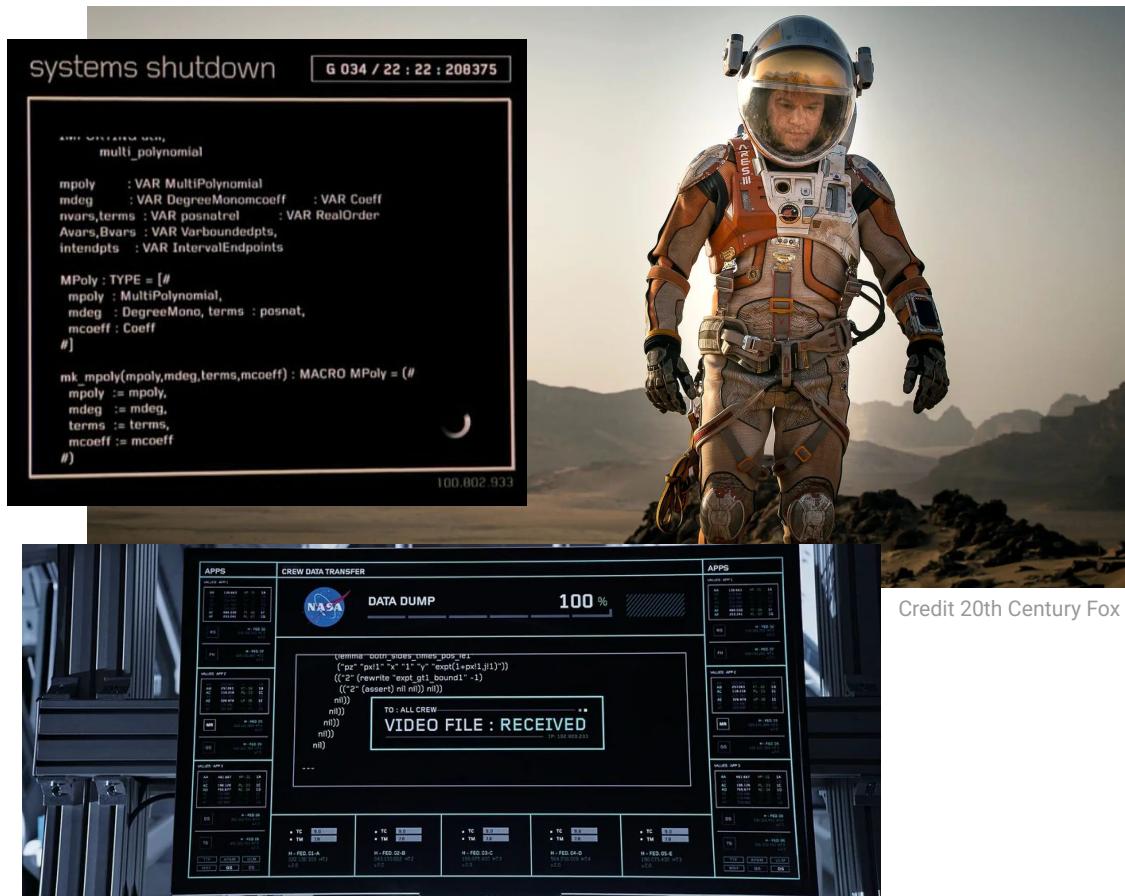


# Why PVS?

 NASA PVS Library has been developed for about 4 decades and contains 79 libraries and 40K theorems

 PVS is the only theorem prover featured in a Hollywood movie

 PVS actually used to verify Mars Rovers' plans (PLEXIL-V)



Credit 20th Century Fox

# CPR\*

✓ CPR\* = Formally verified C implementation of CPR

- double-precision floating-point
- 32 bits unsigned integers

Reference implementation in the revised ADS-B standard document (RTCA DO-260B/Eurocae ED-102A)

🏆 NASA group achievement award 2020

Available at <https://github.com/nasa/cpr>

The screenshot shows a GitHub repository page for 'cpr'. The commit history lists several commits by 'Cesar Munoz' from 5 years ago, including adding summaries, test procedures, and initial releases for C, Gappa, and LICENSES. The README file contains the title 'CPR\*: Formally Verified Compact Position Reporting Algorithm' and a description of the RTCA-DO-260B/Eurocae ED-102A standard.

File / Commit	Description	Date
Cesar Munoz Added summaries	ce509f8 · 5 years ago	17 Commits
C	Add new test procedures from MOPS.	6 years ago
Gappa	Initial release	6 years ago
LICENSES	Initial release	6 years ago
PVS	Added summaries	5 years ago
README.md	Update README.md	6 years ago

**CPR\*: Formally Verified Compact Position Reporting Algorithm**

The Compact Position Reporting (CPR) algorithm consists of a set of functions defined in the standard RTCA-DO-260B/Eurocae ED-102A, Minimum Operational Performance Standards for 1090 MHz extended squitter Automatic Dependent Surveillance - Broadcast (ADS-B) and Traffic Information Services - Broadcast (TIS-B). These function encode and decode aircraft positions. CPR\* is a formally verified C implementation of CPR's functions that use computer arithmetic in fixed- and floating-point formats.



RTCA

EUROCAE

# CPR\* Verification approach

 Successful combination of different formal methods tools

 The CPR verification pipeline required a lot of expertise and manual effort!

 **Manual annotation** of the C code

 **Gappa** needs **hints** to verify a tight rounding error

 A lot of **time and effort** spend in writing proofs in PVS

 Expertise in floating-point arithmetic and ITP



Goal: Make this process **fully automatic** 

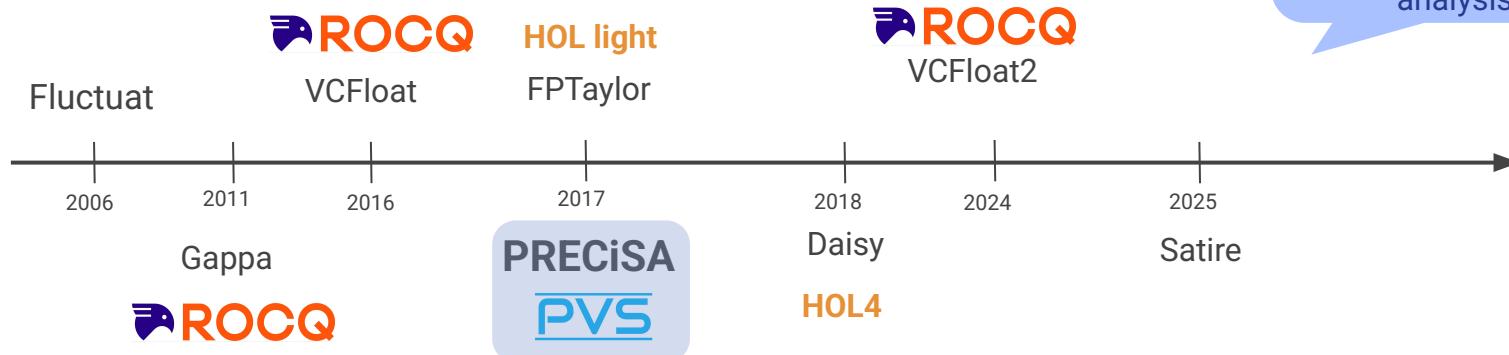
# Automatic floating-point rounding error analysis

✓ Formally verified and sound

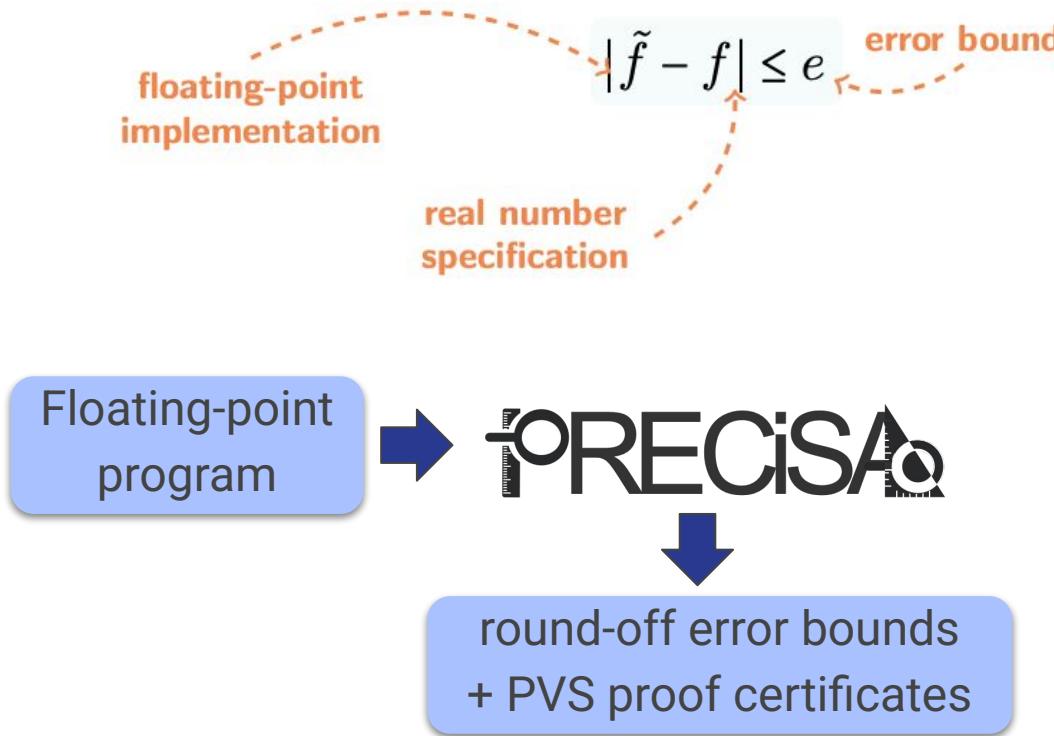
✗ Sound treatment of conditionals

🔍 Accurate

Now many tools are available for floating-point round-off error analysis



# PRECiSA static analyzer



Moscato M., Titolo L., Dutle A., Muñoz C.  
Automatic Estimation of Verified Floating-Point Round-Off Errors via Static Analysis. SAFECOMP 2017.

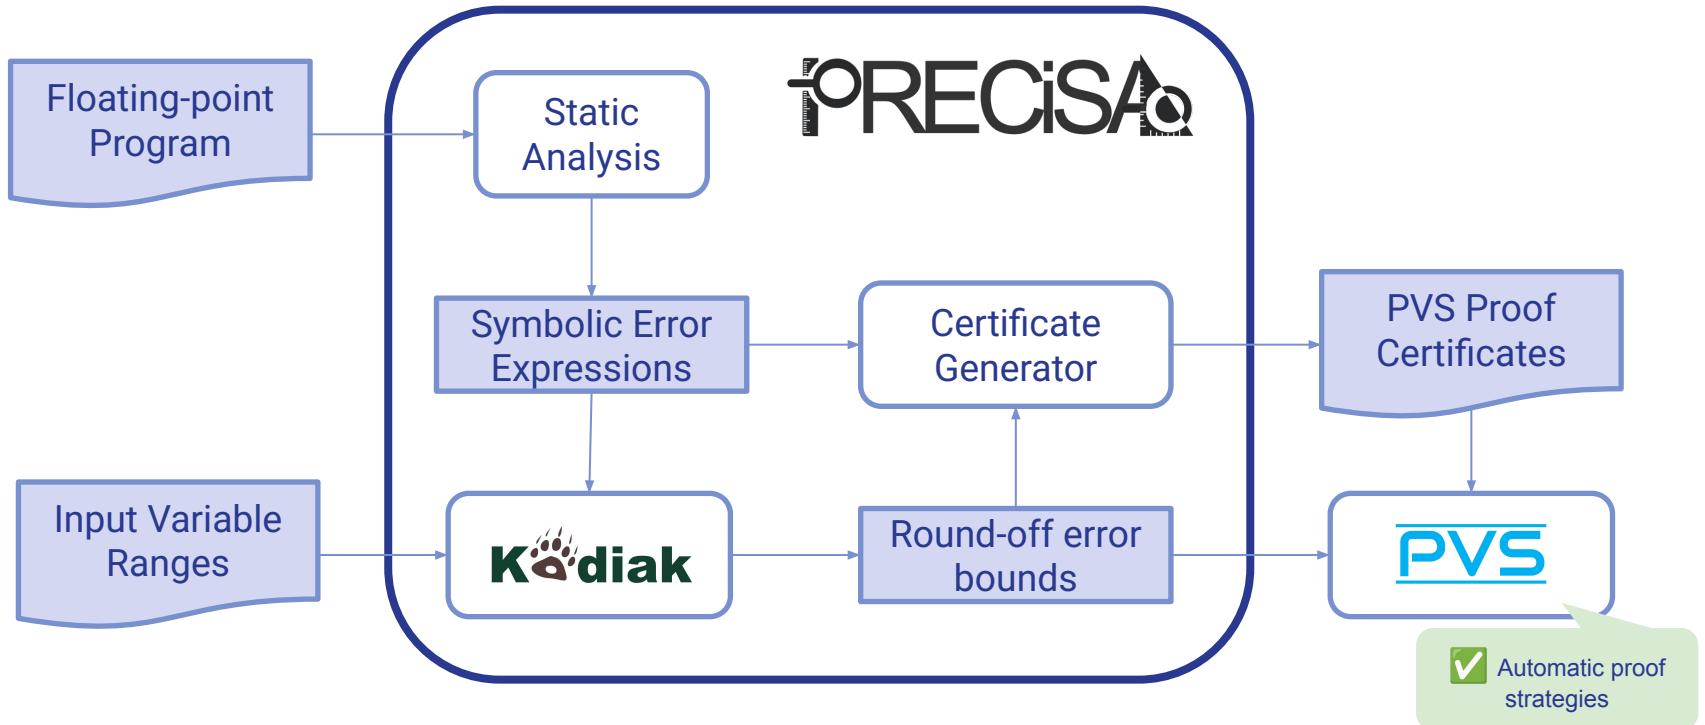


Titolo L., Feliú M., Moscato M., Muñoz C.  
*An Abstract Interpretation Framework for the Round-Off Error Analysis of Floating-Point Programs.* VMCAI 2018.



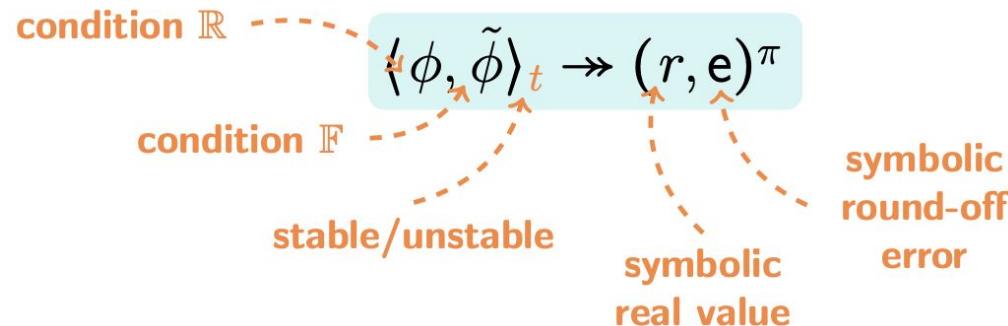
Titolo L., Moscato M., Feliú M., Masci P., Muñoz C.  
*Rigorous Floating-Point Round-Off Error Analysis in PRECiSA 4.0.* FM 2024.

# PRECiSA Workflow



# PRECiSA - Step I: Static Analysis

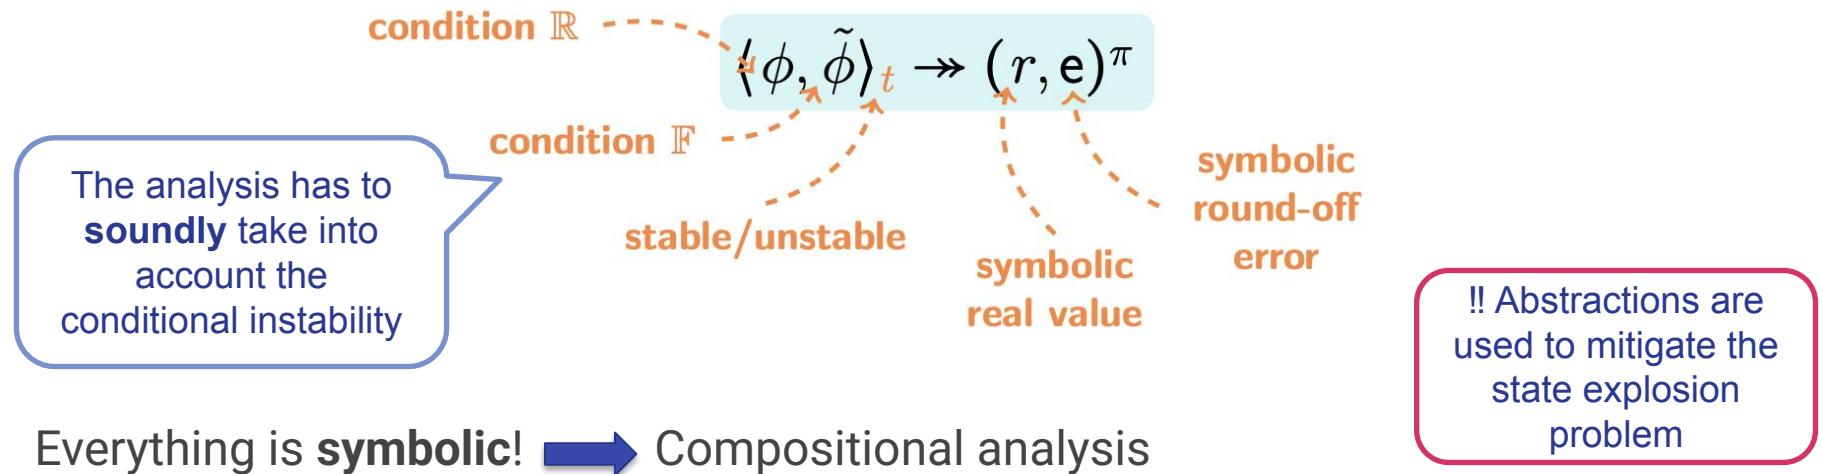
For each function declaration PRECiSA computes a set of **conditional error bounds**



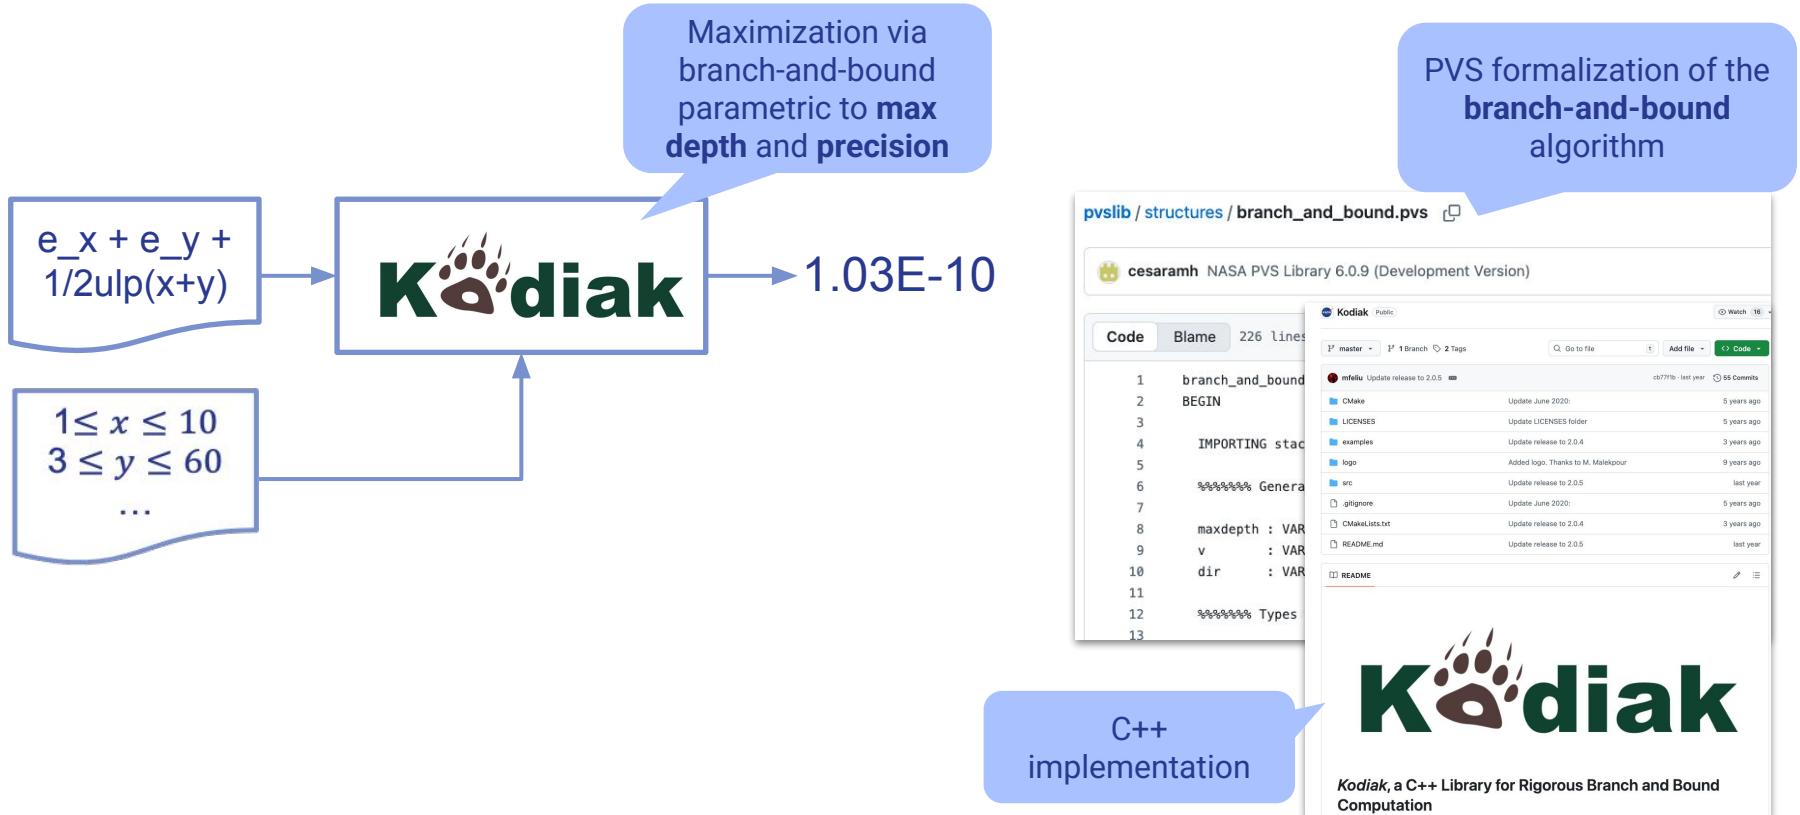
Everything is **symbolic!** → Compositional analysis

# PRECiSA - Step I: Static Analysis

For each function declaration PRECiSA computes a set of **conditional error bounds**



# PRECiSA - Step II: Global Optimization



# PRECiSA Step III: Proof certificate generation

```
mul_aerr_bound(r1,e1,r2,e2)
: nonneg_real
= abs(r1)*e2+abs(r2)*e1+e1*e2
```

Accumulated round-off error

```
% Error of the (correctly-rounded) function w.r.t. the real operation on
% the (real version of the) floating point numbers.
```

```
prove | discharge-tccs | status-proofchain | show-prooflite
```

```
mul_accum_err: LEMMA
```

```
abs(FtoR(f1)-r1) <= e1 AND
abs(FtoR(f2)-r2) <= e2
IMPLIES abs((FtoR(f1)*FtoR(f2))-(r1*r2)) <= mul_aerr_bound(r1,e1,r2,e2)
```

```
mul_ulp_bound(r1,e1,r2,e2)
```

```
: real
```

```
= abs(r1)*abs(r2) + abs(r1)*e2 + e1*abs(r2) + e1*e2
```

```
prove | discharge-tccs | status-proofchain | show-prooflite
```

```
Fmul_accum_err_bound: LEMMA
```

```
abs(FtoR(f1)-r1) <= e1 AND
abs(FtoR(f2)-r2) <= e2
IMPLIES abs(FtoR(f1)*FtoR(f2)) <= mul_ulp_bound(r1,e1,r2,e2)
```

```
prove | discharge-tccs | status-proofchain | show-prooflite
```

```
accum_err_bound: LEMMA
```

```
abs(FtoR(f1)-r1) <= e1 AND
abs(FtoR(f2)-r2) <= e2
IMPLIES abs(FtoR(Fmul(f1,f2)) - (r1*r2))
    <= mul_aerr_bound(r1,e1,r2,e2)
    + ulp(b)(mul_ulp_bound(r1,e1,r2,e2)) / 2
```

Rounding of the result

PRECiSA / PVS / PRECiSA /

lauratitolo v4.0 · pvslib / float / README.md

Cesar Munoz [nasalib] Updated some broken URL links

Preview Code Blame 179 lines (122 loc) · 9 KB

## Floating-Point Library

This library contains several formalizations of floating-point numbers.

- float\_bounded\_axiomatic: An axiomatic formalization that complies with the IEEE 754 standard. It provides a foundational formalization following the IEEE 854 Standard. This is a level foundational specification which can be instantiated into a level foundational formalization representing floating-point numbers. This specification is equivalent to the one for IEEE 854.

PRECiSA strategies  
<https://github.com/nasa/PRECiSA/tree/master/PVS/PRECiSA>

IEEE 745 and 854 + round-off error formalization  
<https://github.com/nasa/pvslib/blob/master/float>

# PRECiSA Step III: Proof certificate generation

Users > Ititolo > Desktop > precisa > public > benchmarks > code-generation > daidalus > tcoa > tcoa\_num\_cert.pvs > tcoa\_num\_ce

```
1 % This file is automatically generated by PRECiSA
2
3 % maxDepth: 7 , prec: 10^-14
4
5 typecheck-file | evaluate-in-pvsi
6 tcoa_num_cert: THEORY
7 BEGIN
8 IMPORTING cert_tcoa, PRECiSA@bbiasp, PRECiSA@bbiadp, PRECiSA@strategies
9
10 %|- *_TCC* : PROOF
11 %|- (precisa-gen-cert-tcc)
12 %|- QED
13
14 % Floating-Point Results: 0, neg_double(div_double(sz, vz))
15 % Real Results: -(r_sz / r_vz), 0
16 % Control Flow: Stable
17 prove | status-proofchain | show-prooffile
18 tcoa_fp_c_0 : LEMMA
19 FORALL(r_sz, r_vz: real, sz: double, vz: double):
20 abs(safe_prjct_double(sz) - r_sz)<=ulp_dp(r_sz)/2 AND abs(safe_prjct_double(vz) - r_vz)<=ulp_dp(r_vz)/2
21 AND (((r_vz /= 0) AND ((r_sz * r_vz) < 0)) AND ((vz /= integerToDouble(0)) AND (mul_double(sz, vz) <
integerToDouble(0))) OR (NOT((r_sz * r_vz) < 0)) AND NOT((mul_double(sz, vz) < integerToDouble(0))))
22 AND r_sz ## [|1,1000|] AND r_vz ## [|1,1000|] AND
finite_double?(tcoa_fp(sz, vz)) AND finite_double?(sz) AND finite_double?(vz) AND finite_double?(mul_double
(sz, vz)) AND finite_double?(integerToDouble(0))
23 IMPLIES|
24 abs(safe_prjct_double(tcoa_fp(sz, vz)) - tcoa(r_sz, r_vz))<=6876345720111303 / 2475880078570760549798248448
25
26 %|- tcoa_fp_c_0 : PROOF
27 %|- (prove-concrete-lemma tcoa_fp_0 14 7)
28 %|- QED
```

Initial ranges

Rounding error

Automatic strategy

# VSCode-PRECiSA: Round-off Error Analysis GUI

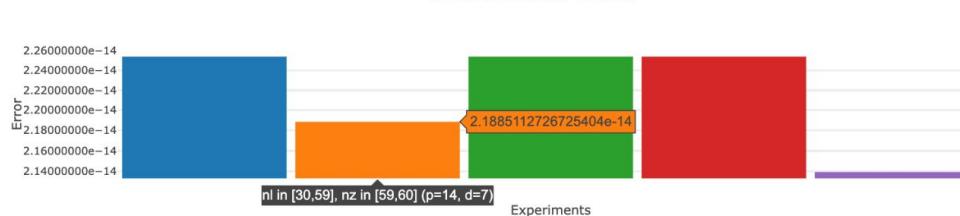
```
nl_comp.pvs > nl_comp
typecheck-file | evaluate-in-psio
nl_comp: THEORY
BEGIN
IMPORTING float@double64
%
% @fp-function
% @fp-range nl in [2,59], nz in [59,60]
estimate-error-bounds | compare-error-bounds
nl_comp(nl,nz: double): double =
(180/3.14) * acos(sqrt((1-cos(3.14/(2*nz)))/(1-cos(2*3.14/nl))))
END nl_comp
```

INPUT	MIN	MAX
nl	50	59
nz	60	

Analyze

Save Results

Accumulated Round-Off Error



- Interval Analysis
- Sensitivity Analysis
- Comparative Analysis

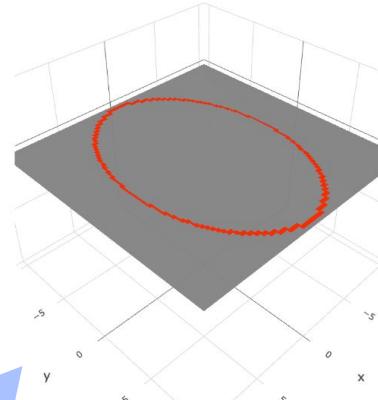
Experiment	Accumulated Round-Off Error
nl in [2,59], nz in [59,60] (p=14, d=7)	2.2540032743878307e-14
nl in [30,59], nz in [59,60] (p=14, d=7)	2.1885112726725404e-14
nl in [2,30], nz in [59,60] (p=14, d=7)	2.2540013434075158e-14
nl in [2,5], nz in [60,60] (p=14, d=7)	2.2540007151689644e-14
nl in [50,59], nz in [60,60] (p=14, d=7)	2.139337646609609e-14

# VSCode-PRECiSA: Instability Analysis GUI

```
pvs ellipse.pvs > ...
typecheck-file | evaluate-in-pvsi
1 ellipse: THEORY
2 BEGIN
3 IMPORTING float@aerr754dp
4
5 % @fp-function
6 % @fp-range x in [1,300], y in [1,300]
estimate-error-bounds | compare-error-bounds
7 pointInEllipse(x,y: double): double =
8   IF x*x/4 + y*y/9 <= 10 THEN 1 ELSE -1 ENDIF
9
10 END ellipse
```

INPUT	MIN	MAX
x	-10	10
y	-10	10

Analyze



Visualize the values that may cause divergences in the control flow using Kodiak's paving functionality

INPUT	MIN	MAX
x	0	5
y	0	5

Analyze

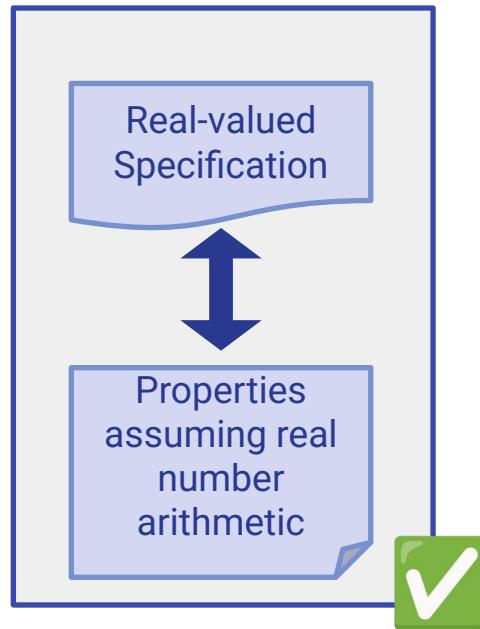
Paving results indicate all Ok!

— This means that the function does not contain unstable guards

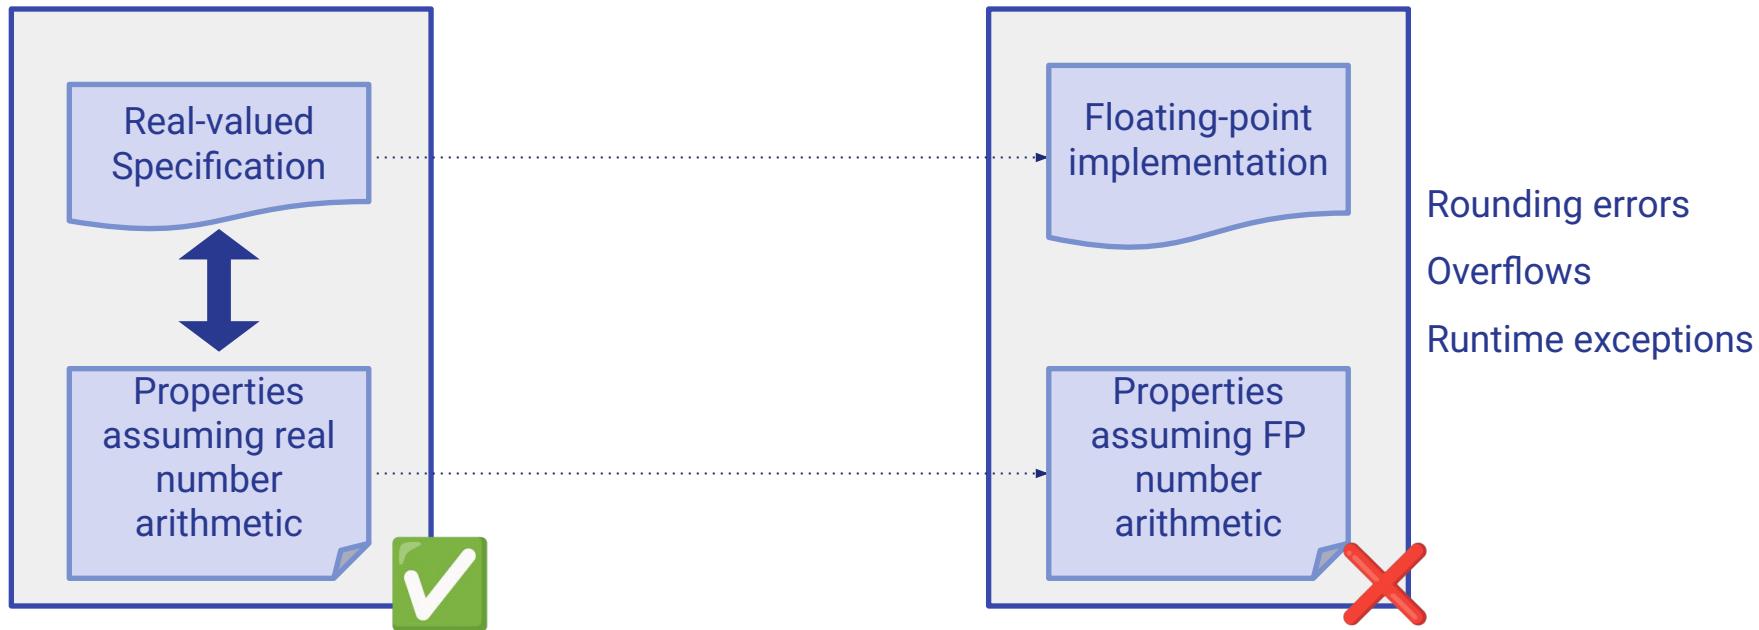


Check that no conditional instability occurs

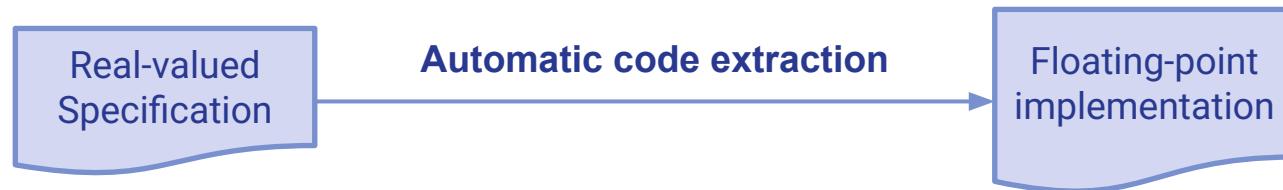
# Real number $\neq$ Floating-point arithmetic



# Real number $\neq$ Floating-point arithmetic



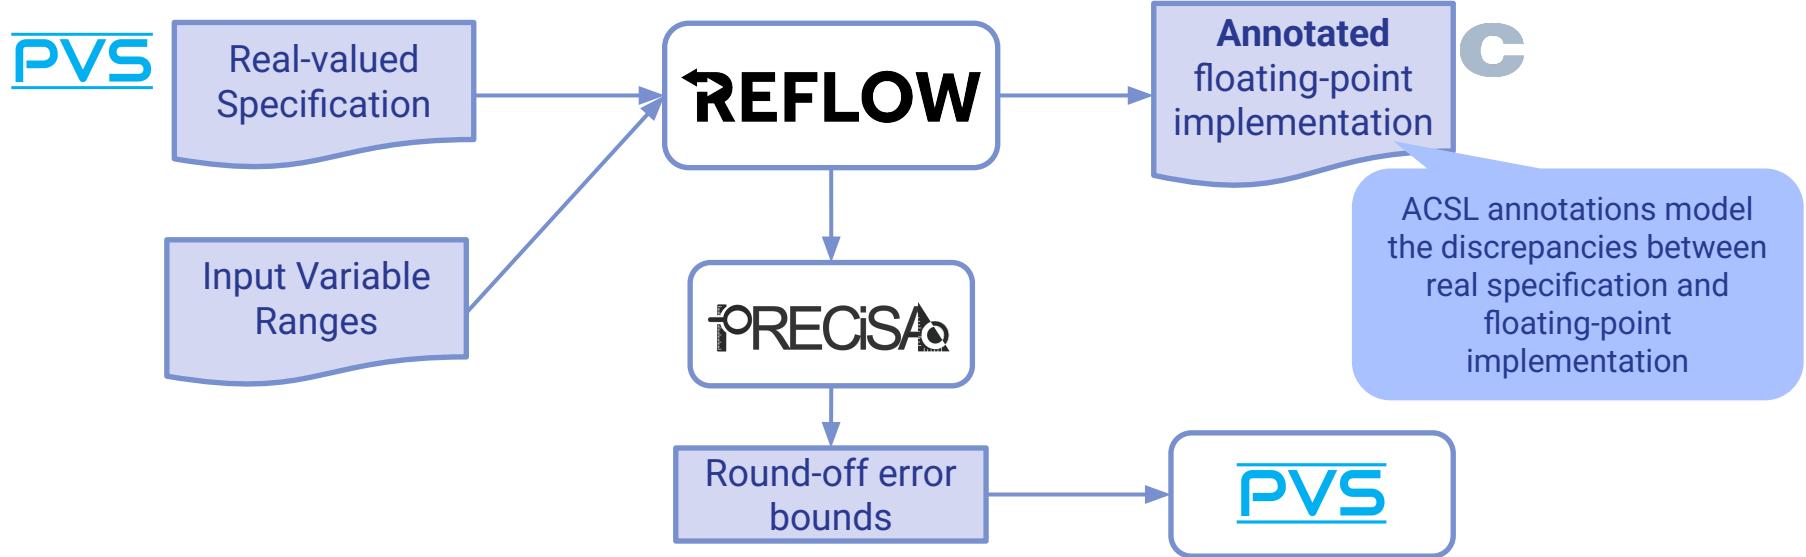
Idea: automatically extracting FP code with formal guarantees on the rounding error



ReFlow automatically generates a C floating-point implementation from a PVS real number specification



# ReFlow workflow



# Example: eps\_line code extraction

```
eps_line(sx,vx,sy,vy: real): real = ((sx*vx) + (sy*vy)) * ((sx*vx) - (sy*vy))
```



```
/*@
logic real eps_line (real sx, real vx, real sy, real vy) =
((sx * vx) + (sy * vy)) * ((sx * vx) - (sy * vy));

ensures \forall real sx, real vx, real sy, real vy;
-60000 <= sx && sx <= 60000 && -60000 <= vx && vx <= 60000 &&
-60000 <= sy && sy <= 60000 && -60000 <= vy && vy <= 60000 &&
\abs(sx_d - sx) <= ulp_dp(sx)/2 && \abs(vx_d - vx) <= ulp_dp(vx)/2 &&
\abs(sy_d - sy) <= ulp_dp(sy)/2 && \abs(vy_d - vy) <= ulp_dp(vy)/2
==> \abs(\result - eps_line(sx, vx, sy, vy)) <= 0x1.db070f4580002p14 ;
*/
double eps_line_fp (double sx_d, double vx_d, double sy_d, double vy_d) {
| return ((sx_d * vx_d) + (sy_d * vy_d)) * ((sx_d * vx_d) - (sy_d * vy_d));
}
```

RTCA/FAA Minimum  
Operational Performance  
Standards (MOPS) DO365  
for detect and avoid of  
unmanned aircraft systems  
(UAS)



# Example: eps\_line code extraction

```
eps_line(sx,vx,sy,vy: real): real = ((sx*vx) + (sy*vy)) * ((sx*vx) - (sy*vy))
```



```
/*@
logic real eps_line (real sx, real vx, real sy, real vy) =  
((sx * vx) + (sy * vy)) * ((sx * vx) - (sy * vy));
```

Axiomatic real-valued program

```
ensures \forall real sx, real vx, real sy, real vy;  
-60000 <= sx && sx <= 60000 && -60000 <= vx && vx <= 60000 &&  
-60000 <= sy && sy <= 60000 && -60000 <= vy && vy <= 60000 &&  
\abs(sx_d - sx) <= ulp_dp(sx)/2 && \abs(vx_d - vx) <= ulp_dp(vx)/2 &&  
\abs(sy_d - sy) <= ulp_dp(sy)/2 && \abs(vy_d - vy) <= ulp_dp(vy)/2  
==> \abs(\result - eps_line(sx, vx, sy, vy)) <= 0x1.db070f4580002p14 ;  
*/  
double eps_line_fp (double sx_d, double vx_d, double sy_d, double vy_d) {  
| return (sx_d * vx_d) + (sy_d * vy_d)) * ((sx_d * vx_d) - (sy_d * vy_d));  
}
```

# Example: eps\_line code extraction

```
eps_line(sx,vx,sy,vy: real): real = ((sx*vx) + (sy*vy)) * ((sx*vx) - (sy*vy))
```



```
/*@  
logic real eps_line (real sx, real vx, real sy, real vy) =  
((sx * vx) + (sy * vy)) * ((sx * vx) - (sy * vy));
```

Axiomatic real-valued program

```
ensures \forall real sx, real vx, real sy, real vy;  
-60000 <= sx && sx <= 60000 && -60000 <= vx && vx <= 60000 &&  
-60000 <= sy && sy <= 60000 && -60000 <= vy && vy <= 60000 &&  
\abs(sx_d - sx) <= ulp_dp(sx)/2 && \abs(vx_d - vx) <= ulp_dp(vx)/2 &&  
\abs(sy_d - sy) <= ulp_dp(sy)/2 && \abs(vy_d - vy) <= ulp_dp(vy)/2  
==> \abs(\result - eps_line(sx, vx, sy, vy)) <= 0x1.db070f4580002p14 ;  
*/
```

```
double eps_line_fp (double sx_d, double vx_d, double sy_d, double vy_d) {  
| return (sx_d * vx_d) + (sy_d * vy_d)) * ((sx_d * vx_d) - (sy_d * vy_d));  
}
```

Floating-point implementation

# Example: eps\_line code extraction

```
eps_line(sx,vx,sy,vy: real): real = ((sx*vx) + (sy*vy)) * ((sx*vx) - (sy*vy))
```



```
/*@  
logic real eps_line (real sx, real vx, real sy, real vy) =  
((sx * vx) + (sy * vy)) * ((sx * vx) - (sy * vy));
```

Axiomatic real-valued program

```
ensures \forall real sx, real vx, real sy, real vy;  
-60000 <= sx && sx <= 60000 && -60000 <= vx && vx <= 60000 &&  
-60000 <= sy && sy <= 60000 && -60000 <= vy && vy <= 60000 &&  
\abs(sx_d - sx) <= ulp_dp(sx)/2 && \abs(vx_d - vx) <= ulp_dp(vx)/2 &&  
\abs(sy_d - sy) <= ulp_dp(sy)/2 && \abs(vy_d - vy) <= ulp_dp(vy)/2  
==> \abs((\result - eps_line(sx, vx, sy, vy)) <= 0x1.db070f4580002p14 ;  
*/  
double eps_line_fp (double sx_d, double vx_d, double sy_d, double vy_d) {  
| return (sx_d * vx_d) + (sy_d * vy_d)) * ((sx_d * vx_d) - (sy_d * vy_d));  
}
```

Initial ranges

Floating-point implementation

# Example: eps\_line code extraction

```
eps_line(sx,vx,sy,vy: real): real = ((sx*vx) + (sy*vy)) * ((sx*vx) - (sy*vy))
```



```
/*@  
logic real eps_line (real sx, real vx, real sy, real vy) =  
((sx * vx) + (sy * vy)) * ((sx * vx) - (sy * vy));
```

Axiomatic real-valued program

```
ensures \forall real sx, real vx, real sy, real vy;  
-60000 <= sx && sx <= 60000 && -60000 <= vx && vx <= 60000 &&  
-60000 <= sy && sy <= 60000 && -60000 <= vy && vy <= 60000 &&  
\abs(sx_d - sx) <= ulp_dp(sx)/2 && \abs(vx_d - vx) <= ulp_dp(vx)/2 &&  
\abs(sy_d - sy) <= ulp_dp(sy)/2 && \abs(vy_d - vy) <= ulp_dp(vy)/2  
==> \abs((\result - eps_line(sx, vx, sy, vy)) <= 0x1.db070f4580002p14 ;  
*/  
double eps_line_fp (double sx_d, double vx_d, double sy_d, double vy_d) {  
| return (sx_d * vx_d) + (sy_d * vy_d)) * ((sx_d * vx_d) - (sy_d * vy_d));  
}
```

Initial ranges

Vars are round-to the nearest

Floating-point implementation

# Example: eps\_line code extraction

```
eps_line(sx,vx,sy,vy: real): real = ((sx*vx) + (sy*vy)) * ((sx*vx) - (sy*vy))
```



```
/*@  
logic real eps_line (real sx, real vx, real sy, real vy) =  
((sx * vx) + (sy * vy)) * ((sx * vx) - (sy * vy));
```

Axiomatic real-valued program

```
ensures \forall real sx, real vx, real sy, real vy;  
-60000 <= sx && sx <= 60000 && -60000 <= vx && vx <= 60000 &&  
-60000 <= sy && sy <= 60000 && -60000 <= vy && vy <= 60000 &&  
\abs(sx_d - sx) <= ulp_dp(sx)/2 && \abs(vx_d - vx) <= ulp_dp(vx)/2 &&  
\abs(sy_d - sy) <= ulp_dp(sy)/2 && \abs(vy_d - vy) <= ulp_dp(vy)/2  
==> \abs((\result - eps_line(sx, vx, sy, vy)) <= 0x1.db070f4580002p14  
*/  
double eps_line_fp (double sx_d, double vx_d, double sy_d, double vy_d) {  
| return (sx_d * vx_d) + (sy_d * vy_d)) * ((sx_d * vx_d) - (sy_d * vy_d));  
}
```

Initial ranges

Vars are round-to the nearest

Maximum round-off error

Floating-point implementation

# Example: Unstable conditional instrumentation

```
tcoa(sz,vz:real): real = IF (sz*vz < 0) THEN -(sz/vz) ELSE -1 ENDIF
```



```
/*@
requires (0 <= E) ;
ensures \forall real sz, real vz; (\abs(Dmul(sz_dp, vz_dp)) - (sz * vz)) <= E
==> \result.isValid ==> (sz * vz < 0 && Dmul(sz_dp, vz_dp) < 0)
|| (sz * vz >= 0 && Dmul(sz_dp, vz_dp) >= 0) ;
*/
struct maybeDouble tcoa_fp (double sz_dp, double vz_dp, double E) {
    if (sz_dp * vz_dp < - E)
        { return someDouble(-sz_dp / vz_dp); }
    } else { if (sz_dp * vz_dp >= E)
        { return someDouble(-1); }
    } else { return instability_warning(); }
}
/*@
ensures \forall real sz, real vz; (0 <= sz) && (sz <= 1000) && (400 <= vz) && (vz <= 600) &&
\abs(sz_dp - sz) <= ulp_dp(sz)/2 && \abs(vz_dp - vz) <= ulp_dp(vz)/2 &&
\result.isValid ==> \abs(\result.value - tcoa(sz, vz)) <= 0x1.c7ae147ae147dp-51;
*/
struct maybeDouble tcoa_num (double sz_dp, double vz_dp) {
    return tcoa_fp (sz_dp, vz_dp, 0x1.4800000000001p-33);
}
```



# Example: Unstable conditional instrumentation

```
tcoa(sz,vz:real): real = IF (sz*vz < 0) THEN -(sz/vz) ELSE -1 ENDIF
```



```
/*@
requires (0 <= E) ;
ensures \forall real sz, real vz; (\abs(Dmul(sz_dp, vz_dp)) - (sz * vz)) <= E
==> \result.isValid ==> (sz * vz < 0 && Dmul(sz_dp, vz_dp) < 0)
|| (sz * vz >= 0 && Dmul(sz_dp, vz_dp) >= 0) ;
*/
struct maybeDouble tcoa_fp (double sz_dp, double vz_dp, double E) {
    if (sz_dp * vz_dp < - E)
        { return someDouble(-sz_dp / vz_dp); }
    } else { if (sz_dp * vz_dp >= E)
        { return someDouble(-1); }
    } else { return instability_warning(); }
}
/*@
ensures \forall real sz, real vz; (0 <= sz) && (sz <= 1000) && (400 <= vz) && (vz <= 600) &&
\abs(sz_dp - sz) <= ulp_dp(sz)/2 && \abs(vz_dp - vz) <= ulp_dp(vz)/2 &&
\result.isValid ==> \abs(\result.value - tcoa(sz, vz)) <= 0x1.c7ae147ae147dp-51;
*/
struct maybeDouble tcoa_num (double sz_dp, double vz_dp) {
    return tcoa_fp (sz_dp, vz_dp, 0x1.4800000000001p-33);
}
```

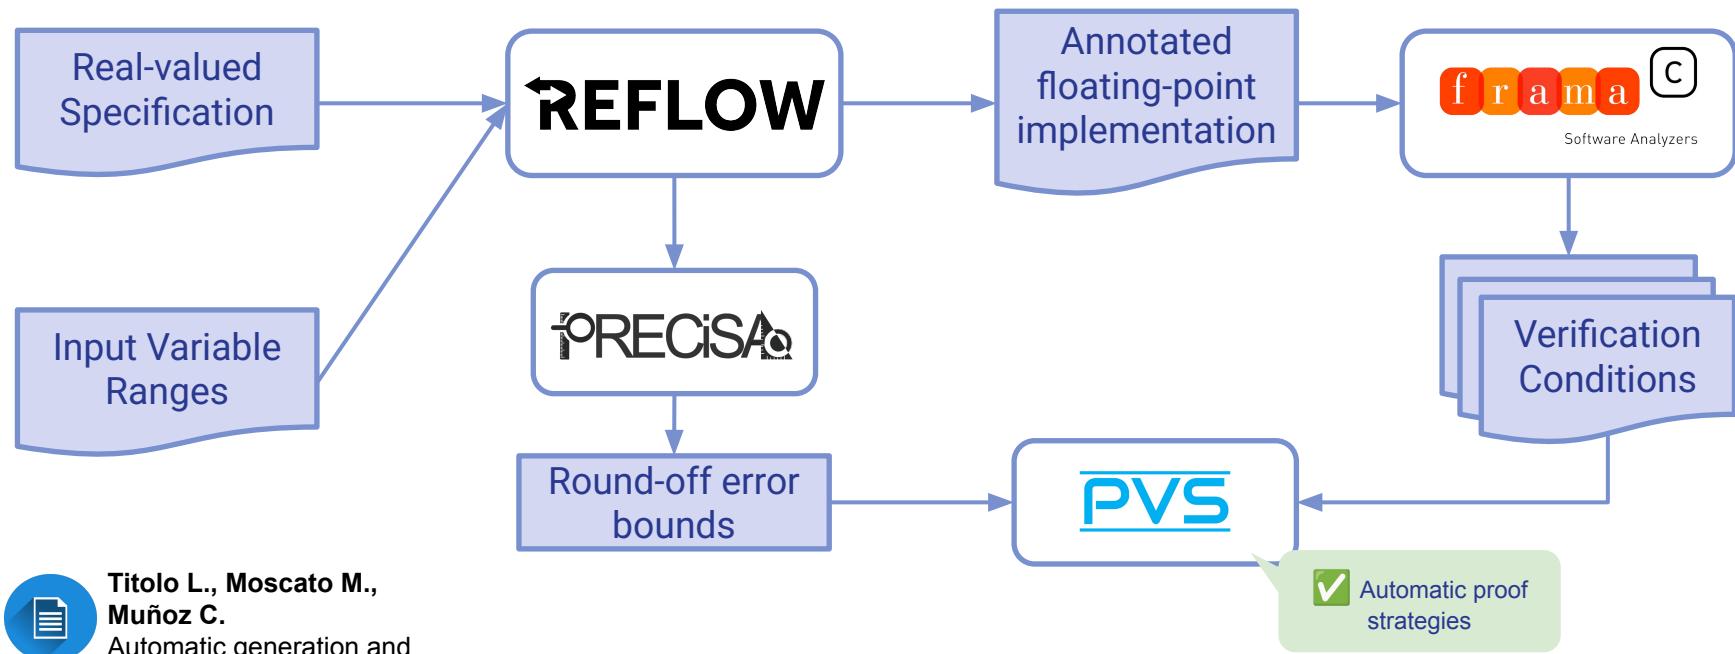
Code instrumentation to  
detect unstable  
conditionals  
(real ≠ fp control flow)

Instrumented program



**Titolo L., Muñoz C.,  
Feliú M., Moscato M.,**  
Eliminating unstable tests in  
floating-point programs.  
LOPSTR 2018.

# C code automatically verified in Frama-C and PVS

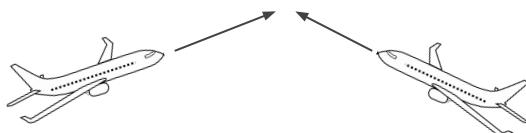


**Titolo L., Moscato M.,  
Muñoz C.**  
Automatic generation and  
verification of test-stable  
floating-point code. iFM 2020.

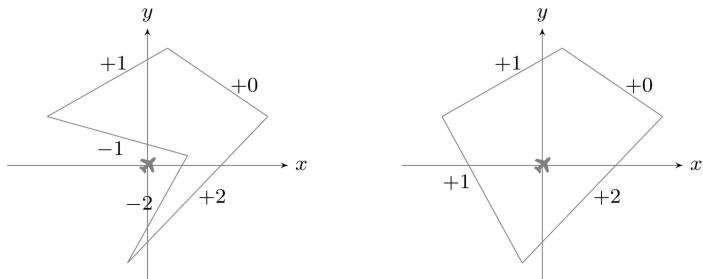
✓ Automatic proof  
strategies

# Applications of ↪REFLOW

## DAIDALUS - Detect-and-avoid

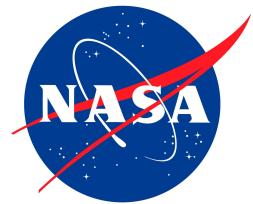


## POLYCARP - Geofencing



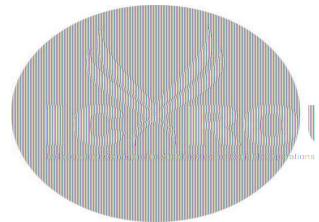
Bernardes N., Moscato M., Titolo L., Ayala M.

A provably correct floating-point implementation of Well Clear Avionics Concepts. FMCAD 2023



Moscato M., Titolo L., Feliú M., Muñoz C.

Provably Correct Floating-Point Implementation of a Point-in-Polygon Algorithm. FM 2019

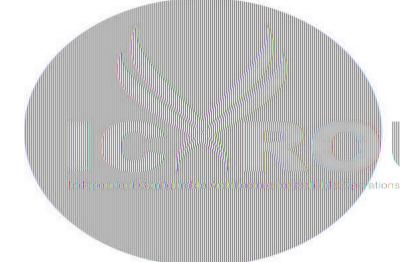


# Interactive Theorem Proving at NASA

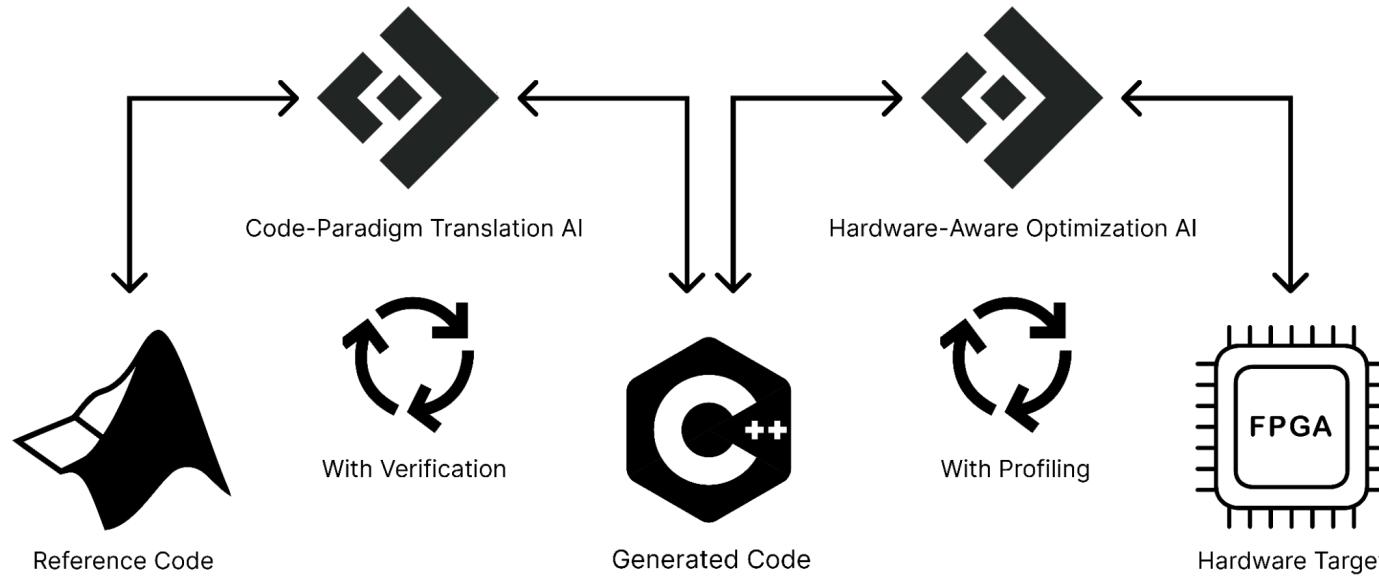
- Static Analysis
- Global optimization
- Hybrid Systems Verification (dDL)
- Requirements elicitation (FRET)
- Some applications:
  - Unmanned Aerial Vehicles
  - Detect-and-avoid
  - Mars Rovers (PLEXIL V)



REFLOW

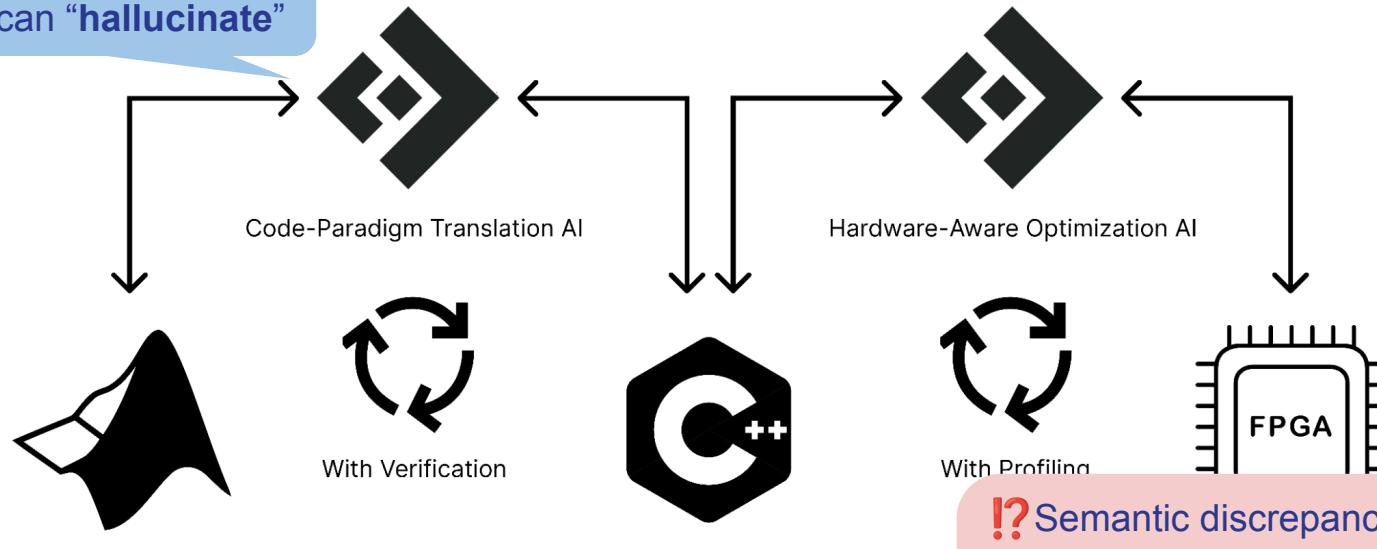


# Proofs for LLM-guided transpilation in industry



# Proofs for LLM-guided transpilation in industry

- ! LLMs are Black Boxes!
- 🌀 LLMs can “hallucinate”

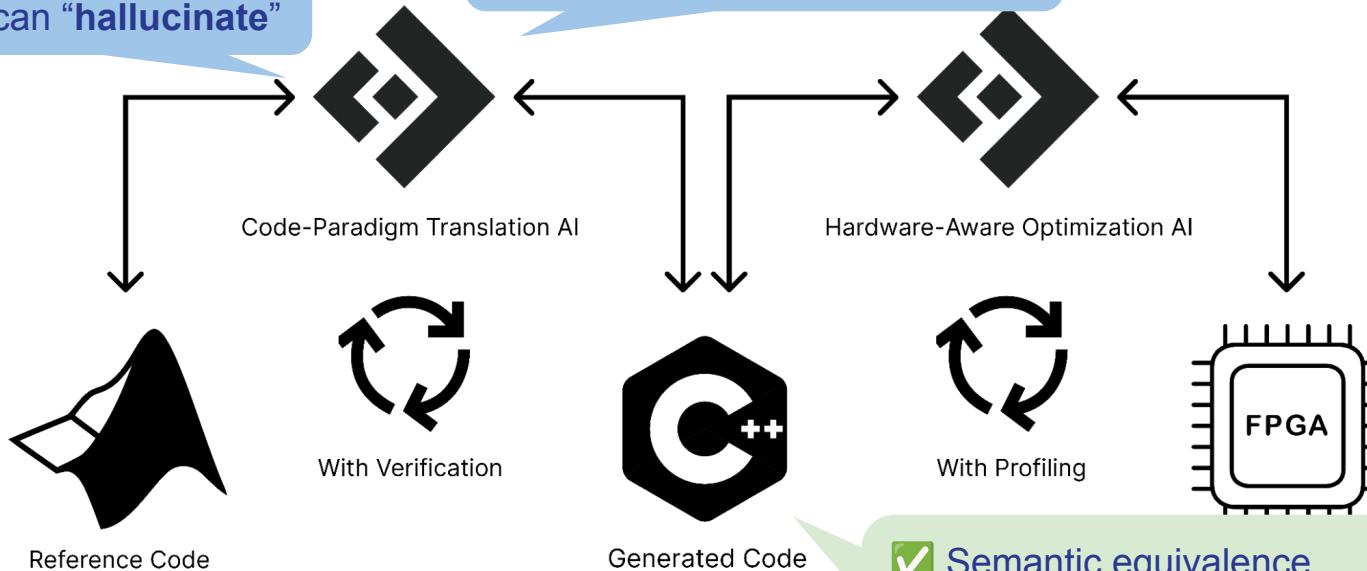


- ! Semantic discrepancies
- ! Code vulnerabilities
- ! Specification violation

# Proofs for LLM-guided transpilation in industry

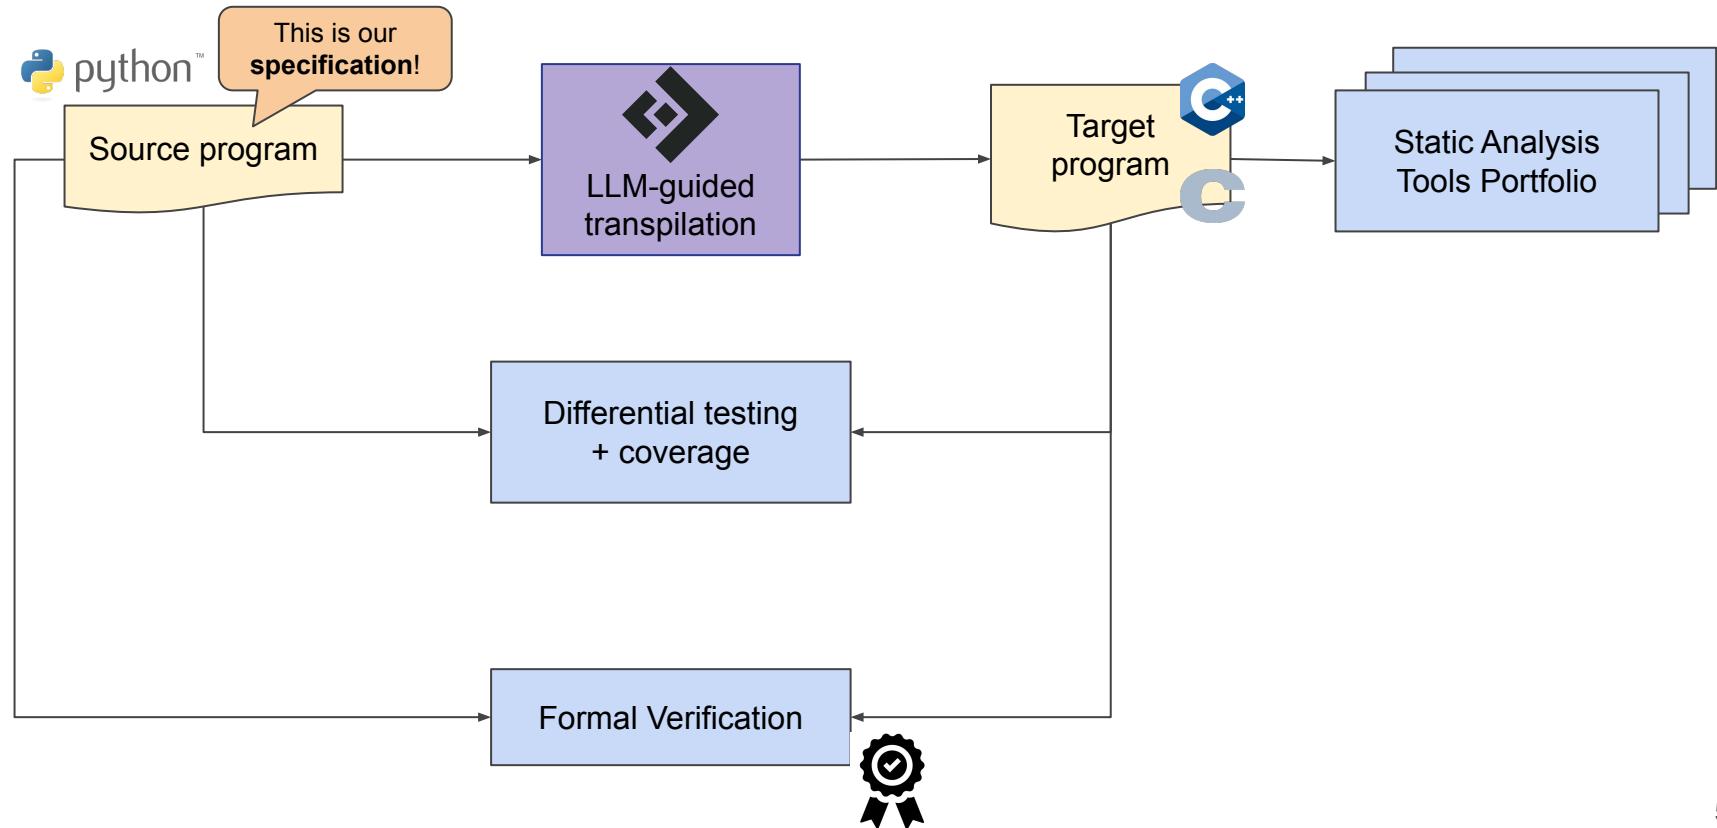
! LLMs are **Black Boxes!**  
🌀 LLMs can “hallucinate”

💻 ✅ LLM generated code must be validated and verified

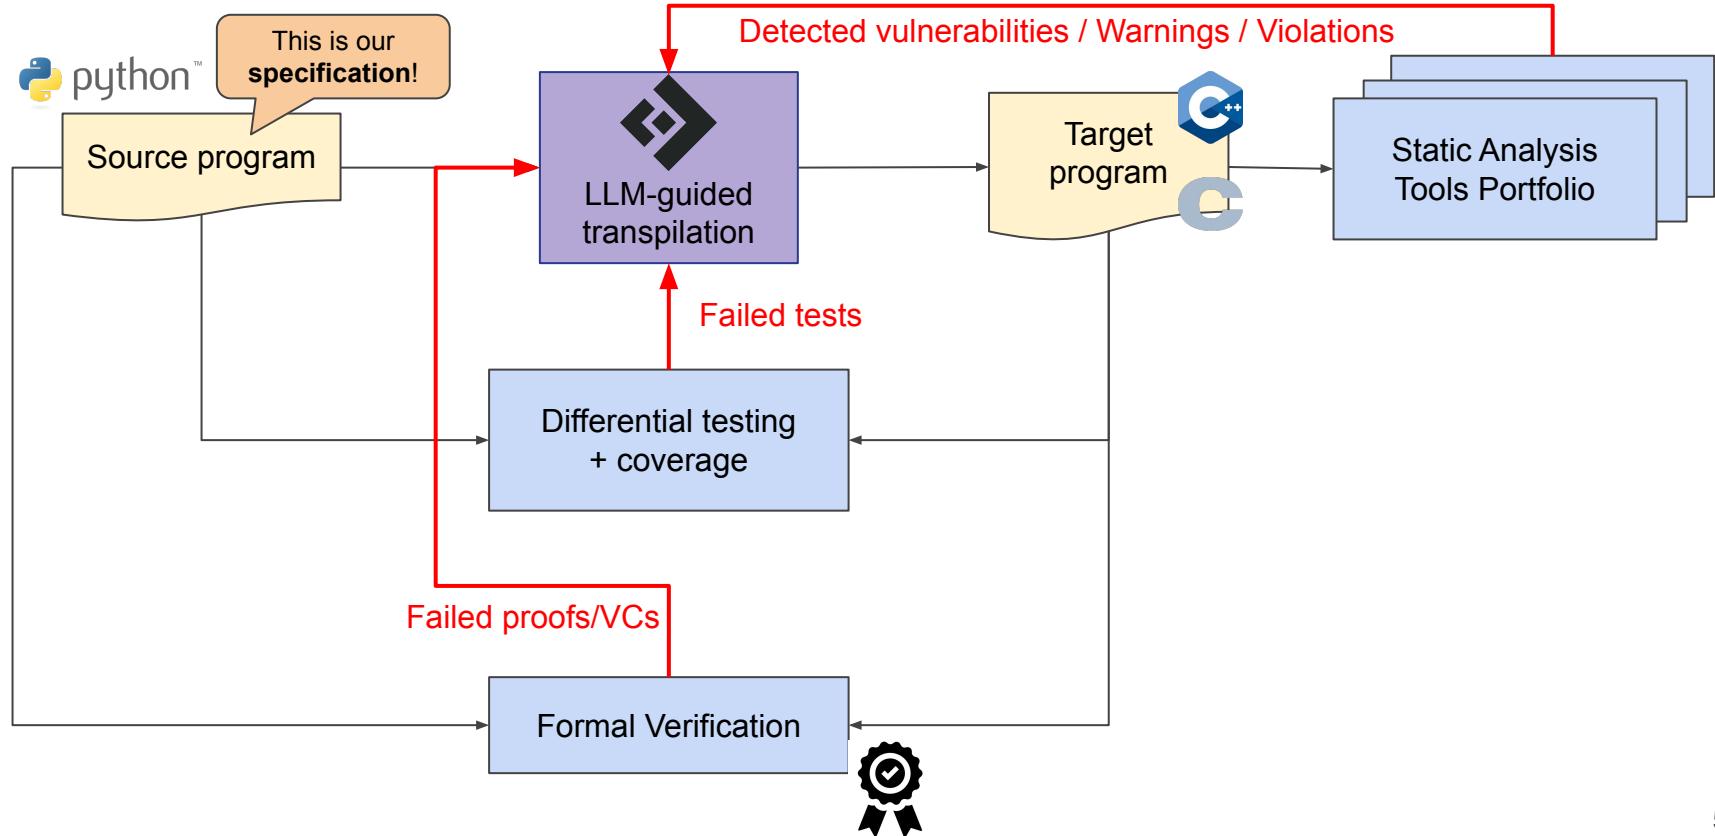


- ✓ Semantic equivalence  
🛡️ No code vulnerabilities  
📘 Standard compliance (ex MISRA)

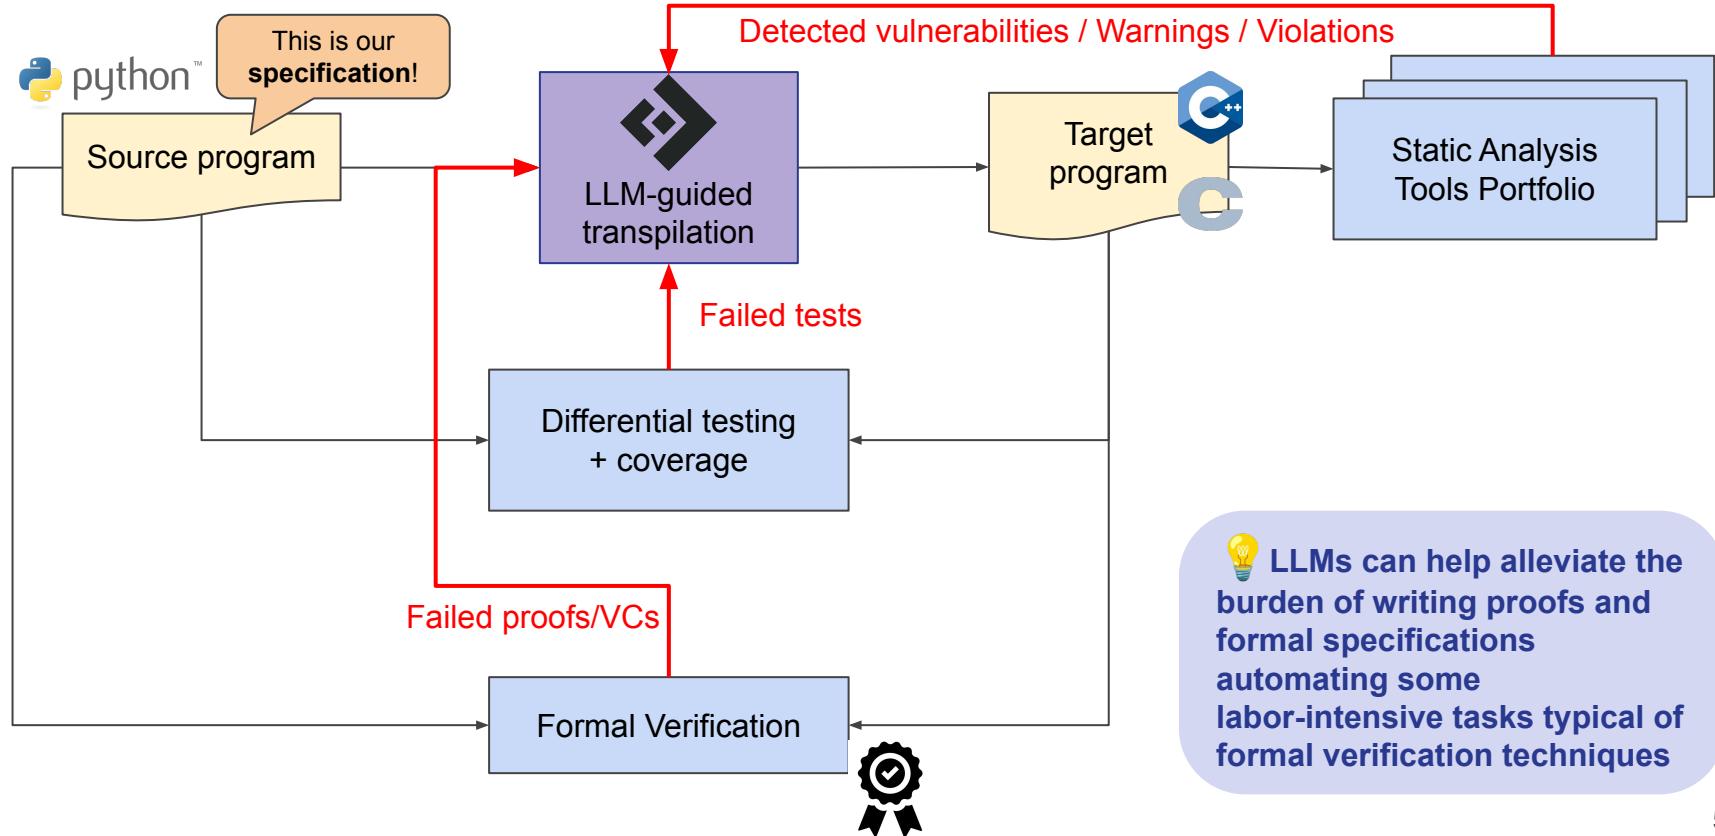
# LLM-generated code must be validated and verified!



# LLM-generated code must be validated and verified!



# LLM-generated code must be validated and verified!



# Conclusion

- 🚀 Interactive theorem provers are heavily used at NASA for safety critical applications
- 💻 ITP for AI - LLM-guided code generation greatly benefits from formal proofs
- 🤖 AI for ITP
  - 🔄 Proof translation for importing/exporting libraries across different ITP platforms
  - 🔧 Proof repair and simplification
  - 🛠️ Proof generation/automation
- 📝🌟 Exciting research and topics to explore at the intersection of ITP and AI

# Thanks for your attention!



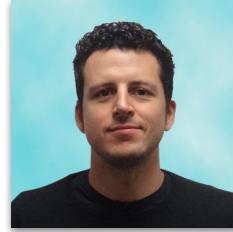
Aaron Dutle  
NASA



Cesar Muñoz  
NASA



Mariano Moscato  
NASA/AMA



Marco A. Feliu  
NASA/AMA



Paolo Masci  
x-NASA



[laura@codemetal.ai](mailto:laura@codemetal.ai)  
<https://lauratitolo.github.io/>

# Backup slides

# Example: Unstable conditional instrumentation

```
tcoa(sz,vz:real): real = IF (sz*vz < 0) THEN -(sz/vz) ELSE -1 ENDIF
```



```
/*@
requires (0 <= E) ;
ensures \forall real sz, real vz; ((\abs(Dmul(sz_dp, vz_dp)) - (sz * vz)) <= E
==> \result.isValid ==> (sz * vz < 0 && Dmul(sz_dp, vz_dp) < 0)
|| (sz * vz >= 0 && Dmul(sz_dp, vz_dp) >= 0) ;
*/
struct maybeDouble tcoa_fp (double sz_dp, double vz_dp, double E) {
    if (sz_dp * vz_dp < - E)
        { return someDouble(-sz_dp / vz_dp); }
    else { if (sz_dp * vz_dp >= E)
        { return someDouble(-1); }
        else { return instability_warning(); } }
}
/*@
ensures \forall real sz, real vz; (0 <= sz) && (sz <= 1000) && (400 <= vz) && (vz <= 600) &&
\abs(sz_dp - sz) <= ulp_dp(sz)/2 && \abs(vz_dp - vz) <= ulp_dp(vz)/2 &&
\result.isValid ==> \abs(\result.value - tcoa(sz, vz)) <= 0x1.c7ae147ae147dp-51;
*/
struct maybeDouble tcoa_num (double sz_dp, double vz_dp) {
    return tcoa_fp (sz_dp, vz_dp, 0x1.4800000000001p-33);
}
```

ACSL contract

Instrumented program

symbolic

# Example: Unstable conditional instrumentation

```
tcoa(sz,vz:real): real = IF (sz*vz < 0) THEN -(sz/vz) ELSE -1 ENDIF
```



```
/*@
requires (0 <= E) ;
ensures \forall real sz, real vz; (\abs(Dmul(sz_dp, vz_dp)) - (sz * vz)) <= E
    ==> \result.isValid ==> (sz * vz < 0 && Dmul(sz_dp, vz_dp) < 0)
        || (sz * vz >= 0 && Dmul(sz_dp, vz_dp) >= 0) ;
*/
struct maybeDouble tcoa_fp (double sz_dp, double vz_dp, double E) {
    if (sz_dp * vz_dp < - E)
        { return someDouble(-sz_dp / vz_dp); }
    } else { if (sz_dp * vz_dp >= E)
        { return someDouble(-1); }
        } else { return instability_warning(); }
}
/*@
ensures \forall real sz, real vz; (0 <= sz) && (sz <= 1000) && (400 <= vz) && (vz <= 600) &&
\abs(sz_dp - sz) <= ulp_dp(sz)/2 && \abs(vz_dp - vz) <= ulp_dp(vz)/2 &&
\result.isValid ==> \abs(\result.value - tcoa(sz, vz)) <= 0x1.c7ae147ae147dp-51;
*/
struct maybeDouble tcoa_num (double sz_dp, double vz_dp) {
    return tcoa_fp (sz_dp, vz_dp, 0x1.4800000000001p-33);
}
```

ACSL contract

instrumented program

function call to  
instrumented tcoa\_fp

symbolic

# Example: Unstable conditional instrumentation

```
tcoa(sz,vz:real): real = IF (sz*vz < 0) THEN -(sz/vz) ELSE -1 ENDIF
```



```
/*@
requires (0 <= E) ;
ensures \forall real sz, real vz; (\abs(Dmul(sz_dp, vz_dp)) - (sz * vz)) <= E
    ==> \result.isValid ==> (sz * vz < 0 && Dmul(sz_dp, vz_dp) < 0)
        || (sz * vz >= 0 && Dmul(sz_dp, vz_dp) >= 0) ;
*/
struct maybeDouble tcoa_fp (double sz_dp, double vz_dp, double E) {
```

```
    if (sz_dp * vz_dp < - E)
    { return someDouble(-sz_dp / vz_dp);
    } else { if (sz_dp * vz_dp >= E)
        { return someDouble(-1);
        } else { return instability_warning();}}
    }
```

```
/*@
ensures \forall real sz, real vz; (0 <= sz) && (sz <= 1000) && (400 <= vz) && (vz <= 600) &&
\abs(sz_dp - sz) <= ulp_dp(sz)/2 && \abs(vz_dp - vz) <= ulp_dp(vz)/2 &&
\result.isValid ==> \abs(\result.value - tcoa(sz, vz)) <= 0x1.c7ae147ae147dp-51;
*/
struct maybeDouble tcoa_num (double sz_dp, double vz_dp) {
```

```
    return tcoa_fp (sz_dp, vz_dp, 0x1.4800000000001p-33);
}
```

ACSL contract

Instrumented program

ACSL contract

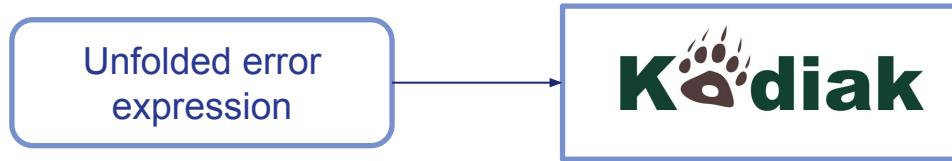
function call to  
instrumented tcoa\_fp

symbolic

numeric

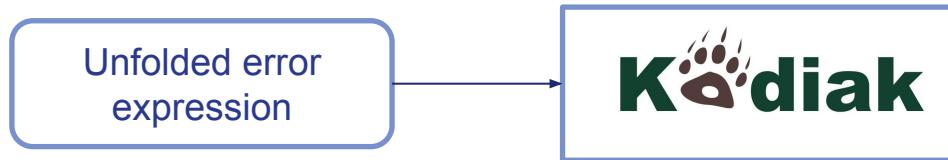
# Function calls abstract analysis

Before:

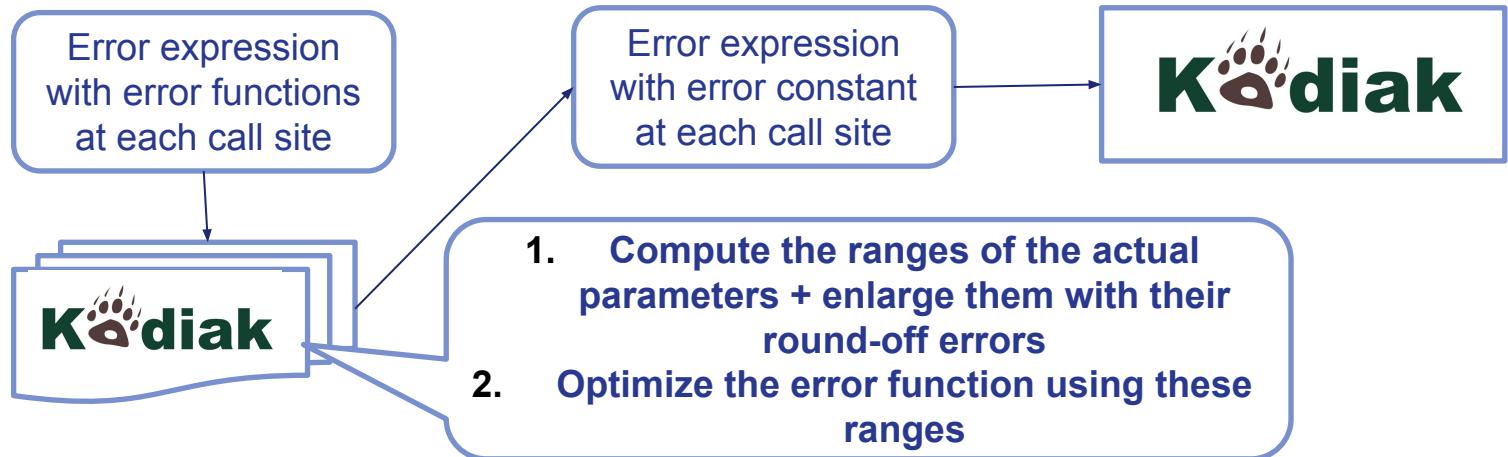


# Function calls abstract analysis

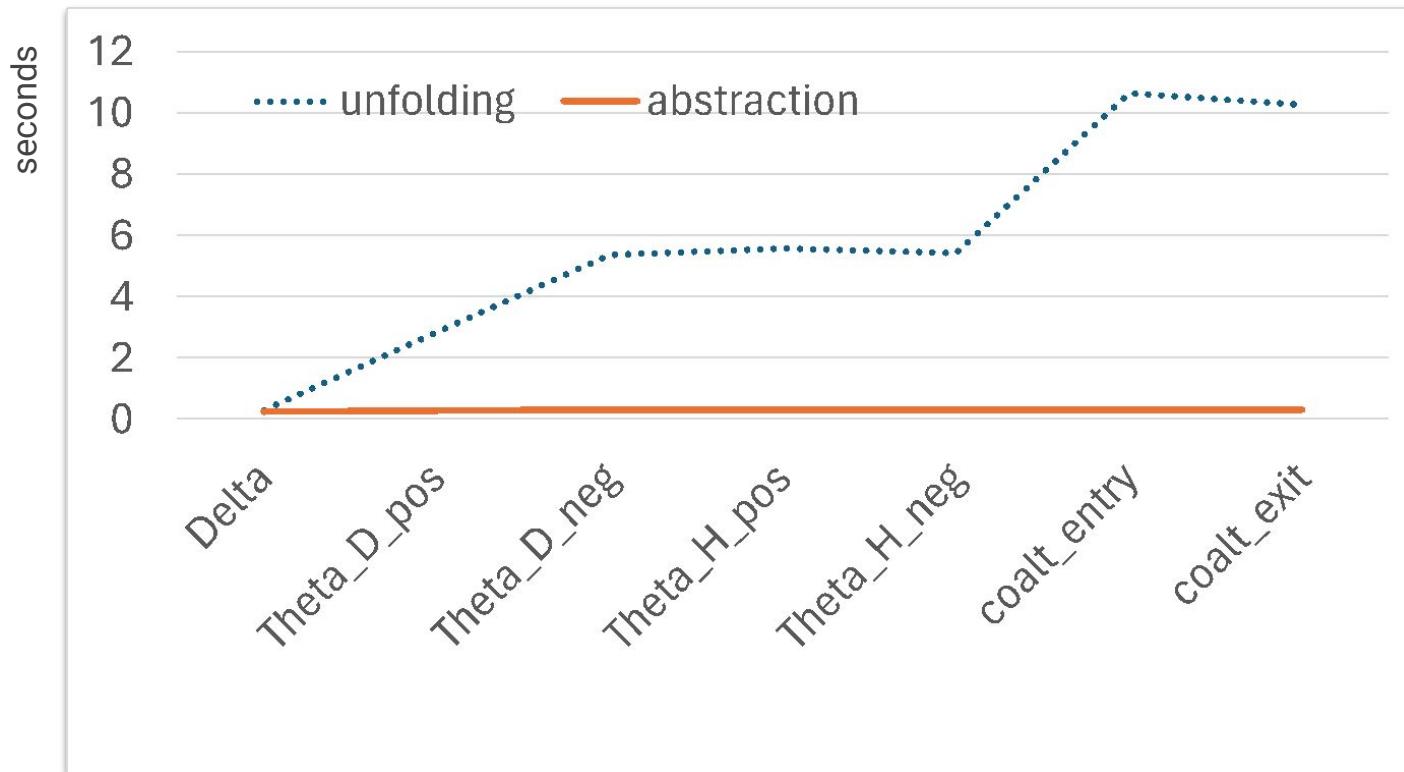
Before:



Now:



# Function calls abstract analysis experiments



# What's next for PRECiSA and Reflow?

- ReFlow is currently under revision for **NASA open-source** release
- Currently limited support for a class of for **loops**  
⇒ Add support for more complex loops
- Improve the precision in the **instability** analysis for conditional and loops
- Reduce the **complexity** of the ACSL annotation
- **Integration** with precision **optimization** tools (Herbie)

# Example: detect-and-avoid coordination

eps\_line(vx, vy, sx, sy) =

if  $(sx \cdot vx + sy \cdot vy) * (sx \cdot vx - sy \cdot vy) > 0$  then

1 // right turn

else

-1 // left turn



# Instrumentation to detect instability

```
eps_line'(vx, vy, sx, sy) =
```

```
if (sx*vx + sy*vy) * (sx*vx - sy*vy) > ε then
```

1 // right turn

```
elsif (sx*vx + sy*vy) * (sx*vx - sy*vy) ≤ -ε then
```

-1 // left turn

```
else ω // warning!
```

ε is a sound overestimation  
of the error of the expression



Strengthen the guards

Cases in which the rounding error may  
affect the evaluation of the guard

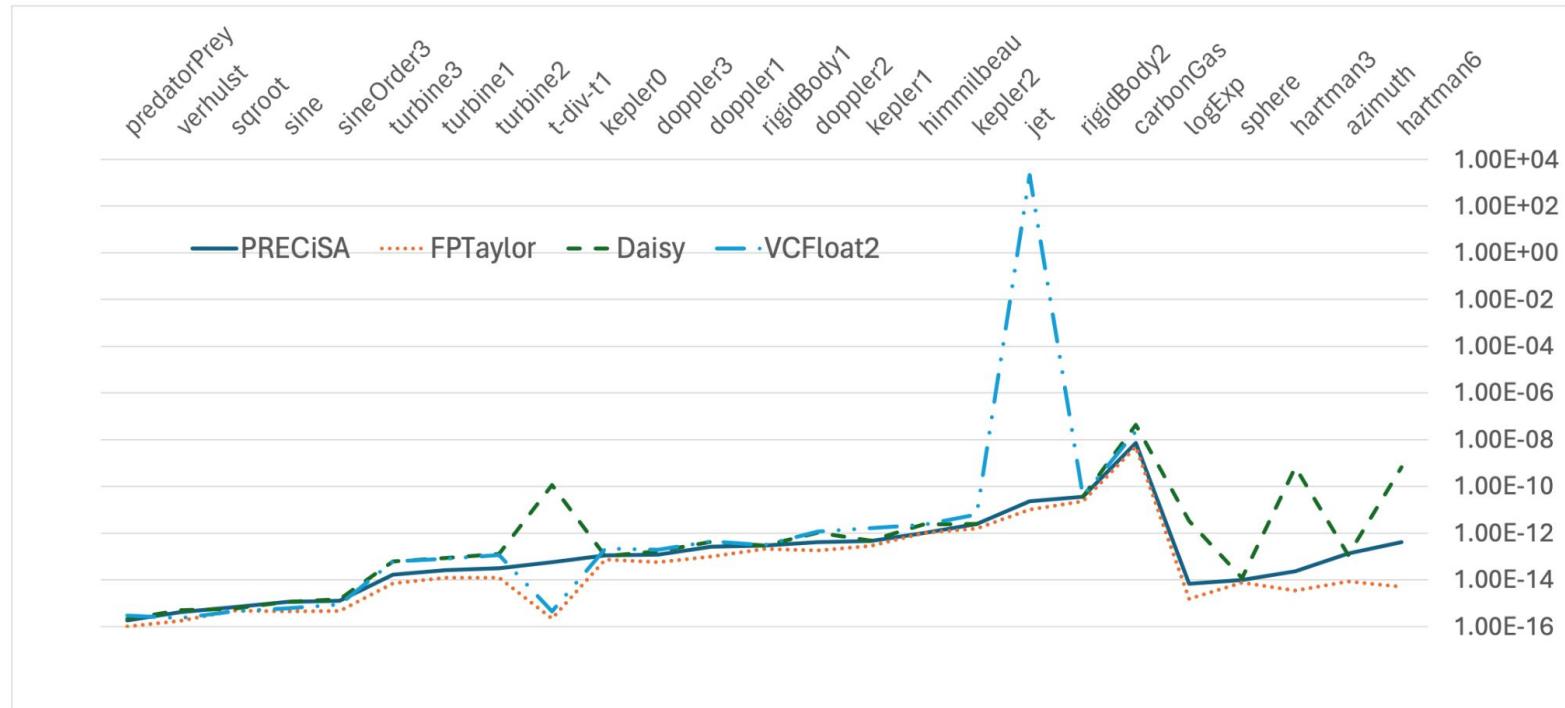


**Titolo L., Muñoz C.,  
Feliú M., Moscato M.,**  
Eliminating unstable tests in  
floating-point programs.  
LOPSTR 2018.

# Tool comparison

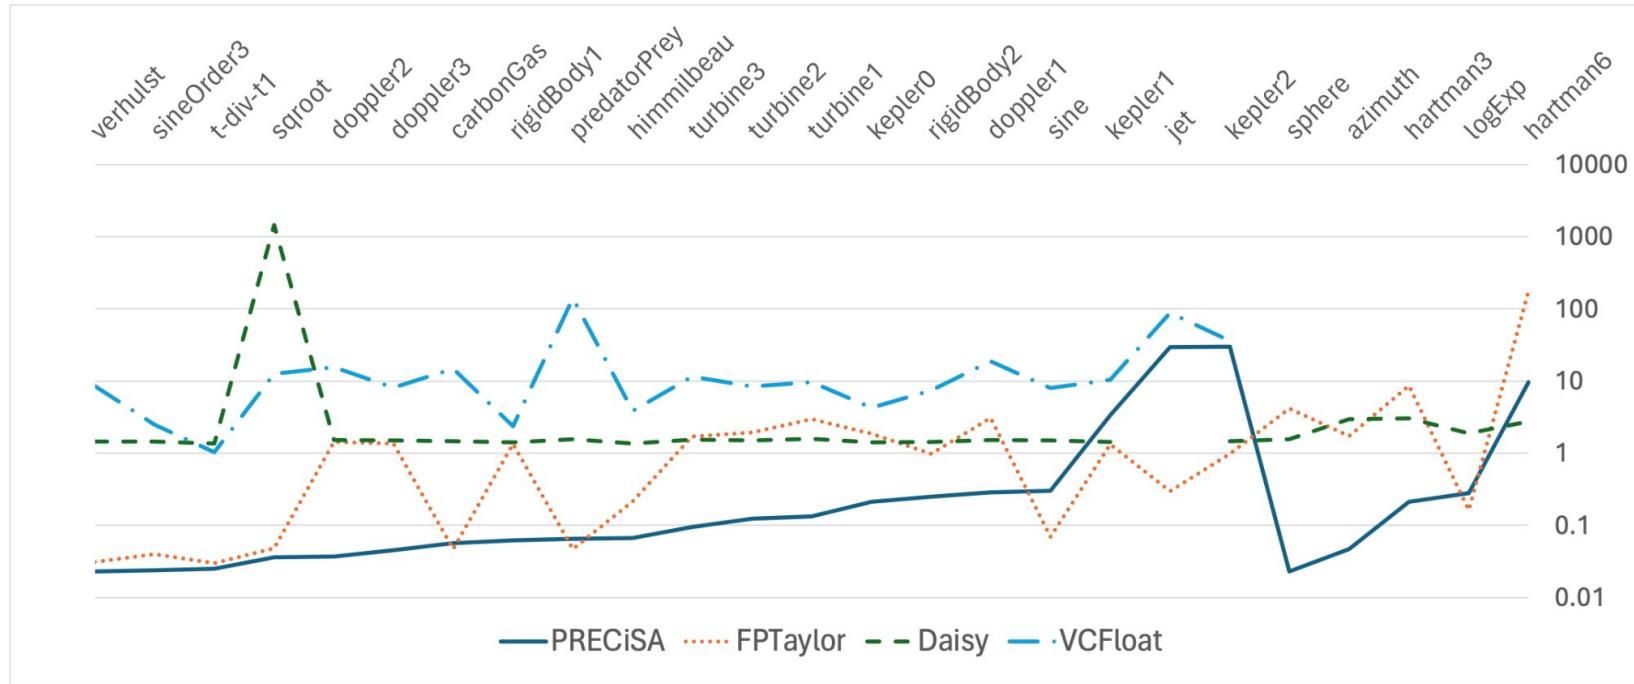
	PRECiSA	FPTaylor	Daisy	VCFloat	Fluctuat	Gappa
proof certificates	✓	✓	✗	✓	✗	✓
conditionals	✓	✗	✗	✗	✓	✗
instability detection	✓	✗	✗	✗	✓	✗
instability analysis	✓	✗	✗	✗	✗	✗
function calls	✓	✗	✗	✗	✓	✗
bounded loops	✓	✗	✗	✗	✓	✗
widening	✓	✗	✗	✗	✓	✗
data collections	✓	✗	✓	✗	✓	✗
rounding modes	✗	✓	✗	✗	✗	✗
fixed-point arith.	✗	✗	✓	✗	✗	✓

# Tool comparison



Experimental results for absolute round-off error bounds.

# Tool comparison



Times in seconds for the generation of round-off error bounds.