



Verifying an Efficient Algorithm to Compute Bernoulli Numbers

Manuel Eberl

Peter Lammich

Faulhaber's formula

$$\sum_{i=0}^n i^0 = n+1$$

Faulhaber's formula

$$\sum_{i=0}^n i^0 = n + 1$$

$$\sum_{i=0}^n i^1 = \frac{1}{2}(n^2 + n)$$

Faulhaber's formula

$$\sum_{i=0}^n i^0 = n + 1$$

$$\sum_{i=0}^n i^1 = \frac{1}{2}(n^2 + n)$$

$$\sum_{i=0}^n i^2 = \frac{1}{6}(2n^3 + 3n^2 + n)$$

Faulhaber's formula

$$\sum_{i=0}^n i^0 = n+1$$

$$\sum_{i=0}^n i^1 = \frac{1}{2}(n^2 + n)$$

$$\sum_{i=0}^n i^2 = \frac{1}{6}(2n^3 + 3n^2 + n)$$

$$\sum_{i=0}^n i^k = \frac{1}{m+1}(P_{m+1}(n+1) - P_{m+1}(0)) \quad \text{where } P_m(x) = \sum_{0 \leq k \leq m} \binom{m}{k} B_k x^{m-k}$$

Bernoulli Numbers

Definition

Definition (Bernoulli numbers)

The Bernoulli numbers are the sequence of rationals $(B_k)_{k \geq 0}$ with

$$\sum_{k \geq 0} B_k \frac{z^k}{k!} = \frac{z}{\exp(z) - 1}$$

We write $B_k = N_k / D_k$.

Definition

Definition (Bernoulli numbers)

The Bernoulli numbers are the sequence of rationals $(B_k)_{k \geq 0}$ with

$$\sum_{k \geq 0} B_k \frac{z^k}{k!} = \frac{z}{\exp(z) - 1}$$

We write $B_k = N_k / D_k$.

k	0	1	2	3	4	5	6	7	8	9	10
B_k	1	$-\frac{1}{2}$	$\frac{1}{6}$	0	$-\frac{1}{30}$	0	$\frac{1}{42}$	0	$-\frac{1}{30}$	0	$\frac{5}{66}$

Definition

Definition (Bernoulli numbers)

The Bernoulli numbers are the sequence of rationals $(B_k)_{k \geq 0}$ with

$$\sum_{k \geq 0} B_k \frac{z^k}{k!} = \frac{z}{\exp(z) - 1}$$

We write $B_k = N_k / D_k$.

k	0	1	2	3	4	5	6	7	8	9	10
B_k	1	$-\frac{1}{2}$	$\frac{1}{6}$	0	$-\frac{1}{30}$	0	$\frac{1}{42}$	0	$-\frac{1}{30}$	0	$\frac{5}{66}$

$$B_{60} \approx -1.2 \cdot 10^{42} / 56786730$$

$$B_{110} \approx 7.6 \cdot 10^{93} / 1518$$

Some more properties

- $B_k = 0$ for odd $k > 1$

Some more properties

- $B_k = 0$ for odd $k > 1$
- B_2, B_4, B_6, \dots positive; B_4, B_8, B_{12}, \dots negative

Some more properties

- $B_k = 0$ for odd $k > 1$
- B_2, B_4, B_6, \dots positive; B_4, B_8, B_{12}, \dots negative
- Through connection with ζ : $\log|B_{2k}|$ and $\log|N_{2k}|$ are of order $\Theta(k \log k)$

Some more properties

- $B_k = 0$ for odd $k > 1$
- B_2, B_4, B_6, \dots positive; B_4, B_8, B_{12}, \dots negative
- Through connection with ζ : $\log|B_{2k}|$ and $\log|N_{2k}|$ are of order $\Theta(k \log k)$

Theorem (Voronoi's congruence)

$$(a^k - 1)B_k \equiv ka^{k-1} \sum_{1 \leq m < p} m^{k-1} \left\lfloor \frac{ma}{p} \right\rfloor \pmod{p}$$

Allows us to compute $B_k \bmod p$ as a sum.

Some more properties

- $B_k = 0$ for odd $k > 1$
- B_2, B_4, B_6, \dots positive; B_4, B_8, B_{12}, \dots negative
- Through connection with ζ : $\log|B_{2k}|$ and $\log|N_{2k}|$ are of order $\Theta(k \log k)$

Theorem (Voronoi's congruence)

$$(a^k - 1)B_k \equiv ka^{k-1} \sum_{1 \leq m < p} m^{k-1} \left\lfloor \frac{ma}{p} \right\rfloor \pmod{p}$$

Allows us to compute $B_k \bmod p$ as a sum.

Theorem (Kummer's congruence)

$$B_k/k \equiv B_{k'}/k \pmod{p} \quad \text{where } k' = k \bmod (p-1)$$

Thus we can w.l.o.g. assume that $k = 2, 4, \dots, p-3$.

Computation

Algorithms

Slow algorithms:

- Using the recurrence that drops out of the definition: very inefficient.

Algorithms

Slow algorithms:

- Using the recurrence that drops out of the definition: very inefficient.
- Various cubic methods (some only use integer arithmetic)

Algorithms

Slow algorithms:

- Using the recurrence that drops out of the definition: very inefficient.
- Various cubic methods (some only use integer arithmetic)

Modern quadratic algorithms, roughly $O(k^2 \log^{1+o(1)} k)$:

Algorithms

Slow algorithms:

- Using the recurrence that drops out of the definition: very inefficient.
- Various cubic methods (some only use integer arithmetic)

Modern quadratic algorithms, roughly $O(k^2 \log^{1+o(1)} k)$:

- Compute $\sin(z)/\cos(z)$ as formal power series

Algorithms

Slow algorithms:

- Using the recurrence that drops out of the definition: very inefficient.
- Various cubic methods (some only use integer arithmetic)

Modern quadratic algorithms, roughly $O(k^2 \log^{1+o(1)} k)$:

- Compute $\sin(z)/\cos(z)$ as formal power series
- Compute B_{2k} by approximating $\zeta(2k)$

Algorithms

Slow algorithms:

- Using the recurrence that drops out of the definition: very inefficient.
- Various cubic methods (some only use integer arithmetic)

Modern quadratic algorithms, roughly $O(k^2 \log^{1+o(1)} k)$:

- Compute $\sin(z)/\cos(z)$ as formal power series
- Compute B_{2k} by approximating $\zeta(2k)$

Allows computing B_{2k} without B_2, \dots, B_{2k-2} .

Algorithms

Slow algorithms:

- Using the recurrence that drops out of the definition: very inefficient.
- Various cubic methods (some only use integer arithmetic)

Modern quadratic algorithms, roughly $O(k^2 \log^{1+o(1)} k)$:

- Compute $\sin(z)/\cos(z)$ as formal power series
- Compute B_{2k} by approximating $\zeta(2k)$
Allows computing B_{2k} without B_2, \dots, B_{2k-2} .
- [Harvey \(2010\)](#): Multimodular algorithm.

Algorithms

Slow algorithms:

- Using the recurrence that drops out of the definition: very inefficient.
- Various cubic methods (some only use integer arithmetic)

Modern quadratic algorithms, roughly $O(k^2 \log^{1+o(1)} k)$:

- Compute $\sin(z)/\cos(z)$ as formal power series
- Compute B_{2k} by approximating $\zeta(2k)$
Allows computing B_{2k} without B_2, \dots, B_{2k-2} .
- [Harvey \(2010\)](#): Multimodular algorithm.
Compute $B_k \bmod p$ for many primes p independently.

Algorithms

Slow algorithms:

- Using the recurrence that drops out of the definition: very inefficient.
- Various cubic methods (some only use integer arithmetic)

Modern quadratic algorithms, roughly $O(k^2 \log^{1+o(1)} k)$:

- Compute $\sin(z)/\cos(z)$ as formal power series
- Compute B_{2k} by approximating $\zeta(2k)$
Allows computing B_{2k} without B_2, \dots, B_{2k-2} .
- [Harvey \(2010\)](#): Multimodular algorithm.
Compute $B_k \bmod p$ for many primes p independently.
Combine results via *Chinese Remainder Theorem*.

Harvey: High-Level Algorithm

Given input k . compute $B_k = N_k / D_k$ as follows:

- Compute “enough” primes via prime sieve.

Harvey: High-Level Algorithm

Given input k . compute $B_k = N_k / D_k$ as follows:

- Compute “enough” primes via prime sieve.
- Compute denominator via *von Staudt–Clausen*: $D_k = \prod_{(p-1)|k} p$

Harvey: High-Level Algorithm

Given input k . compute $B_k = N_k / D_k$ as follows:

- Compute “enough” primes via prime sieve.
- Compute denominator via *von Staudt–Clausen*: $D_k = \prod_{(p-1)|k} p$
- Compute precise approximation of $\log_2 N_k$

Harvey: High-Level Algorithm

Given input k . compute $B_k = N_k / D_k$ as follows:

- Compute “enough” primes via prime sieve.
- Compute denominator via *von Staudt–Clausen*: $D_k = \prod_{(p-1)|k} p$
- Compute precise approximation of $\log_2 N_k$
- Compute set of primes P such that $\prod P > N_k$

Harvey: High-Level Algorithm

Given input k . compute $B_k = N_k / D_k$ as follows:

- Compute “enough” primes via prime sieve.
- Compute denominator via *von Staudt–Clausen*: $D_k = \prod_{(p-1)|k} p$
- Compute precise approximation of $\log_2 N_k$
- Compute set of primes P such that $\prod P > N_k$
- Compute $N_k \bmod p$ for each $p \in P$

Harvey: High-Level Algorithm

Given input k . compute $B_k = N_k / D_k$ as follows:

- Compute “enough” primes via prime sieve.
- Compute denominator via *von Staudt–Clausen*: $D_k = \prod_{(p-1)|k} p$
- Compute precise approximation of $\log_2 N_k$
- Compute set of primes P such that $\prod P > N_k$
- Compute $N_k \bmod p$ for each $p \in P$
- Use CRT to compute $N_k \bmod \prod P$

Harvey: High-Level Algorithm

Given input k . compute $B_k = N_k / D_k$ as follows:

- Compute “enough” primes via prime sieve.
- Compute denominator via *von Staudt–Clausen*: $D_k = \prod_{(p-1)|k} p$
- Compute precise approximation of $\log_2 N_k$
- Compute set of primes P such that $\prod P > N_k$
- Compute $N_k \bmod p$ for each $p \in P$
- Use CRT to compute $N_k \bmod \prod P$
- Read off N_k

Harvey: Computation modulo p

What dictates the performance: Computing $B_k \bmod p$ for “nice” primes p (where $2^k \not\equiv_p 1$).

Harvey: Computation modulo p

What dictates the performance: Computing $B_k \bmod p$ for “nice” primes p (where $2^k \not\equiv_p 1$).

Each such computation boils down to:

- An outer for-loop with few iterations

Harvey: Computation modulo p

What dictates the performance: Computing $B_k \bmod p$ for “nice” primes p (where $2^k \not\equiv_p 1$).

Each such computation boils down to:

- An outer for-loop with few iterations
- An inner for-loop with many iterations

Harvey: Computation modulo p

What dictates the performance: Computing $B_k \bmod p$ for “nice” primes p (where $2^k \not\equiv_p 1$).

Each such computation boils down to:

- An outer for-loop with few iterations
- An inner for-loop with many iterations
- In the loop body:

Harvey: Computation modulo p

What dictates the performance: Computing $B_k \bmod p$ for “nice” primes p (where $2^k \not\equiv_p 1$).

Each such computation boils down to:

- An outer for-loop with few iterations
- An inner for-loop with many iterations
- In the loop body:
 - Compute next 64 bits of binary expansion of $1/p$

Harvey: Computation modulo p

What dictates the performance: Computing $B_k \bmod p$ for “nice” primes p (where $2^k \not\equiv_p 1$).

Each such computation boils down to:

- An outer for-loop with few iterations
- An inner for-loop with many iterations
- In the loop body:
 - Compute next 64 bits of binary expansion of $1/p$
 - For each 8-bit block b_0, \dots, b_7 , add some value to a table entry $t[i, b_i]$

Harvey: Computation modulo p

What dictates the performance: Computing $B_k \bmod p$ for “nice” primes p (where $2^k \not\equiv_p 1$).

Each such computation boils down to:

- An outer for-loop with few iterations
- An inner for-loop with many iterations
- In the loop body:
 - Compute next 64 bits of binary expansion of $1/p$
 - For each 8-bit block b_0, \dots, b_7 , add some value to a table entry $t[i, b_i]$
- Afterwards: determine final result as a weighted sum of the table entries

Harvey: Computation modulo p

What dictates the performance: Computing $B_k \bmod p$ for “nice” primes p (where $2^k \not\equiv_p 1$).

Each such computation boils down to:

- An outer for-loop with few iterations
- An inner for-loop with many iterations
- In the loop body:
 - Compute next 64 bits of binary expansion of $1/p$
 - For each 8-bit block b_0, \dots, b_7 , add some value to a table entry $t[i, b_i]$
- Afterwards: determine final result as a weighted sum of the table entries

Important: Inner loop body as cheap as possible.

Harvey: Computation modulo p

What dictates the performance: Computing $B_k \bmod p$ for “nice” primes p (where $2^k \not\equiv_p 1$).

Each such computation boils down to:

- An outer for-loop with few iterations
- An inner for-loop with many iterations
- In the loop body:
 - Compute next 64 bits of binary expansion of $1/p$
 - For each 8-bit block b_0, \dots, b_7 , add some value to a table entry $t[i, b_i]$
- Afterwards: determine final result as a weighted sum of the table entries

Important: Inner loop body as cheap as possible.

Size of blocks chosen so that table fits into L1d cache.

Refinement

Refinement

Isabelle Refinement Framework and Isabelle-LLVM by Lammich:

- Start with a high-level view of the algorithm:
 - `nat`, `int` instead of `uint32_t`, `int32_t` etc.
 - Use abstract mathematical notions like “smallest prime factor of n ”, “ $\text{ord}_{\mathbb{Z}/p\mathbb{Z}}(n)$ ” without worrying about how to compute them

Refinement

Isabelle Refinement Framework and Isabelle-LLVM by Lammich:

- Start with a high-level view of the algorithm:
 - `nat`, `int` instead of `uint32_t`, `int32_t` etc.
 - Use abstract mathematical notions like “smallest prime factor of n ”, “ $\text{ord}_{\mathbb{Z}/p\mathbb{Z}}(n)$ ” without worrying about how to compute them
- Refine down to more concrete implementations, e.g. for/while loops to compute prime sieve

Refinement

Isabelle Refinement Framework and Isabelle-LLVM by Lammich:

- Start with a high-level view of the algorithm:
 - `nat`, `int` instead of `uint32_t`, `int32_t` etc.
 - Use abstract mathematical notions like “smallest prime factor of n ”, “ $\text{ord}_{\mathbb{Z}/p\mathbb{Z}}(n)$ ” without worrying about how to compute them
- Refine down to more concrete implementations, e.g. for/while loops to compute prime sieve
- Data refinement to fixed-width machine words (adding assumptions as needed)

Refinement

Isabelle Refinement Framework and Isabelle-LLVM by Lammich:

- Start with a high-level view of the algorithm:
 - `nat`, `int` instead of `uint32_t`, `int32_t` etc.
 - Use abstract mathematical notions like “smallest prime factor of n ”, “ $\text{ord}_{\mathbb{Z}/p\mathbb{Z}}(n)$ ” without worrying about how to compute them
- Refine down to more concrete implementations, e.g. for/while loops to compute prime sieve
- Data refinement to fixed-width machine words (adding assumptions as needed)

Example applications:

- IsaSAT by Fleury: fully verified SAT solver
- Lammich: verified sorting algorithms on par with C++ standard library

Refinement example: Factor Cache

On the abstract level:

definition "smallest_divisor n =
(if $n < 2 \vee \text{prime } n$ then 0 else LEAST d . $d \neq 1 \wedge d \text{ dvd } n$)"

Refinement example: Factor Cache

On the abstract level:

definition "smallest_divisor n =
(if $n < 2 \vee \text{prime } n$ then 0 else LEAST d . $d \neq 1 \wedge d \text{ dvd } n$)"

definition factor_cache_impl :: "nat \Rightarrow (nat \Rightarrow nat) nres" **where** ...

Refinement example: Factor Cache

On the abstract level:

definition "smallest_divisor n =

(if $n < 2 \vee \text{prime } n$ then 0 else LEAST d . $d \neq 1 \wedge d \text{ dvd } n$)"

definition factor_cache_impl :: "nat \Rightarrow (nat \Rightarrow nat) nres" **where** ...

lemma "factor_cache_impl $N \leq \text{SPEC } (\lambda a. \forall k < N. a \ k = \text{smallest_divisor } k)$ "

Refinement example: Factor Cache

On the abstract level:

definition "smallest_divisor n =

(if $n < 2 \vee \text{prime } n$ then 0 else LEAST d . $d \neq 1 \wedge d \text{ dvd } n$)"

definition factor_cache_impl :: "nat \Rightarrow (nat \Rightarrow nat) nres" **where** ...

lemma "factor_cache_impl $N \leq \text{SPEC } (\lambda a. \forall k < N. a \ k = \text{smallest_divisor } k)"$

First refinement: Replace function with list; record only entries for odd indices

definition factor_cache_impl2 :: "nat \Rightarrow nat list nres" **where** ...

Refinement example: Factor Cache

On the abstract level:

definition "smallest_divisor n =
(if $n < 2 \vee \text{prime } n$ then 0 else LEAST d . $d \neq 1 \wedge d \text{ dvd } n$)"

definition factor_cache_impl :: "nat \Rightarrow (nat \Rightarrow nat) nres" **where** ...

lemma "factor_cache_impl $N \leq \text{SPEC } (\lambda a. \forall k < N. a \ k = \text{smallest_divisor } k)"$

First refinement: Replace function with list; record only entries for odd indices

definition factor_cache_impl2 :: "nat \Rightarrow nat list nres" **where** ...

lemma " $3 \leq N \implies \text{factor_cache_impl2 } N \leq \Downarrow(\text{fc_rel } N) (\text{factor_cache_impl } N)"$

Refinement example: Factor Cache

On the abstract level:

definition "smallest_divisor n =

(if $n < 2 \vee \text{prime } n$ then 0 else LEAST d . $d \neq 1 \wedge d \text{ dvd } n$)"

definition factor_cache_impl :: "nat \Rightarrow (nat \Rightarrow nat) nres" **where** ...

lemma "factor_cache_impl $N \leq \text{SPEC } (\lambda a. \forall k < N. a \ k = \text{smallest_divisor } k)"$

First refinement: Replace function with list; record only entries for odd indices

definition factor_cache_impl2 :: "nat \Rightarrow nat list nres" **where** ...

lemma " $3 \leq N \implies \text{factor_cache_impl2 } N \leq \Downarrow(\text{fc_rel } N) (\text{factor_cache_impl } N)"$

Second refinement: Replace lists with arrays and *nat* with *word*:

sepref_def factor_cache_ll_impl :: "'w word \Rightarrow ('w word ptr \times 'w word) l1m" **where** ...

Refinement example: Factor Cache

On the abstract level:

definition "smallest_divisor $n =$

(if $n < 2 \vee \text{prime } n$ then 0 else LEAST $d. d \neq 1 \wedge d \text{ dvd } n$)"

definition factor_cache_impl :: "nat \Rightarrow (nat \Rightarrow nat) nres" **where** ...

lemma "factor_cache_impl $N \leq \text{SPEC } (\lambda a. \forall k < N. a \ k = \text{smallest_divisor } k)$ "

First refinement: Replace function with list; record only entries for odd indices

definition factor_cache_impl2 :: "nat \Rightarrow nat list nres" **where** ...

lemma " $3 \leq N \implies \text{factor_cache_impl2 } N \leq \Downarrow(\text{fc_rel } N) (\text{factor_cache_impl } N)$ "

Second refinement: Replace lists with arrays and *nat* with *word*:

sepref_def factor_cache_ll_impl :: "'w word \Rightarrow ('w word ptr \times 'w word) l1m" **where** ...

lemma "(factor_cache_ll_impl, factor_cache_impl2)

$\in \text{unat_assn}^k \rightarrow_a \text{array_assn unat_assn } \times_a \text{unat_assn}"$

Final correctness theorem

```
llvm_htriple
(
  ↑ll_pto okX okp ∧* ↑ll_pto numX nump ∧* ↑ll_pto denomX denomp ∧*
  ssize_assn thr thri ∧* ssize_assn depth depthi ∧* numnat_assn k ki
)

(bern_crt_impl_wrapper okp nump denomp thri depthi ki)

(λ_. EXS oki numi denomi num denom.
  ↑ll_pto oki okp ∧* ↑ll_pto numi nump ∧* ↑ll_pto denomi denomp ∧*
  mpzb_assn num numi ∧* mpzb_assn denom denomi ∧*
  ↑((oki ≠ 0 → denom = int (bernoulli_denom k) ∧ num = bernoulli_num k) ∧
    (k ≤ 105946388 → oki ≠ 0))
)
```

Final correctness theorem

```
llvm_htriple
(
  1ll_pto okX okp  $\wedge^*$  1ll_pto numX nump  $\wedge^*$  1ll_pto denomX denomp  $\wedge^*$ 
  ssize_assn thr thri  $\wedge^*$  ssize_assn depth depthi  $\wedge^*$  numnat_assn k ki
)

(bern_crt_impl_wrapper okp nump denomp thri depthi ki)

( $\lambda\_.$  EXS oki numi denomi num denom.
  1ll_pto oki okp  $\wedge^*$  1ll_pto numi nump  $\wedge^*$  1ll_pto denomi denomp  $\wedge^*$ 
  mpz_b_assn num numi  $\wedge^*$  mpz_b_assn denom denomi  $\wedge^*$ 
   $\uparrow((\text{oki} \neq 0 \rightarrow \text{denom} = \text{int}(\text{bernoulli\_denom } k) \wedge \text{num} = \text{bernoulli\_num } k) \wedge$ 
     $(k \leq 105946388 \rightarrow \text{oki} \neq 0))$ 
)
```

Challenges

Lots of algorithmic components

- Prime sieve, factoring integers

Lots of algorithmic components

- Prime sieve, factoring integers Done.

Lots of algorithmic components

- Prime sieve, factoring integers Done.
- Group-theoretic computations in $\mathbb{Z}/p\mathbb{Z}$: find generators, compute $\text{ord}(x)$

Lots of algorithmic components

- Prime sieve, factoring integers Done.
- Group-theoretic computations in $\mathbb{Z}/p\mathbb{Z}$: find generators, compute $\text{ord}(x)$ Done. (using factorisation)

Lots of algorithmic components

- Prime sieve, factoring integers Done.
- Group-theoretic computations in $\mathbb{Z}/p\mathbb{Z}$: find generators, compute $\text{ord}(x)$ Done. (using factorisation)
- Efficient computations modulo p (e.g. Montgomery form)

Lots of algorithmic components

- Prime sieve, factoring integers **Done.**
- Group-theoretic computations in $\mathbb{Z}/p\mathbb{Z}$: find generators, compute $\text{ord}(x)$ **Done.** (using factorisation)
- Efficient computations modulo p (e.g. Montgomery form) **Done.**

Lots of algorithmic components

- Prime sieve, factoring integers **Done.**
- Group-theoretic computations in $\mathbb{Z}/p\mathbb{Z}$: find generators, compute $\text{ord}(x)$ **Done.** (using factorisation)
- Efficient computations modulo p (e.g. Montgomery form) **Done.**
- Floating-point computations for bounds etc.

Lots of algorithmic components

- Prime sieve, factoring integers **Done.**
- Group-theoretic computations in $\mathbb{Z}/p\mathbb{Z}$: find generators, compute $\text{ord}(x)$ **Done.** (using factorisation)
- Efficient computations modulo p (e.g. Montgomery form) **Done.**
- Floating-point computations for bounds etc.
Replaced with fixed-point.

Lots of algorithmic components

- Prime sieve, factoring integers **Done.**
- Group-theoretic computations in $\mathbb{Z}/p\mathbb{Z}$: find generators, compute $\text{ord}(x)$ **Done.** (using factorisation)
- Efficient computations modulo p (e.g. Montgomery form) **Done.**
- Floating-point computations for bounds etc.
Replaced with fixed-point.
- Computing the binary fraction expansion of $1/p$

Lots of algorithmic components

- Prime sieve, factoring integers **Done.**
- Group-theoretic computations in $\mathbb{Z}/p\mathbb{Z}$: find generators, compute $\text{ord}(x)$ **Done.** (using factorisation)
- Efficient computations modulo p (e.g. Montgomery form) **Done.**
- Floating-point computations for bounds etc.
Replaced with fixed-point.
- Computing the binary fraction expansion of $1/p$ **Done.**

Lots of algorithmic components

- Prime sieve, factoring integers **Done.**
- Group-theoretic computations in $\mathbb{Z}/p\mathbb{Z}$: find generators, compute $\text{ord}(x)$ **Done.** (using factorisation)
- Efficient computations modulo p (e.g. Montgomery form) **Done.**
- Floating-point computations for bounds etc.
Replaced with fixed-point.
- Computing the binary fraction expansion of $1/p$ **Done.**
- Fast Chinese Remaindering

Lots of algorithmic components

- Prime sieve, factoring integers **Done.**
- Group-theoretic computations in $\mathbb{Z}/p\mathbb{Z}$: find generators, compute $\text{ord}(x)$ **Done.** (using factorisation)
- Efficient computations modulo p (e.g. Montgomery form) **Done.**
- Floating-point computations for bounds etc.
Replaced with fixed-point.
- Computing the binary fraction expansion of $1/p$ **Done.**
- Fast Chinese Remaindering
Done. Using remainder trees.

Lots of algorithmic components

- Prime sieve, factoring integers **Done.**
- Group-theoretic computations in $\mathbb{Z}/p\mathbb{Z}$: find generators, compute $\text{ord}(x)$ **Done.** (using factorisation)
- Efficient computations modulo p (e.g. Montgomery form) **Done.**
- Floating-point computations for bounds etc.
Replaced with fixed-point.
- Computing the binary fraction expansion of $1/p$ **Done.**
- Fast Chinese Remaindering
Done. Using remainder trees.
- Arbitrary-precision integers

Lots of algorithmic components

- Prime sieve, factoring integers **Done.**
- Group-theoretic computations in $\mathbb{Z}/p\mathbb{Z}$: find generators, compute $\text{ord}(x)$ **Done.** (using factorisation)
- Efficient computations modulo p (e.g. Montgomery form) **Done.**
- Floating-point computations for bounds etc.
Replaced with fixed-point.
- Computing the binary fraction expansion of $1/p$ **Done.**
- Fast Chinese Remaindering
Done. Using remainder trees.
- Arbitrary-precision integers
Out of scope. We use GMP as trusted component.

Lots of algorithmic components

- Prime sieve, factoring integers **Done.**
- Group-theoretic computations in $\mathbb{Z}/p\mathbb{Z}$: find generators, compute $\text{ord}(x)$ **Done.** (using factorisation)
- Efficient computations modulo p (e.g. Montgomery form) **Done.**
- Floating-point computations for bounds etc.
Replaced with fixed-point.
- Computing the binary fraction expansion of $1/p$ **Done.**
- Fast Chinese Remaindering
Done. Using remainder trees.
- Arbitrary-precision integers
Out of scope. We use GMP as trusted component. **Future work?**

Lots of mathematical background

- Definition of Bernoulli numbers and basic properties

Lots of mathematical background

- Definition of Bernoulli numbers and basic properties *Already there.*

Lots of mathematical background

- Definition of Bernoulli numbers and basic properties *Already there.*
- Bounds for Bernoulli numbers

Lots of mathematical background

- Definition of Bernoulli numbers and basic properties *Already there.*
- Bounds for Bernoulli numbers *Easy.*

Lots of mathematical background

- Definition of Bernoulli numbers and basic properties *Already there.*
- Bounds for Bernoulli numbers *Easy.*
- Kummer/Voronoi congruence and Harvey's tweaks

Lots of mathematical background

- Definition of Bernoulli numbers and basic properties *Already there.*
- Bounds for Bernoulli numbers *Easy.*
- Kummer/Voronoi congruence and Harvey's tweaks *Done.*

Lots of mathematical background

- Definition of Bernoulli numbers and basic properties *Already there.*
- Bounds for Bernoulli numbers *Easy.*
- Kummer/Voronoi congruence and Harvey's tweaks *Done.*
- Concrete bounds for the Chebyshev ϑ function:

$$\vartheta(x) = \sum_{p \leq x} \ln p \geq 0.82x \quad \text{for } x \geq 97$$

Lots of mathematical background

- Definition of Bernoulli numbers and basic properties *Already there.*
- Bounds for Bernoulli numbers *Easy.*
- Kummer/Voronoi congruence and Harvey's tweaks *Done.*
- Concrete bounds for the Chebyshev ϑ function:

$$\vartheta(x) = \sum_{p \leq x} \ln p \geq 0.82x \quad \text{for } x \geq 97$$

Needed for our a-priori estimate of how many primes to sieve.

Lots of mathematical background

- Definition of Bernoulli numbers and basic properties *Already there.*
- Bounds for Bernoulli numbers *Easy.*
- Kummer/Voronoi congruence and Harvey's tweaks *Done.*
- Concrete bounds for the Chebyshev ϑ function:

$$\vartheta(x) = \sum_{p \leq x} \ln p \geq 0.82x \quad \text{for } x \geq 97$$

Needed for our a-priori estimate of how many primes to sieve.
Done.

Lines of Code

Component	LOC
Voronoi/Kummer	2300
Prime bounds	1800

Lines of Code

Component	LOC
Voronoi/Kummer	2300
Prime bounds	1800
Fixed-point \log_2	1000
Binary fraction expansion	900
Montgomery multiplication	2300
Prime sieve, order, generators	2700
Fast Chinese Remaindering	3800
Other	900

Lines of Code

Component	LOC
Voronoi/Kummer	2300
Prime bounds	1800
Fixed-point \log_2	1000
Binary fraction expansion	900
Montgomery multiplication	2300
Prime sieve, order, generators	2700
Fast Chinese Remaindering	3800
Other	900

Component	LOC
Additions to sepref package	4300
GMP bindings	1700

Lines of Code

Component	LOC
Voronoi/Kummer	2300
Prime bounds	1800
Fixed-point \log_2	1000
Binary fraction expansion	900
Montgomery multiplication	2300
Prime sieve, order, generators	2700
Fast Chinese Remaindering	3800
Other	900

Component	LOC
Additions to sepref package	4300
GMP bindings	1700
Abstract algorithm	1500
Concrete algorithm	8500
Total	31700

Low-level optimisations

We can compute the binary fraction expansion of $1/p$ in 64-bit chunks by letting `bitbuf = 1` and then repeating

```
output ((bitbuf << 64) / p)
bitbuf = ((bitbuf << 64) % p)
```

Low-level optimisations

We can compute the binary fraction expansion of $1/p$ in 64-bit chunks by letting `bitbuf = 1` and then repeating

```
output ((bitbuf << 64) / p)
bitbuf = ((bitbuf << 64) % p)
```

But: Division is expensive. Therefore, we instead precompute a 128-bit fixed-point approximation `invp` of $1/p$ and compute

- `quotient (bitbuf << 64) / p` via `bitbuf_new = (invp * bitbuf) >> 64`

Low-level optimisations

We can compute the binary fraction expansion of $1/p$ in 64-bit chunks by letting `bitbuf = 1` and then repeating

```
output ((bitbuf << 64) / p)
bitbuf = ((bitbuf << 64) % p)
```

But: Division is expensive. Therefore, we instead precompute a 128-bit fixed-point approximation `invp` of $1/p$ and compute

- quotient `(bitbuf << 64) / p` via `bitbuf_new = (invp * bitbuf) >> 64`
- remainder `(bitbuf << 64) % p` via `-p * bitbuf_new`.

Low-level optimisations

We can compute the binary fraction expansion of $1/p$ in 64-bit chunks by letting `bitbuf = 1` and then repeating

```
output ((bitbuf << 64) / p)
bitbuf = ((bitbuf << 64) % p)
```

But: Division is expensive. Therefore, we instead precompute a 128-bit fixed-point approximation `invp` of $1/p$ and compute

- quotient `(bitbuf << 64) / p` via `bitbuf_new = (invp * bitbuf) >> 64`
- remainder `(bitbuf << 64) % p` via `-p * bitbuf_new`.

There is a small chance that the result is off-by-one, which we have to detect and correct accordingly.

Low-level optimisations

```
// bitbuf := 2 ^ 64 * x / p,    x := 2 ^ 64 * x % p
uint64_t bitbuf = (invp * (__uint128_t) x) >> 64;
x = -p * bitbuf;

// There is a small chance that our quotient is actually too small by 1; we detect this here.
if (__builtin_expect(x >= p, 0)) [[unlikely]] {
    bitbuf++;
    x -= p;
}
```


Low-level optimisations

Proving such low-level code correct requires

- understanding what the right high-level model is

Low-level optimisations

Proving such low-level code correct requires

- understanding what the right high-level model is
- figuring out preconditions, both abstractly and regarding overflow etc.

Low-level optimisations

Proving such low-level code correct requires

- understanding what the right high-level model is
- figuring out preconditions, both abstractly and regarding overflow etc.

Lessons learnt:

- Proving absence of overflow can be painful.

Low-level optimisations

Proving such low-level code correct requires

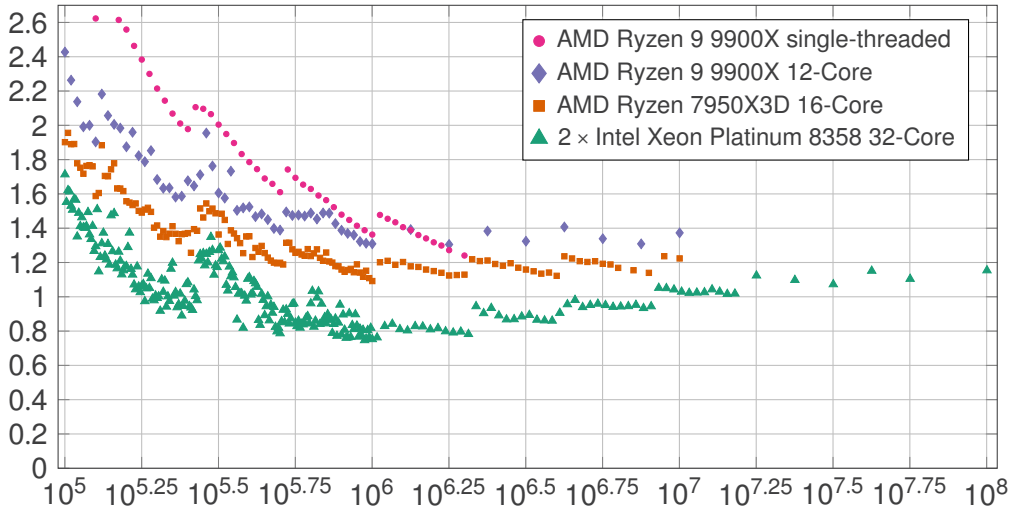
- understanding what the right high-level model is
- figuring out preconditions, both abstractly and regarding overflow etc.

Lessons learnt:

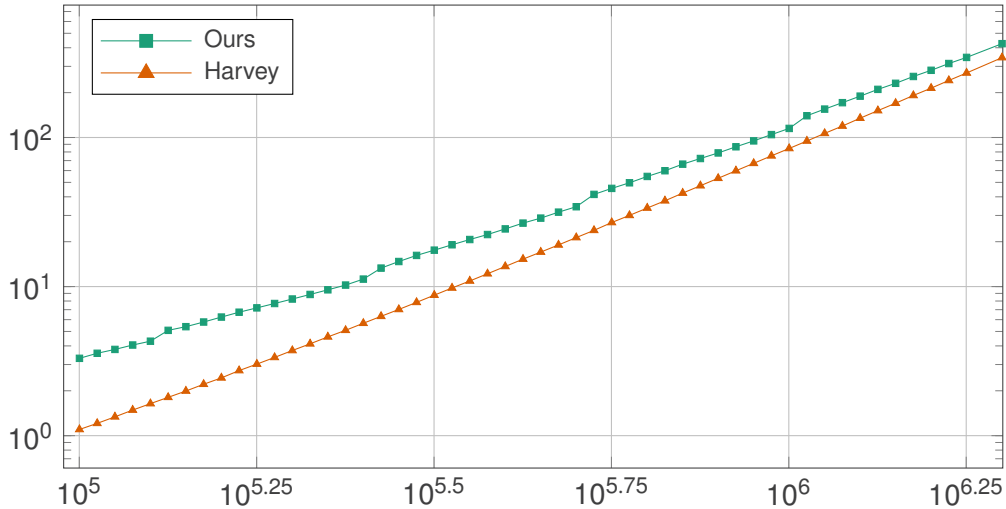
- Proving absence of overflow can be painful.
- **Advantage:** Using an ITP helps you figure out the range in which the algorithm does not produce overflow.

Evaluation

Time ratio (Ours / Harvey)



Time (s)



Conclusion

Conclusion

We verified a complex and challenging mathematical algorithm all the way down to LLVM code.

Conclusion

We verified a complex and challenging mathematical algorithm all the way down to LLVM code.

Made various additions to Isabelle-LLVM; exposed some weak points (e.g. sharing read-only access among parallel threads).

Conclusion

We verified a complex and challenging mathematical algorithm all the way down to LLVM code.

Made various additions to Isabelle-LLVM; exposed some weak points (e.g. sharing read-only access among parallel threads).

Performance of resulting LLVM code not quite on par with Harvey's unverified C++ code, but quite close (especially for large inputs).

Conclusion

We verified a complex and challenging mathematical algorithm all the way down to LLVM code.

Made various additions to Isabelle-LLVM; exposed some weak points (e.g. sharing read-only access among parallel threads).

Performance of resulting LLVM code not quite on par with Harvey's unverified C++ code, but quite close (especially for large inputs).

Closing the gap would require in-depth microbenchmarking.

Conclusion

We verified a complex and challenging mathematical algorithm all the way down to LLVM code.

Made various additions to Isabelle-LLVM; exposed some weak points (e.g. sharing read-only access among parallel threads).

Performance of resulting LLVM code not quite on par with Harvey's unverified C++ code, but quite close (especially for large inputs).

Closing the gap would require in-depth microbenchmarking.

We are already *much* faster than Mathematica's `BernoulliB` algorithm!