



# A Certified Proof Checker for Deep Neural Network Verification in Imandra

Remi Desmartin\*, **Omri Isac\***, Grant Passmore, Ekaterina  
Komendantskaya, Kathrin Stark, Guy Katz

ITP, October 2025



**Remi Desmartin, PhD student**

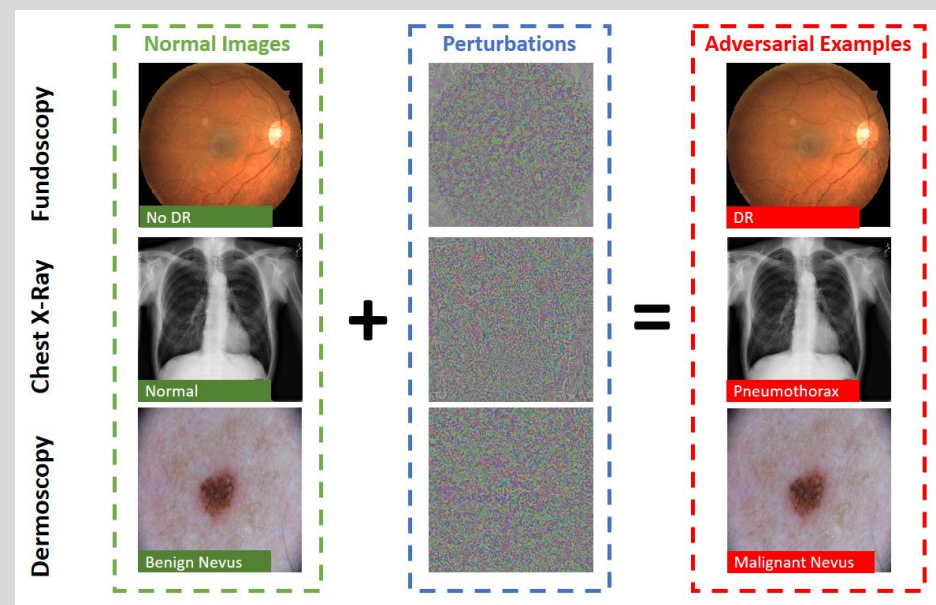
**Heriot-Watt University**

Deep Neural Networks (DNNs) accomplish groundbreaking results in many fields, **including safety-critical.**



Unlike traditional software, **DNNs are opaque functions,** **trained** over a large set of input and output examples.

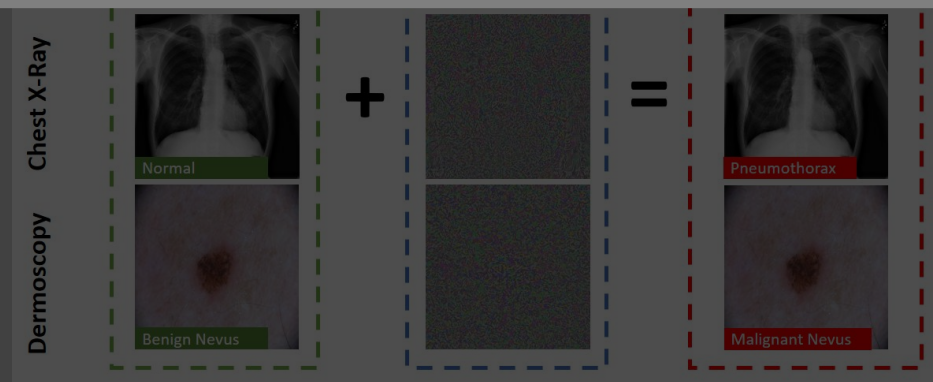
This may lead to undesirable behavior of DNNs, which is **hard to predict and potentially dangerous.**



Understanding Adversarial Attacks on Deep Learning Based Medical Image Analysis Systems, Ma et al., 2019

This may lead to undesirable behavior of DNNs, which is **hard to predict** and **potentially dangerous**.

Can we **rule out** possibility of (some) errors?  
Or **find hidden** errors, if exist




Understanding Adversarial Attacks on Deep Learning Based Medical Image Analysis Systems, Ma et al., 2019

The verification community developed **various DNN verifiers**, based on SMT, abstract interpretation, MILP and more.




Verifier



DNN

DNN verifiers are prone to bugs and **may be numerically unstable, which can be exploited** [Jia & Rinard, 2021].



They are complex software and performance-oriented, thus **their verification is very difficult.**


## Produce certificates of the verifiers' results:

- Checked using an independent, trusted and simple checker.
- Witness the correctness of the verifier or reveal errors.



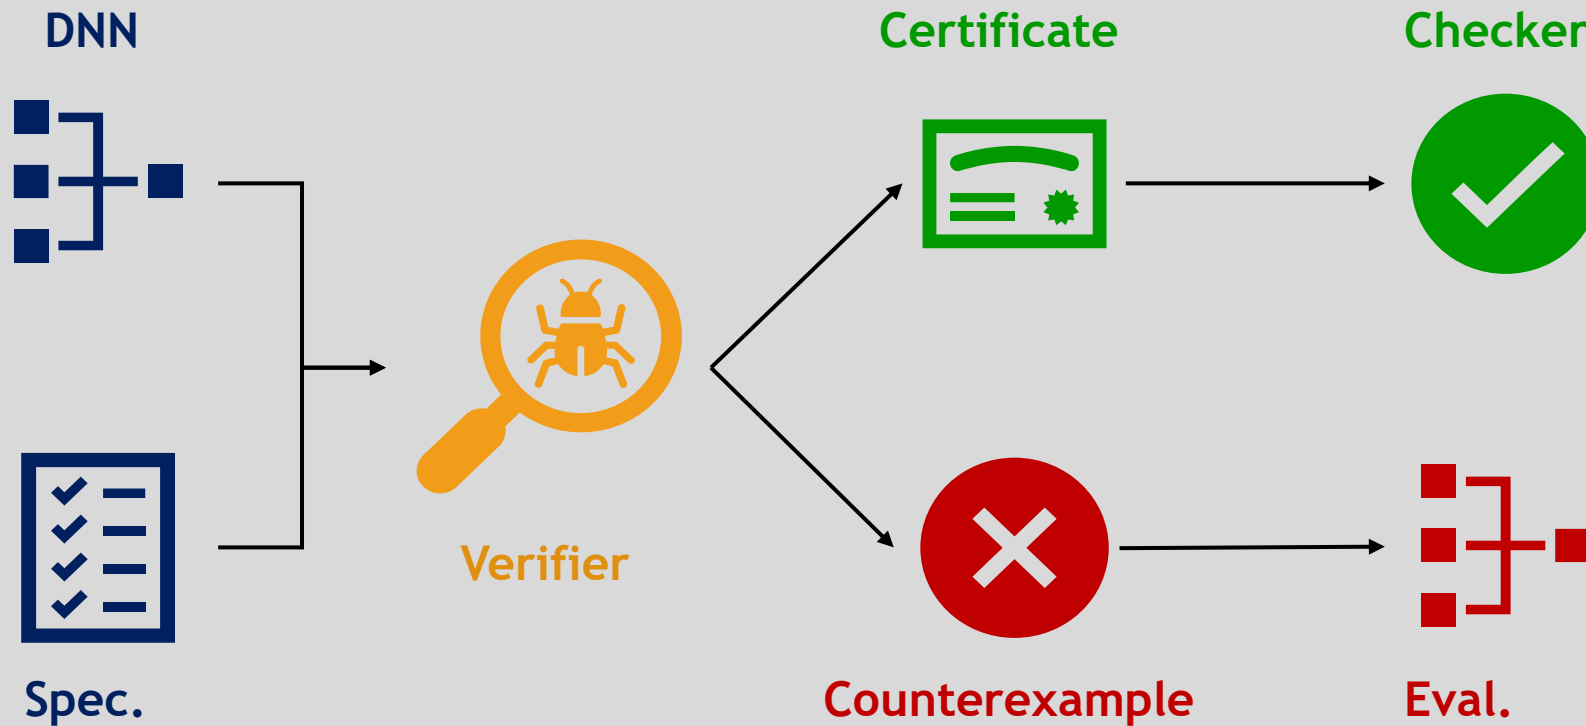
A common practice in SAT and SMT communities.






Trusting this scheme requires implementing a **reliable checker**, e.g., **formally verified in an ITP**.


# The Model of Trust



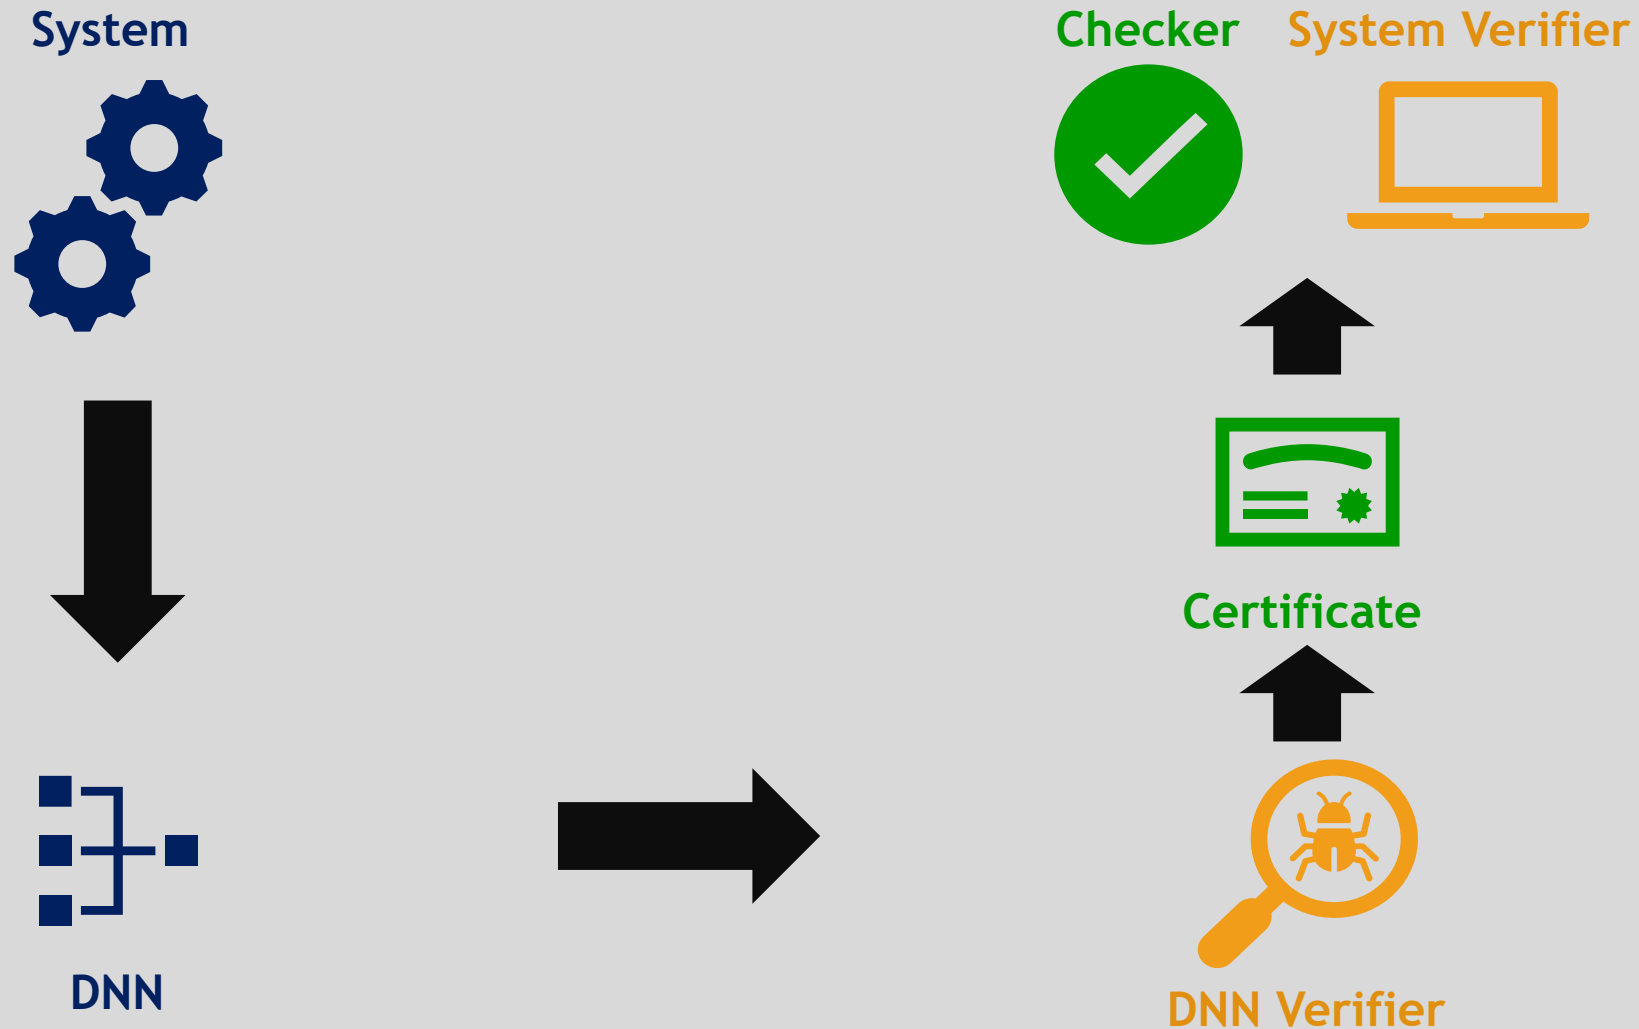
# The Model of Trust



# The Model of Trust




# Integration in System Verification



# Background


# Deep Neural Networks

- Layered function:
  - **Affine:**  $b_i = \sum_j \mathbf{w}_{i,j} \cdot f_j + c_i$
  - **Nonlinear activations:**  $f_i = \mathbf{a}_i(b_i)$
- The ReLU activation:



# DNN Verification

- Given a DNN  $\mathcal{N}$  an input property  $\varphi(x)$  and an output property  $\psi(y)$ , **decide whether there exists an input vector  $x$  such that  $\varphi(x) \rightarrow \psi(\mathcal{N}(x))$ .**
  - If exists, the problem is **satisfiable (SAT)**;
  - Otherwise **unsatisfiable (UNSAT)**.





# DNN Verification

- Assuming QF linear properties, and ReLU activations → **linear + piecewise linear constraints:**
- Find assignment for  $V \in \mathbb{R}^n$


- **Linear:**

$$A \cdot V = \bar{0} \quad [A \in \mathbb{R}^{m \times n}]$$
$$l \leq V \leq u \quad [l, u \in \mathbb{R}^n]$$

- **Piecewise Linear:**

$$f_i = \text{ReLU}(b_i) \quad [f_i, b_i \in V]$$


# Verification Query Example



$$\varphi(x) := x \in [0, 1]^2$$

$$\neg\psi(y) := y \in [4, 5]$$

# Verification Query Example



$$\varphi(x) := x \in [0, 1]^2$$

$$\neg\psi(y) := y \in [4, 5]$$

Variables:

$$V = [x_1 \quad x_2 \quad b_1 \quad b_2 \quad f_1 \quad f_2 \quad aux_1 \quad aux_2 \quad y]$$

Tableau:

$$A = \begin{bmatrix} 2 & 1 & -1 & 0 & 0 & 0 & 0 & 0 & 0 \\ -1 & 1 & 0 & -1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & -1 & 2 & 0 & 0 & -1 \\ 0 & 0 & -1 & 0 & 1 & 0 & -1 & 0 & 0 \\ 0 & 0 & 0 & -1 & 0 & 1 & 0 & -1 & 0 \end{bmatrix}$$

Bounds:

$$u = [1 \quad 1 \quad 3 \quad 1 \quad 3 \quad 1 \quad 3 \quad 2 \quad 5]$$

$$l = [0 \quad 0 \quad 0 \quad -1 \quad 0 \quad 0 \quad 0 \quad 0 \quad 4]$$


ReLU's:

$$f_1 = \text{ReLU}(b_1)$$

$$f_2 = \text{ReLU}(b_2)$$

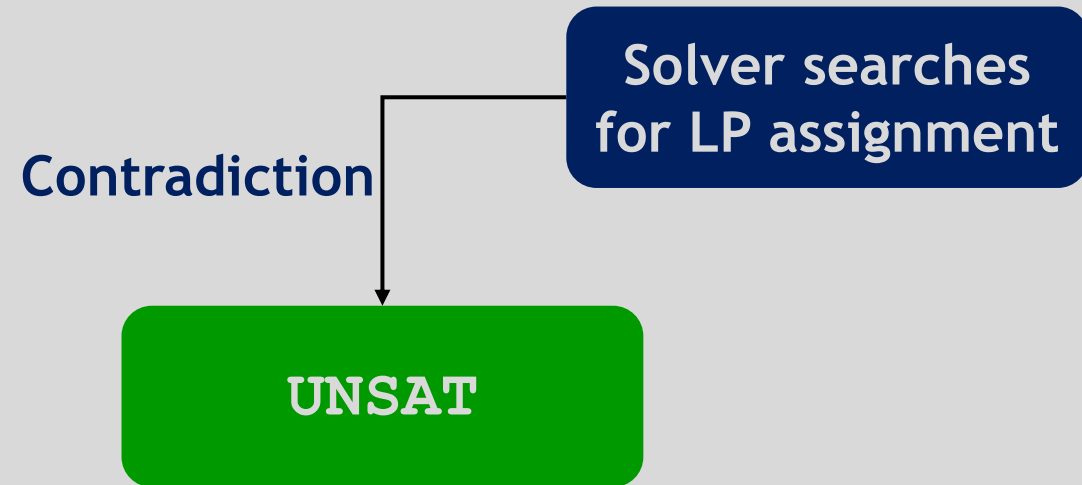
# DNN Verification Algorithms

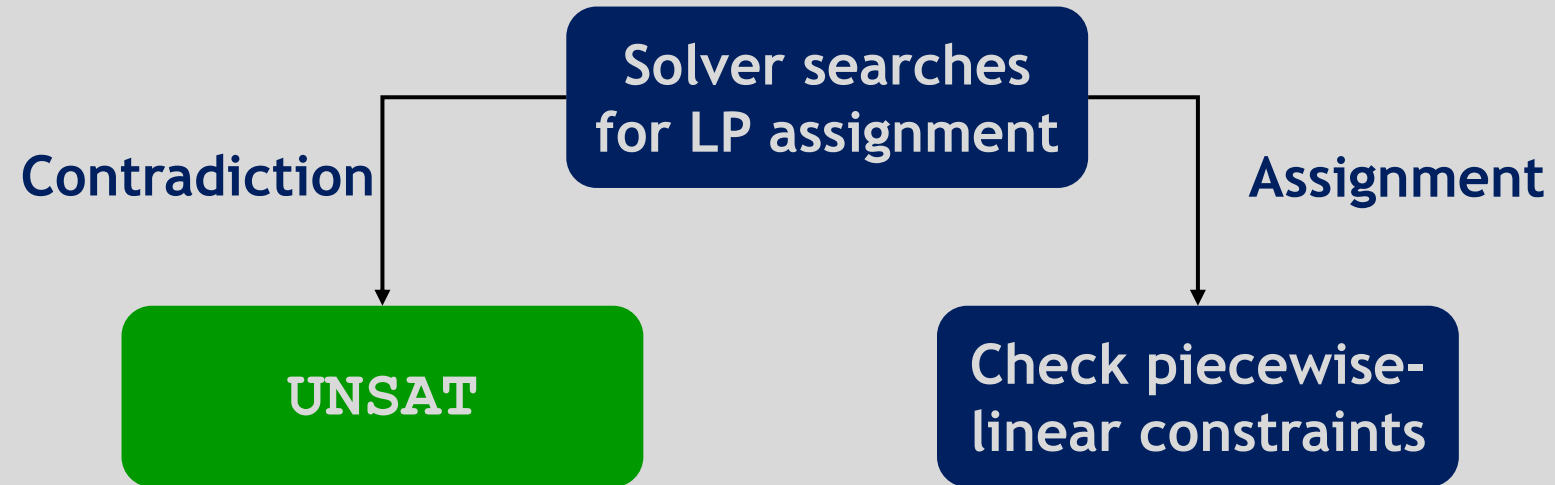
- Many verification algorithms **search for a satisfying assignment** with LP solvers + splitting approach.

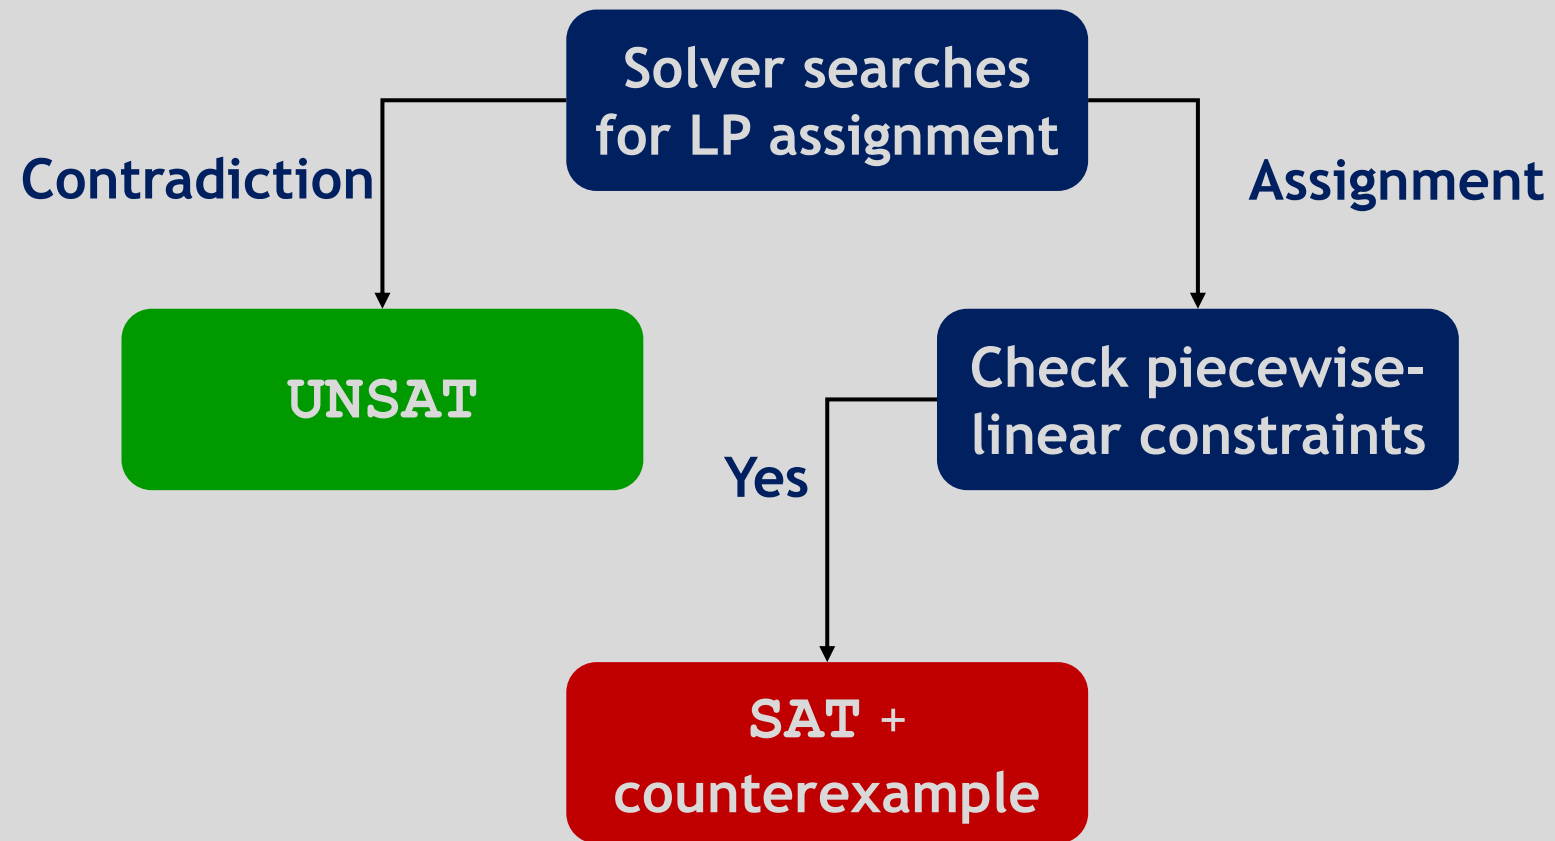


$$\begin{aligned} A \cdot V &= \bar{0} \\ l &\leq V \leq u \\ f_i &= \text{ReLU}(b_i) \end{aligned}$$

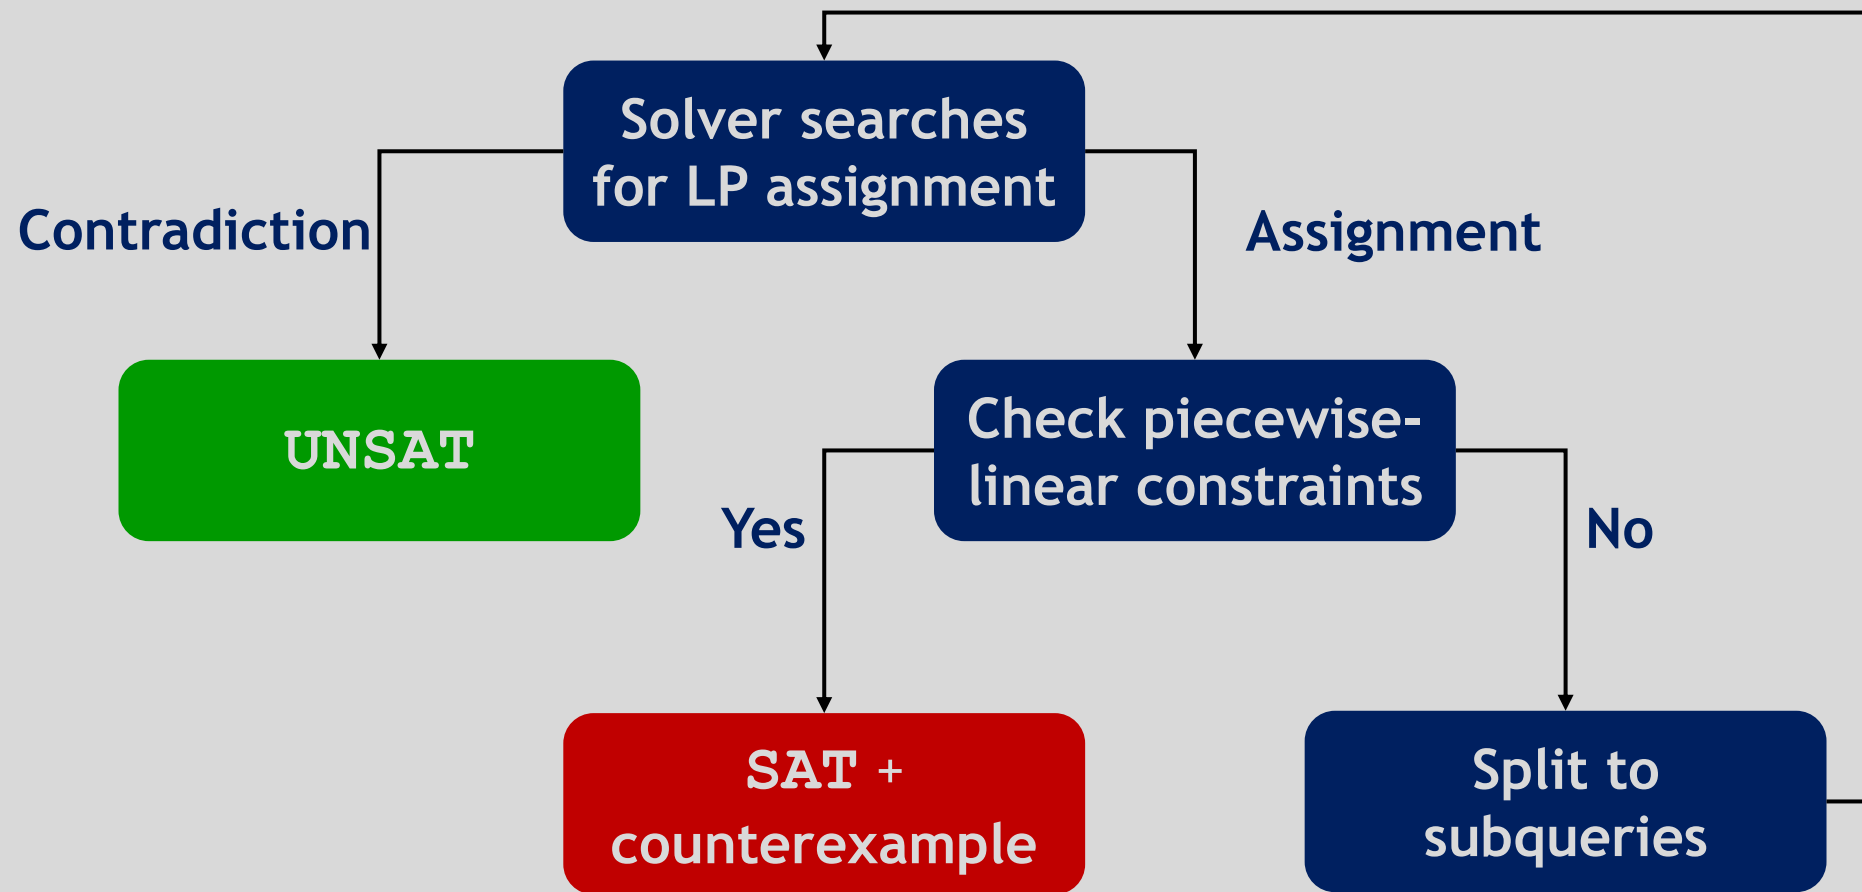
**Solver searches  
for LP assignment**

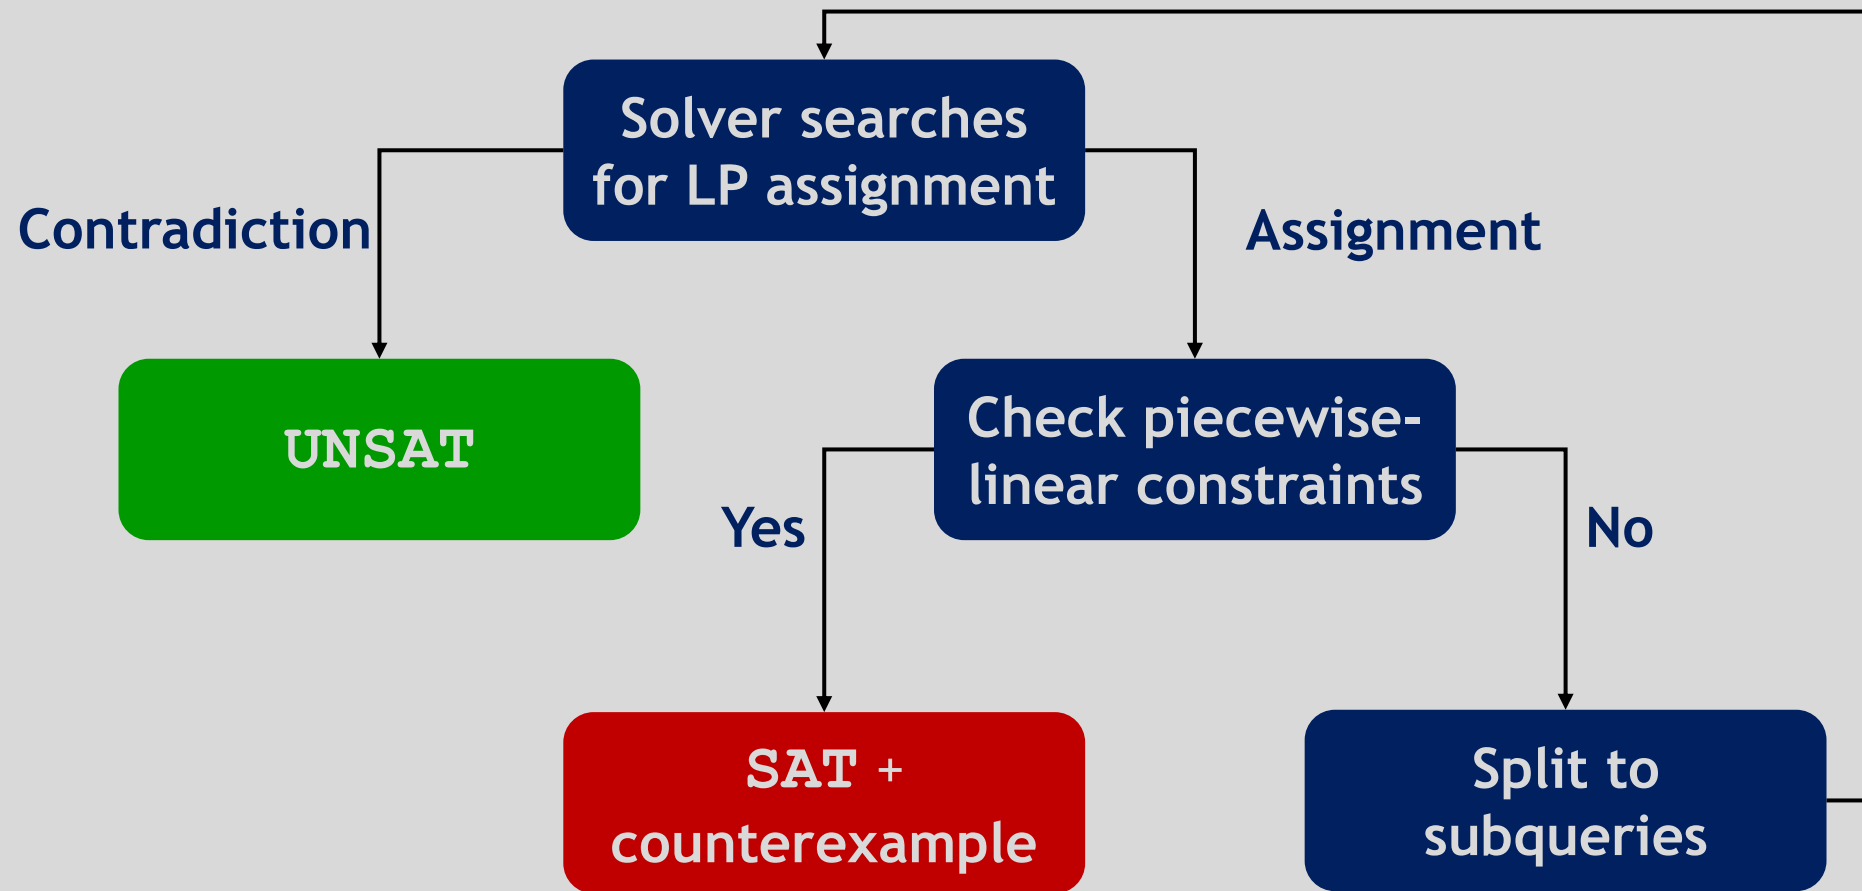













Splitting creates a **tree structure** to the search.  
The query is UNSAT if and only if all leaves are UNSAT.

# Certificate Production

- Certifying SAT is straightforward; **Certifying UNSAT is more complicated**, due to NP-Hardness [Katz et. al, '17; Sälzer and Lange, '21].
- **Marabou DNN verifier** produces and checks UNSAT proofs (C++).

# The Certificate Tree



In the **certificate-tree**, nodes represent case-splits; leaves contain proofs of contradictions, which are represented by a vector.

# The Farkas Lemma (Variant)

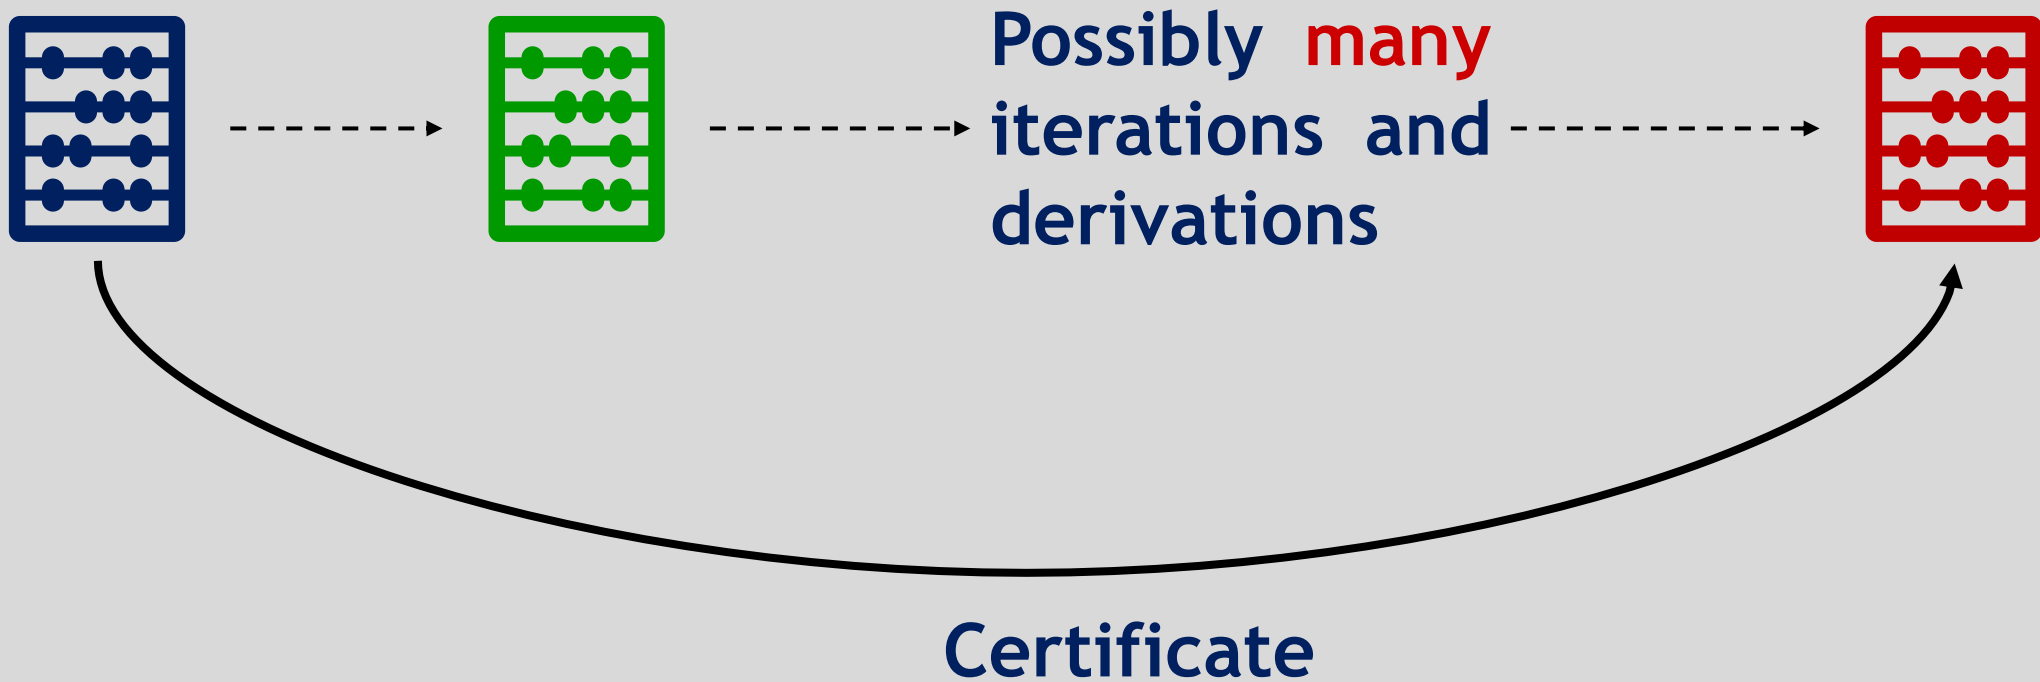


Gyula Farkas


For any set of linear equations  $AV = 0$  where  $l \leq V \leq u$  **exactly one holds**:

- There **exists a solution** (row vector)  $l \leq V \leq u$  s.t.  $AV = 0$ .
- There **exists a refutation** (column vector)  $y$  s.t.  $\forall l \leq V \leq u: y^T AV < 0$ .

Refutation can be **constructed during execution of DNN verifier**.



# The Certificate Tree



Certificate checking consists of **traversing the tree**, **checking splits are covering** and **checking leaves** using the Farkas lemma.

# Certified Checking




# Imandra



- An OCaml-based **FPL** and an **ITP**.
- Supports **arbitrary precision** real arithmetic.

```
theorem farkas_unsat (s : system) (x : var_vect) (c : certificate) =  
  well_formed s x && check_cert s c  
  ==>  
  eval_system s x = false  
[@@by [% use cert_is_neg s c x]  
@> [% use solution_is_not_neg s c x]  
@> auto] [@@fc]
```

# Data Representation



**Variables:**

$$V = [x_1 \ x_2 \ b_1 \ b_2 \ f_1 \ f_2 \ aux_1 \ aux_2 \ y]$$

**Tableau:**

$$A = \begin{bmatrix} 2 & 1 & -1 & 0 & 0 & 0 & 0 & 0 & 0 \\ -1 & 1 & 0 & -1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & -1 & 2 & 0 & 0 & -1 \\ 0 & 0 & -1 & 0 & 1 & 0 & -1 & 0 & 0 \\ 0 & 0 & 0 & -1 & 0 & 1 & 0 & -1 & 0 \end{bmatrix}$$

**Bounds:**

$$u = [1 \ 1 \ 3 \ 1 \ 3 \ 1 \ 3 \ 2 \ 5]$$

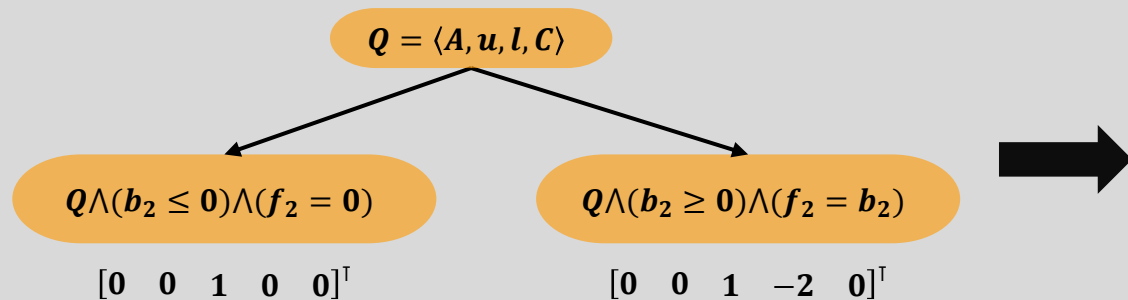
$$l = [0 \ 0 \ 0 \ -1 \ 0 \ 0 \ 0 \ 0 \ 4]$$

**ReLU:**

$$f_1 = \text{ReLU}(b_1)$$

$$f_2 = \text{ReLU}(b_2)$$

# Data Representation



```

type proofTree =
    | Leaf of real list
    | Node of Split * proofTree * proofTree
type Split = Relu of int * int * int
    
```

Variables:

$V = [x_1 \ x_2 \ b_1 \ b_2 \ f_1 \ f_2 \ aux_1 \ aux_2 \ y]$

Tableau:

$A = \begin{bmatrix} 2 & 1 & -1 & 0 & 0 & 0 & 0 & 0 & 0 \\ -1 & 1 & 0 & -1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & -1 & 2 & 0 & 0 & -1 \\ 0 & 0 & -1 & 0 & 1 & 0 & -1 & 0 & 0 \\ 0 & 0 & 0 & -1 & 0 & 1 & 0 & -1 & 0 \end{bmatrix}$



Bounds:

$u = [1 \ 1 \ 3 \ 1 \ 3 \ 1 \ 3 \ 2 \ 5]$   
 $l = [0 \ 0 \ 0 \ -1 \ 0 \ 0 \ 0 \ 0 \ 4]$

ReLU:

$f_1 = \text{ReLU}(b_1)$   
 $f_2 = \text{ReLU}(b_2)$

# Data Representation



```

type proofTree =
  | Leaf of real list
  | Node of Split * proofTree * proofTree
type Split = Relu of int * int * int
  
```

Variables:

$V = [x_1 \ x_2 \ b_1 \ b_2 \ f_1 \ f_2 \ aux_1 \ aux_2 \ y]$

Tableau:


$$A = \begin{bmatrix} 2 & 1 & -1 & 0 & 0 & 0 & 0 & 0 & 0 \\ -1 & 1 & 0 & -1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & -1 & 2 & 0 & 0 & -1 \\ 0 & 0 & -1 & 0 & 1 & 0 & -1 & 0 & 0 \\ 0 & 0 & 0 & -1 & 0 & 1 & 0 & -1 & 0 \end{bmatrix}$$

Bounds:

$u = [1 \ 1 \ 3 \ 1 \ 3 \ 1 \ 3 \ 2 \ 5]$   
 $l = [0 \ 0 \ 0 \ -1 \ 0 \ 0 \ 0 \ 0 \ 4]$

ReLU's:

$f_1 = \text{ReLU}(b_1)$   
 $f_2 = \text{ReLU}(b_2)$



```

type poly = Real.t list
type expr =
  | Eq of poly
  | Geq of poly
type system = expr list
  
```


# The Checking Function

```
let rec check_tree tableau upper_bounds lower_bounds constraints
proof_node =
match proof_node with
| Proof_tree.Leaf contradiction ->
    check_contradiction ... (* applies the Farkas lemma *)
| Proof_tree.Node ( split, left, right ) ->
    let valid_split = ... in
    let valid_children = ... in (* includes recursive calls *)
    valid_split && valid_children
```


# Proof Main Components



**Farkas lemma  
(UNSAT direction)**



**Type reduction  
soundness**



**ReLU splits are  
covering**

# The Farkas Lemma (Generalized)




Gyula Farkas

Given a polynomial system  $\bigwedge_{i=1}^N (p_i(x) = 0) \wedge \bigwedge_{i=1}^K (q_i(x) \geq 0)$  **exactly one holds:**

- There **exists a solution.**
- There **exists a refutation:**  $\mathbb{I} \in \mathbb{R}^N, \mathbb{C} \in \mathbb{R}_{\geq 0}^K$  with  $\mathbb{I} \cdot p_i + \mathbb{C} \cdot q_j < 0$ .

# The Farkas Lemma (Generalized)



Given a polynomial system  $\bigwedge_{i=1}^N (p_i(x) = 0) \wedge \bigwedge_{i=1}^K (q_i(x) \geq 0)$  **exactly one holds:**

```
theorem farkas_unsat (s : system) (x : var_vect) (c : certificate) =  
  well_formed s x && check_cert s c  
  ==>  
  eval_system s x = false
```



Gyula Farkas



# Data Representation

## Matrix and vectors

Variables:

$$V = [x_1 \ x_2 \ b_1 \ b_2 \ f_1 \ f_2 \ aux_1 \ aux_2 \ y]$$

Tableau:

$$A = \begin{bmatrix} 2 & 1 & -1 & 0 & 0 & 0 & 0 & 0 & 0 \\ -1 & 1 & 0 & -1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & -1 & 2 & 0 & 0 & -1 \\ 0 & 0 & -1 & 0 & 1 & 0 & -1 & 0 & 0 \\ 0 & 0 & 0 & -1 & 0 & 1 & 0 & -1 & 0 \end{bmatrix}$$

Bounds:

$$u = [1 \ 1 \ 3 \ 1 \ 3 \ 1 \ 3 \ 2 \ 5] \\ l = [0 \ 0 \ 0 \ -1 \ 0 \ 0 \ 0 \ 0 \ 4]$$

mk\_system



## System of polynomial equations and inequalities

Tableau:

$$2x_1 + x_2 - b_1 = 0$$

...

Bounds:

$$1 - x_1 \geq 0$$

$$x_1 \geq 0$$

...

If the **system** is UNSAT, so is the LP using the matrix and vectors.

# Data Representation

Matrix and vectors

System of polynomial Equations and

```
lemma soundness_check_cert_composition tableau upper_bounds lower_bounds x =  
  let sys = mk_system (mk_eq_constraints tableau) upper_bounds lower_bounds in  
  well_formed_tableau_bounds tableau upper_bounds lower_bounds &&  
  List.length x = List.length (List.hd tableau) &&  
  not (eval_system sys x)  
==>  
not (is_in_kernel tableau x) || not (bounded x upper_bounds lower_bounds)
```

...

If the **system** is UNSAT, so is are the input matrix and vectors.


# Soundness of Leaf Checking

If **leaf checking** returns true, then there is **no solution** to underlying **linear query**

# Soundness of Leaf Checking

```
lemma soundness_leaf tableau upper_bounds lower_bounds relu_constraints x
proof_tree =
  match proof_tree with
  | Node _ -> true
  | Leaf contradiction ->
    well_formed_tableau_bounds tableau upper_bounds lower_bounds &&
    List.length x = List.length (List.hd tableau) &&
    check_tree (mk_eq_constraints tableau) upper_bounds lower_bounds
      relu_constraints proof_tree
==>
  unsat tableau upper_bounds lower_bounds relu_constraints x
```

# Soundness of ReLU Splits



If the split corresponds to a ReLU constraint in  $Q$ , and both  $Q_L$  and  $Q_R$  are UNSAT, so is  $Q$ .

# Soundness of ReLU Splits

```
theorem soundness_relu_split_matching tableau us ls constraints xs split =  
  let (lb_left, ub_left), (lb_right, ub_right) = update_bounds_from_split ls  
    us split in  
  List.length xs = List.length ls  
  && List.length xs = List.length us  
  ==>  
  match split with  
  | ReLUSplit (b,f,aux) ->  
    List.mem (Relu (b,f,aux)) constraints  
    && sat tableau us ls constraints xs  
    ==>  
    (sat tableau ub_left lb_left constraints xs  
     || sat tableau ub_right lb_right constraints xs)  
  | _ -> true
```

If the split corresponds to a ReLU constraint in  $Q$ , and both  $Q_L$  and  $Q_R$  are UNSAT, so is  $Q$ .

# Main Result

```
let rec check_tree tableau upper_bounds lower_bounds constraints
proof_node =
match proof_node with
| Proof_tree.Leaf contradiction ->
    check_contradiction ...
| Proof_tree.Node ( split, left, right ) ->
    let valid_split = ... in
    let valid_children = ... in
    valid_split && valid_children
```

# Main Result

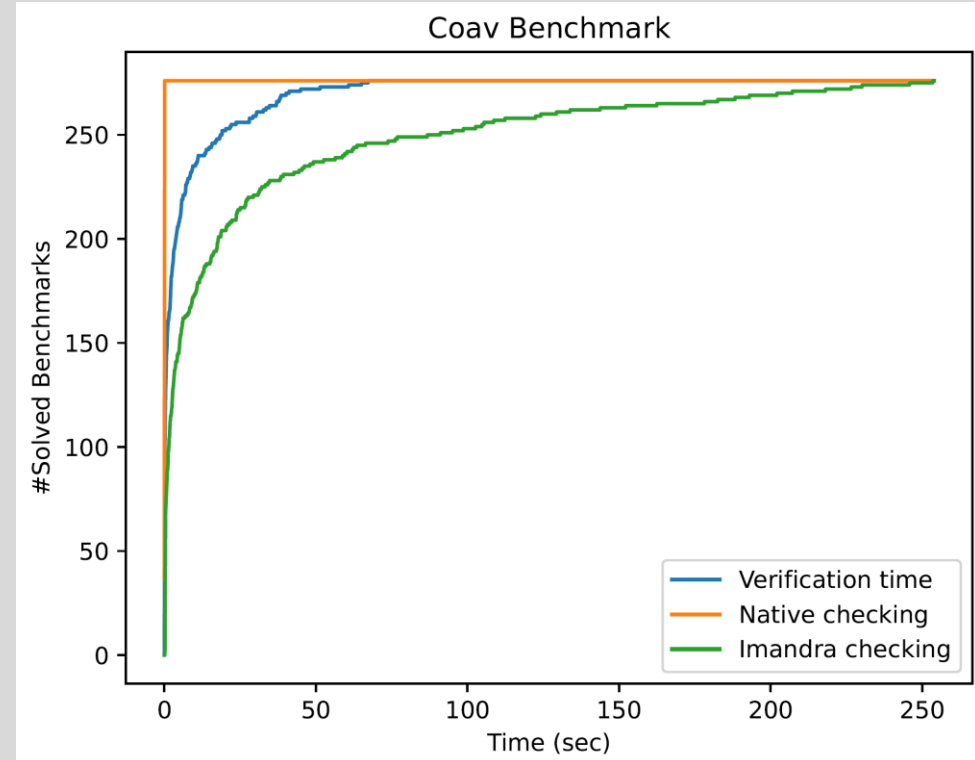

```
let rec check_tree tableau upper_bounds lower_bounds constraints  
proof_node =
```

```
theorem check_tree_soundness (tableau: real list list) (upper_bounds: real list)  
(lower_bounds: real list)  
  (constraints: Constraint.t list) (tree: Proof_tree.t) (x: real list) =  
  valid_proof tableau upper_bounds lower_bounds constraints tree  
  && well_formed_vector tableau x  
=>  
  unsat tableau upper_bounds lower_bounds constraints x
```

```
let valid_children = ... in  
valid_split && valid_children
```



# Performance Evaluation




Evaluation shows a  $\times \sim 4.7$  checking-to-verification delay, suggesting a clear tradeoff between reliability and scalability.

# L.o.C. Evaluation

Result	Module name	L.O.C	Aux. Lem.	Library Dependencies (accumulating)
Poly Farkas lemma (Theorem 5)	farkas.iml	194	21	Imandra Standard Libraries: Real, List, Polynomials
Sound application of DNN polynomial	well_formed_reduction.iml	41	13	farkas.iml, certificate.iml,
Farkas lemma (Theorem 7)	bound_reduction.iml	359		arithmetic.iml, util.iml,
	tableau_reduction.iml	356		tightening.iml, constraint.iml, proof_tree.iml, checker.iml, bound_reduct_g.iml, mk_bound_poly.iml
Soundness of leaf checking (Lemma 10)	leaf_soundness.iml	145	40	sat.iml, split.iml bound_reduction.iml, well_formed_reduction.iml, tableau_reduction.iml
Single variable splits are covering (Lemma 12)	single_var_split_soundness.iml	81	10	
ReLU splits are covering (Lemma 14)	relu_split_soundness.iml	113	18	relu.iml
	relu_case_1_bounded.iml	337		
	relu_case_2_bounded.iml	338		
Soundness of node checking	node_soundness.iml	78	19	relu_split_soundness.iml, sin- gle_var_split_soundness.iml
Soundness (Theorem 15)	checker_soundness.iml	71	146	leaf_soundness.iml, node_soundness.iml
<b>Total:</b>		<b>2113</b>	<b>267</b>	

■ **Table 2** Summary of the entire formalisation. The Table reads as follows: a result  $A$  is proven in module  $A\_mod$ , which is  $N$  lines long, calls  $M$  auxiliary lemmas, and depends on libraries as listed.


# Future Work



**Integrate** in a  
system verifier

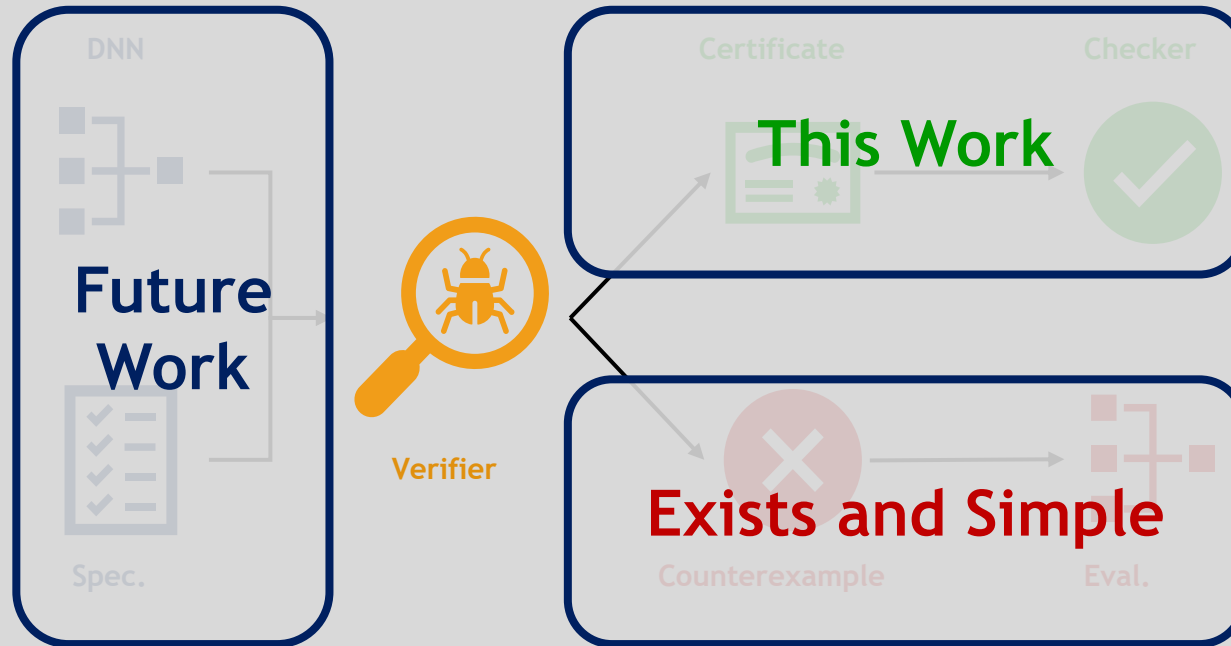


**Support** certificate  
optimizations



**Bridge gap** between  
DNN and query

# Conclusion



Feel **free to contact:**  
[omri.isac@mail.huji.ac.il](mailto:omri.isac@mail.huji.ac.il)



paper + code **available**