



# Gödel Mirror

**A Paraconsistent Calculus that Metabolizes Contradictions Mechanized in Lean 4**

**Jhet Chan**

Independent Researcher, Malaysia

**ITP 2025 Lean Workshop**

2 October 2025

Reykjavík, Iceland





# Motivation

## Contradictions are Everywhere

- **Large Language Models (LLMs)** often contradict themselves.
- In classical logic, a contradiction ( $P \wedge \neg P$ ) leads to **explosion**, anything can be proven.
- **Gödel's Incompleteness Theorems** showed that paradox is inevitable in any system powerful enough for self-reference.

## Our Question:

Instead of letting contradictions explode, can we **metabolize** them as part of the computation?



# The Gödel Mirror Mechanism

## From Explosion to a Controlled Cycle

The Gödel Mirror turns a contradictions into a stable, structured object through a three-step cycle.

- **Paradox:** A self-referential term is identified.
- **cap(...):** The paradox is encapsulated, preventing explosion.
- **enter(...):** The capsule re-enters the system in a controlled way.
- **node(...):** It stabilizes into a structured "node", a value.

The key idea:

**Contradiction becomes a stable object in computation (read: if  $P \wedge \neg P \rightarrow$  recursion)**



## Encoding in Lean

**The entire system is captured in a simple inductive data type.**

```
inductive MirrorSystem where
  | base
  | node    : MirrorSystem → MirrorSystem
  | self_ref -- Represents a paradoxical term
  | cap     : MirrorSystem → MirrorSystem
  | enter   : MirrorSystem → MirrorSystem
  | named   : String → MirrorSystem → MirrorSystem
```



## Live Demo: The Liar Paradox (Demo.lean)

Let's define the Liar paradox: "This statement is false", represented by a named self-reference.


```
def liar := MirrorSystem.named "Liar" MirrorSystem.self_ref
```

```
-- Run it for 3 steps
```













```
#eval run liar 3
```

**Demo repo (branch 'itp'):**

<https://github.com/jhetchan/godel-mirror/tree/itp>















## Demo 1: Liar Paradox (named "Liar" self\_ref)

Step 0:	 'Liar'	← Start: paradoxical term
Step 1:	 cap(  'Liar')	← Encapsulated (paradox wrapped)
Step 2:	 enter(  cap(  'Liar'))	← Reentering
Step 3:	 node(  enter(  cap(  'Liar')))	← Stabilized as node
Step 4:	 node(  node(...))	← Keeps wrapping in nodes

**Starting point:** named "Liar" self\_ref

**State:** Paradox (self-referential statement)













## Demo 1: Liar Paradox (named "Liar" self\_ref)

Step 0:	 'Liar'	← Start: paradoxical term
Step 1:	 cap(  'Liar')	← Encapsulated (paradox wrapped)
Step 2:	 enter(  cap(  'Liar'))	← Reentering
Step 3:	 node(  enter(  cap(  'Liar')))	← Stabilized as node
Step 4:	 node(  node(...))	← Keeps wrapping in nodes

**Transformation:** Paradox detected → Encapsulated

**Rule applied:**  $\text{classify}(\text{liar}) = \text{Paradox} \rightarrow \text{step}(\text{liar}) = \text{cap}(\text{liar})$

## Demo 1: Liar Paradox (named "Liar" self\_ref)













Step 0:	 'Liar'	← Start: paradoxical term
Step 1:	 cap(  'Liar')	← Encapsulated (paradox wrapped)
Step 2:	 enter(  cap(  'Liar'))	← Reentering
Step 3:	 node(  enter(  cap(  'Liar')))	← Stabilized as node
Step 4:	 node(  node(...))	← Keeps wrapping in nodes

**Transformation:** Capsule reenters the system

**Rule applied:**  $\text{classify}(\text{cap}(\text{liar})) = \text{Integrate} \rightarrow \text{step} = \text{enter}(\text{cap}(\text{liar}))$















## Demo 1: Liar Paradox (named "Liar" self\_ref)

Step 0:	 'Liar'	← Start: paradoxical term
Step 1:	 cap(  'Liar')	← Encapsulated (paradox wrapped)
Step 2:	 enter(  cap(  'Liar'))	← Reentering
Step 3:	 node(  enter(  cap(  'Liar')))	← Stabilized as node
Step 4:	 node(  node(...))	← Keeps wrapping in nodes

The Liar paradox goes through the **3-step cycle**: paradox → cap → enter → node, then continues wrapping in more nodes (stable state).


## Demo 1: Liar Paradox (named "Liar" self\_ref)



Step 0:	 'Liar'	← Start: paradoxical term
Step 1:	 cap(  'Liar')	← Encapsulated (paradox wrapped)
Step 2:	 enter(  cap(  'Liar'))	← Reentering
Step 3:	 node(  enter(  cap(  'Liar')))	← Stabilized as node
Step 4:	 node(  node(...))	← Keeps wrapping in nodes




**Transformation:** Stabilizes as a node





**Rule applied:**  $\text{classify}(\text{enter}(\text{cap}(\text{liar}))) = \text{Reentry} \rightarrow \text{step} = \text{node}(\dots)$



## Demo 2: Nested paradox processes outer-first (fuel = 4) - Line 39

Step 0:  'Outer' ← named "Outer" (named "Inner" self\_ref)

Step 1:  cap( 'Outer') ← System detects the OUTER label

Step 2:  enter( cap( 'Outer'))




Step 3:  node( enter( cap( 'Outer')))

Step 4:  node( node(...)) ← Now in stable "node wrapping" phase

### Key insight:

Even though there's a nested `named "Inner" self_ref` inside, the system processes the **outermost name first**.  
**Step 4 is extra:** After reaching `node(...)`, any further steps just keep wrapping in more nodes (stable absorption).

## The 3-Step Cycle as Structural Invariant

1. Detect paradox →  Encapsulate (cap)
2. Encapsulated paradox →  Reentry (enter)
3. Reentering capsule →  Stabilize (node)



## Key Properties Demonstrated

- Deterministic: Always the same transformation
- Controlled: No logical explosion
- Universal: Works for simple, named, and nested paradoxes
- Verified: Assertions prove the expected behavior



# Formal Verification



## Completion (Examples.lean)

These use `completeFuel` which tries to resolve ALL paradoxes recursively in one go (not step-by-step).

### Completion Example 1: Liar

Before: <code>named 'Liar' (self_ref)</code>	← Original paradox
After: <code>node (enter (cap (named 'Liar' (self_ref))))</code>	← Wrapped once
Contains paradox (before): <code>true</code>	← Has <code>self_ref</code>
Contains paradox (after): <code>true</code>	← STILL has <code>self_ref</code> inside!

**Why still paradox?** The `completeFuel` function wraps the *entire term* once as `node (enter (cap (...)))`, but the inner `self_ref` is still there. The wrapping doesn't eliminate the paradox, it just "contains" it.



## Completion Example 2: Nested Paradox (`cap(self_ref)`)

```
Before: cap(self_ref)           ← Paradox already inside cap
After:  cap(node(enter(cap(self_ref)))) ← Inner self_ref resolved
Contains paradox (before): true
Contains paradox (after): true   ← Outer cap still wraps a paradox
```

**What happened?** `completeFuel` detected `self_ref` inside the `cap`, wrapped it as `node(enter(cap(self_ref)))`, then put that back inside the outer `cap`.  
The structure changed but `self_ref` is still there.





## Key Insights

1. **step (step-by-step)**: Processes the **outermost** paradox one step at a time
  - Paradox → cap → enter → node → (keeps wrapping in nodes)
2. **completeFuel (recursive completion)**: Tries to resolve **all** paradoxes at once
  - Wraps paradoxes as node(enter(cap(...)))
  - BUT: self\_ref itself never disappears - it's the atomic paradox
  - The wrapping just "stabilizes" the structure



## Key Insights

3. Why "Contains paradox: true" after completion?

- self\_ref is the **irreducible paradox** - you can't eliminate it
- You can only wrap/contain it in stable structure
- contains\_paradox checks if self\_ref appears **anywhere** in the tree
- Even after wrapping, self\_ref is still there (just deeper in the tree)

4. **This is by design:** The Gödel Mirror doesn't "solve" paradoxes, it **metabolizes** them into stable structures that don't explode the logic. The paradox is still there, but controlled.



## Three Proven Theorems

- ❑ **Progress:** Every non-value term can take a computational **step**. The system never gets stuck.
  - ❑ For any term  $t \in \text{MirrorSystem}$ , either  $t$  is a value (i.e.,  $t = \text{base}$ ) or there exists a term  $t'$  such that  $t \rightarrow t'$ .
- ❑ **Non-Explosion:** Any paradoxical term resolves deterministically to a stable **node** in exactly **3** steps.
  - ❑  $t \rightarrow \text{cap}(t) \rightarrow \text{enter}(t) \rightarrow \text{node}(t)$
- ❑ **Label Preservation:** A **named** structure keeps its label throughout the reduction process.

For proofs, can see `MetaTheory.lean` in repo.



# Takeaways



## Why Lean 4?

Lean provided a complete toolkit:

1. **Precision:** Every rule and definition is type-checked. No ambiguity.
2. **Execution:** The fuelled evaluator allowed for rapid testing and live demos.
3. **Proofs:** I could formally verify critical properties like non-explosion.

Without Lean, the Gödel Mirror would be a conceptual sketch. **With Lean, it is a verifiable system.**



# What Was Hard

## The Challenge: Non-Termination

Lean's functions must be **terminating** (total) to ensure logical consistency. But the Gödel Mirror is deliberately cyclic and non-terminating.

## The Solution: A Two-Pronged Approach

- **For Demo:** Use a **fuelled evaluator** (`run n`). This is a total function that simply stops after `n` steps, making it easy to execute.
- **For Proof:** Define the semantics as a **relation** (`step : MirrorSystem → MirrorSystem → Prop`). This allows reasoning about infinite computations without writing non-terminating functions. (still work-in-progress)

Lean's flexibility handled both the executable spec and the meta-theory.



## Why It Matters for Lean Community

Takeaways from Gödel Mirror	Takeaways for Lean Community
<b>Modeling Non-Terminating Calculi</b> Lean can formalize systems that go beyond strong normalization.	<b>Dedicated TRS Library Needed</b> Isabelle/HOL has a mature rewriting library; Lean doesn't yet.
<b>Fuel + Relational Semantics Pattern</b> Combining executable fuel evaluators with relational proofs is a practical technique for non-terminating semantics.	<b>Automation:</b> Avoid reproving confluence/termination lemmas manually. <b>Accessibility:</b> Lower barrier for CS researchers and students. <b>Advanced Topics:</b> Enables systematic work on infinitary and non-terminating rewriting.
<b>Lean as Executable Spec Lab</b> Beyond theorem proving, Lean works as a design studio for new formal systems, where you can prototype, run, <i>and</i> prove in one place.	<b>Path Forward</b> Like mathlib for mathematicians. It'd be great a TRS library for PL/rewriting researchers too.



## Conclusion

1. Contradictions don't have to explode, they can be **metabolized** into structured computational nodes.

Paradox → cap → enter → node

2. **Lean 4 made it real.** The only reason I could test, run, and prove this as an outsider was because Lean integrates precision, execution, and proof in one place.





# Gödel Mirror

**A Paraconsistent Calculus that Metabolizes Contradictions Mechanized in Lean 4**

**Demo repo (branch 'itp'):**

<https://github.com/jhetchan/godel-mirror/tree/itp>

**Jhet Chan**

<https://jhetchan.com>

Independent Researcher, Malaysia





# Questions