


Certifying rlive: A New Proof Strategy for Liveness Model Checking

Giulia Sindoni¹, Alberto Griggio¹, and Stefano Tonetta¹

¹Fondazione Bruno Kessler, Trento, IT

FroCoS 2025, 1 October 2025


Certifying Model Checking (CMC)



- Trust in verification results for safety-critical systems.
- Certification: extend standard MC with certificates: evidence validating the (yes) answer.
- What serves as evidence? A deductive proof.

The Liveness Checking Problem

- Finite state transition systems $\mathcal{M} = \langle I, T, V \rangle$.
- Liveness checking problem: $\mathcal{M} \models \mathbf{FG}q$: for all paths of \mathcal{M} , q eventually holds in all the future states



- Counterexample: an infinite path where $\neg q$ is visited infinitely often ($\mathbf{GF}\neg q$). Finite states: if the property is violated, there always exists a lasso-shaped counterexample.

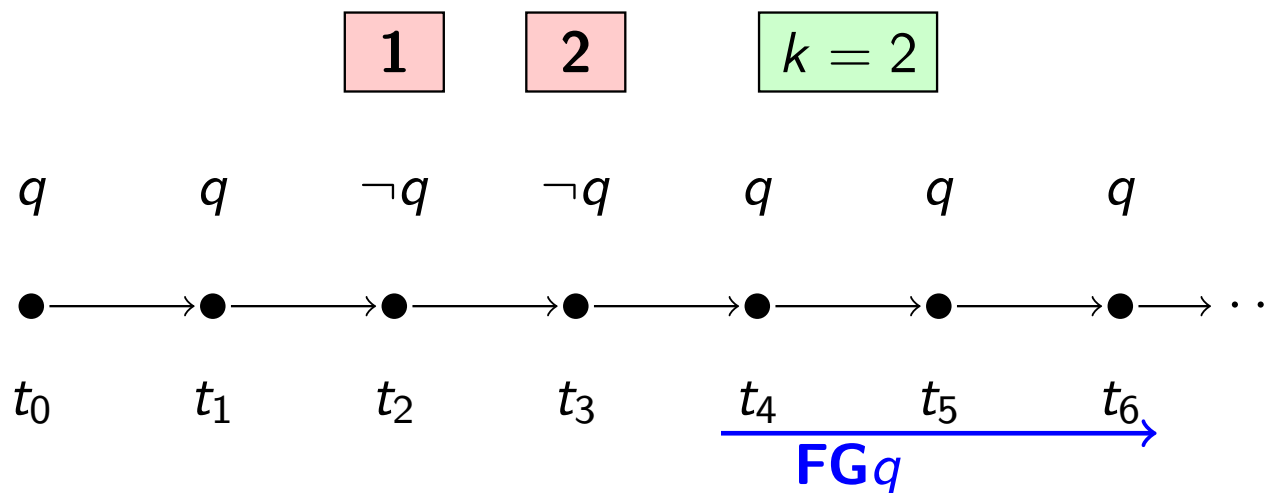


K-liveness: a Liveness Algorithm that Counts (Claessen and Sörensson, 2012)

Key Insight

For any valid liveness property $\mathbf{FG}q$ in a finite-state system, there exists a bound k such that $\neg q$ can become true at most k times in any trace.

Count $\neg q$ occurrences

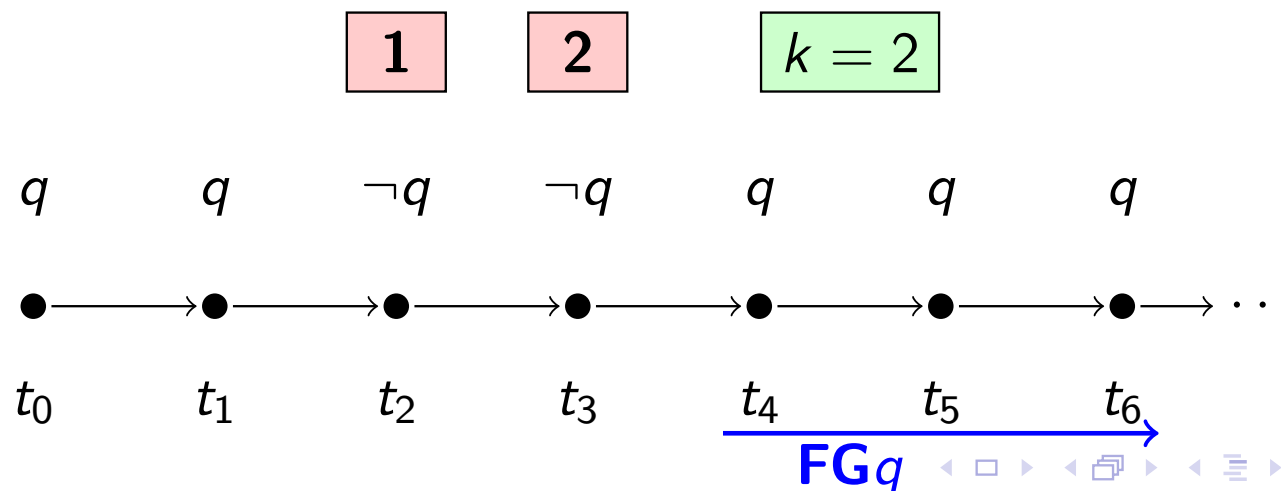


K-liveness: a Liveness Algorithm that Counts

Algorithm Idea

- 1 Start with $k = 0$
- 2 Try to prove: $\neg q$ occurs at most k times
- 3 If successful \Rightarrow property holds
- 4 If failed, increment k and repeat
- 5 Each iteration is a safety check

Count $\neg q$ occurrences



rlive: Avoiding the Shoals (Xia et al., 2024)

Key Insight

Builds counterexamples to **FG** q incrementally through a recursive, depth-first search process.


Shoals

Sets of states that can reach $\neg q$ only finitely many times. When search reaches a dead-end, safety checker provides inductive invariant C representing a newly discovered shoal.

rlive: Shoals Block Future Searches

Algorithm:

- ① Find path from initial states to $\neg q$ states
 - ② Continue searching from successors of each $\neg q$ state
 - ③ Either a previously visited $\neg q$ -state is met again, creating a lasso-shaped counterexample. Or no more $\neg q$ -states can be reached: a shoal is obtained.
- Each shoal blocks part of state space
 - Add constraint $\neg C \wedge \neg C'$ to transition relation
 - Future searches exclude shoal states



Certifying rlive: From the Algorithm to the Temporal Deductive rule

rlive Algorithm

```

1 Procedure rlive( $X, I, T, FGq$ ) begin
2    $C := \perp$ 
3    $B :=$  empty stack of states
4   while  $check\_invariant(X, I, T \wedge (\neg C \wedge \neg C'), T^{-1}(\neg C) \rightarrow q)$  is Unsafe do
5      $s :=$  final state of  $get\_counterexample()$ 
6      $B.push(s)$ 
7     while  $B$  is not empty do
8        $s := B.top()$ 
9       if  $check\_invariant(X, T(s), T \wedge (\neg C \wedge \neg C'), T^{-1}(\neg C) \rightarrow q)$  is
         Unsafe then
10         $t :=$  final state of  $get\_counterexample()$ 
11        if  $t \in B$  then
12          return Unsafe
13         $B.push(t)$ 
14      else
15         $inv := get\_inductive\_invariant()$ 
16         $C := C \vee inv$ 
17         $B.pop()$ 
18   return Safe
    
```

 \Rightarrow

RL rule

$$P_{init} := (I \wedge \mathbf{GT}) \rightarrow \mathbf{FC} \vee \mathbf{G}q$$

$$P_0 := \mathbf{G}(C_0 \leftrightarrow \perp)$$

$$Pk_1 := \mathbf{G}((C_0 \vee C_1) \wedge \mathcal{T} \rightarrow \mathbf{X}(C_0 \vee C_1))$$

$$Pp_1 := \mathbf{G}((C_0 \vee C_1) \wedge \mathcal{T} \wedge \neg q \rightarrow \mathbf{X}(C_0))$$

 \vdots

$$Pk_n := \mathbf{G}((C_0 \vee \dots \vee C_n) \wedge \mathcal{T} \rightarrow \mathbf{X}(C_0 \vee \dots \vee C_n))$$

$$Pp_n := \mathbf{G}((C_0 \vee \dots \vee C_n) \wedge \mathcal{T} \wedge \neg q \rightarrow$$

$$\mathbf{X}(C_0 \vee \dots \vee C_{n-1}))$$

$$\frac{\text{RL}}{I \wedge \mathbf{G}(\mathcal{T}) \rightarrow \mathbf{FG}q} = \mathcal{M} \models \mathbf{FG}q$$

where $C := C_0 \vee C_1 \vee \dots \vee C_n$ is the final set of discovered shoals.

Certifying rlive: the Temporal Deductive rule

- $P_{\text{init}} := (\mathcal{I} \wedge \mathbf{G}\mathcal{T}) \rightarrow \mathbf{F}C \vee \mathbf{G}q$: on any trace either the shoal is eventually entered, or q is an invariant.
- $P_0 := \mathbf{G}(C_0 \leftrightarrow \perp)$: the shoal is empty initially.
- $Pk_i := \mathbf{G}((C_0 \vee \dots \vee C_i) \wedge \mathcal{T} \rightarrow \mathbf{X}(C_0 \vee \dots \vee C_i))$ the invariant C incrementally built is inductive.
- $Pp_i := \mathbf{G}((C_0 \vee \dots \vee C_i) \wedge \mathcal{T} \wedge \neg q \rightarrow \mathbf{X}(C_0 \vee \dots \vee C_{i-1}))$ the search space can be incrementally restricted, as long as we keep visiting a new $\neg q$ -state.

Correctness and Completeness of the RL Rule

Proof of Correctness: by Contradiction

Correctness of rule RL **formally proven** in the theorem prover. The main step is:

$$\text{RLB} \frac{\text{FC} \quad \Pi \quad [\mathbf{GF}\neg q]}{\frac{\text{FC}_0 \quad P_0 := \mathbf{G}\neg C_0}{\frac{\perp}{\mathbf{FG}q}}}$$

where $\Pi = \{Pk_1, Pp_1, \dots, Pk_n, Pp_n\}$

Proof of Completeness

If the algorithm rlive succeeds in establishing the liveness property, then it generates shoals $C = C_0 \vee \dots \vee C_n$ such that the premises of the temporal deductive rule RL are true. So the model will satisfy the necessary premises for RL to be applied successfully.

RL as a Generalization of k-liveness rule

Shared Intuition

Both algorithms prove the same fundamental property: $\neg q$ can occur at most finitely many times in any trace of a finite-state system.

The Generalization

Key insight: In rule for k-liveness (Griggio et al. 2021) we have formulae (inductive invariants) $\alpha_0, \dots, \alpha_{k+1}$, that keep count of the number of times $\neg q$ is reached. There is a mapping $\alpha_i \mapsto C_{k-i+1}$ such that:

- RL rule it can be used to build proofs for k-liveness using this mapping.

The Proof Strategy for Liveness Checking

$$\frac{P_{\text{init}} \quad P_0 \quad Pk_1 \quad Pp_1 \quad \dots \quad Pk_n \quad Pp_n}{\mathcal{I} \wedge \mathbf{G}(\mathcal{T}) \rightarrow \mathbf{FG}q} \text{RL}$$

TP strategy

- ① Assume P_{init}, P_0 and Π ,
- ② Apply RL rule: thus we can conclude the goal: $\vdash \mathcal{I} \wedge \mathbf{G}(\mathcal{T}) \rightarrow \mathbf{FG}q$
- ③ Discharge of proof obligations: P_0, Π discharged using SAT solver.
- ④ Discharge of proof obligations:
 $P_{\text{init}} := \mathcal{I} \wedge \mathbf{GT} \rightarrow \mathbf{FC} \vee \mathbf{G}q = \mathcal{I} \wedge \mathbf{G}(\mathcal{T} \wedge \neg C) \rightarrow \mathbf{G}q$ invariant claim discharged using a subroutine for proving invariants.


Technical foundation: Formalising LTL in PVS

- PVS: specification language with integrated theorem prover.
- Interactive but also supports strategies developments.

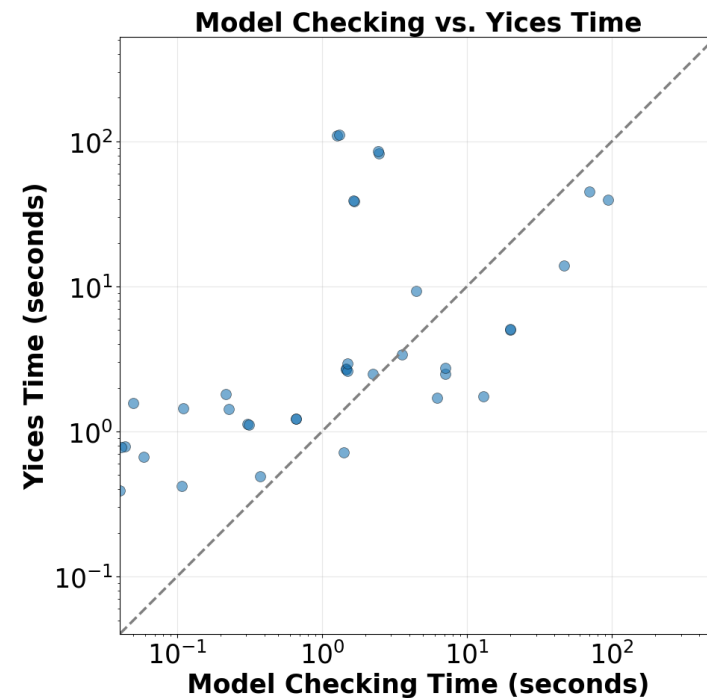

Shallow embedding of LTL

```
shallow_ltl[State: TYPE+]: THEORY
BEGIN
Trace: TYPE = ARRAY[nat -> State]
ltlformula: TYPE = [Trace -> [nat -> bool]]
...
NOT(P)(trace: Trace)(t: nat): bool = NOT P(trace)(t);
NEXT(P)(trace: Trace)(t: nat): bool = P(trace)(t+1);
...
valid(P): bool = FORALL (trace: Trace): P(trace)(0)
...
```

Certification Flow



Experimental Setup



- Benchmarks Source: Hardware Model Checking Competition.
- 53 problems tested, 41 successfully certified within time and memory limit.
- Demonstrates feasibility but highlights performance gap.
- Bottleneck: PVS internal bookkeeping and definition management.
- Insight: performance gap primarily due to theorem prover infrastructure

Achievements and Future Work

Key Contributions:

- Novel proof strategy for certifying liveness checking results.
- Despite the complexity of rlive, shoals provided by the model checker are sufficient to generate proofs.
- Minimal model checker modifications - only output shoal.
- Progress in CMC: distribute the trust across more fundamental principles and create redundancy that increases overall confidence.

Future Work:

- Extending certifying model checking approach to other liveness checking algorithms (liveness-to-safety, FAIR): a strategy that encompasses them all?
- Generalisations to the infinite-state transition systems and SMT.

Bibliography

- ① Claessen Koen, and Niklas Sörensson. “A liveness checking algorithm that counts.” Formal Methods in Computer-Aided Design (FMCAD). IEEE, 2012.
- ② Xia Yechuan, et al. “Avoiding the shoals: a new approach to liveness checking.” CAV24, 2024.
- ③ Alberto Griggio, Marco Roveri, and Stefano Tonetta. “Certifying proofs for SAT-based model checking.” Formal Methods in System Design 57.2 (2021): 178-210.
- ④ Giulia Sindoni, et al. “A Theorem Prover Based Approach for SAT-Based Model Checking Certification.” Automated Deduction–CADE 30, 2025.